

政大資科系

作業系統

Operating System

廖峻鋒

cfliao@nccu.edu.tw

Operating System

Processes

Chun-Feng Liao

廖峻鋒

Department of Computer Science

National Chengchi University

Outline

- Process Structure
- Process Scheduling
- Process Operation
- Process Communication

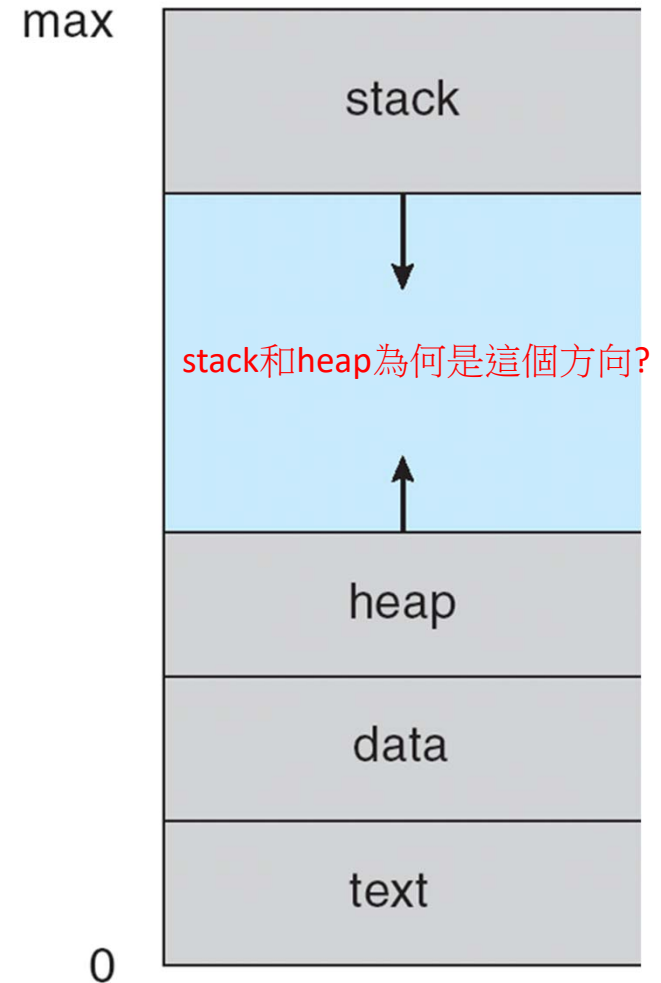
Process and Program

- Process is active
 - A program in execution
- Program is passive
 - A binary file stored in the disk
- Contents
 - Status: PC and other registers' values
 - Text (code)
 - Data (static / global variables)
 - Heap (動態配置的資料型態)
 - Stack (activation records)

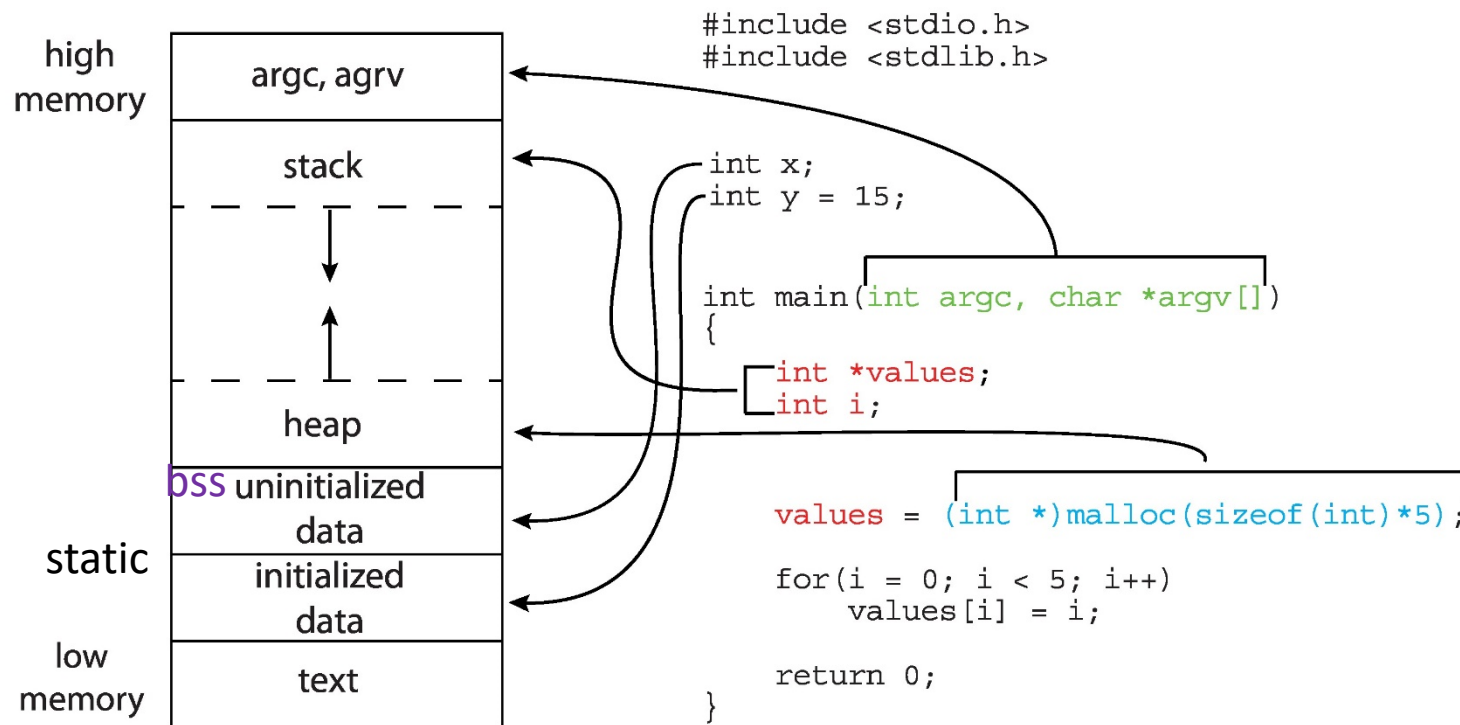


The activation record stores the parameters, local variables, return address of a function call

Memory layout of a process



Memory Layout of a C Program



Demo: size -A

Understanding the Memory Layout of Linux Executables

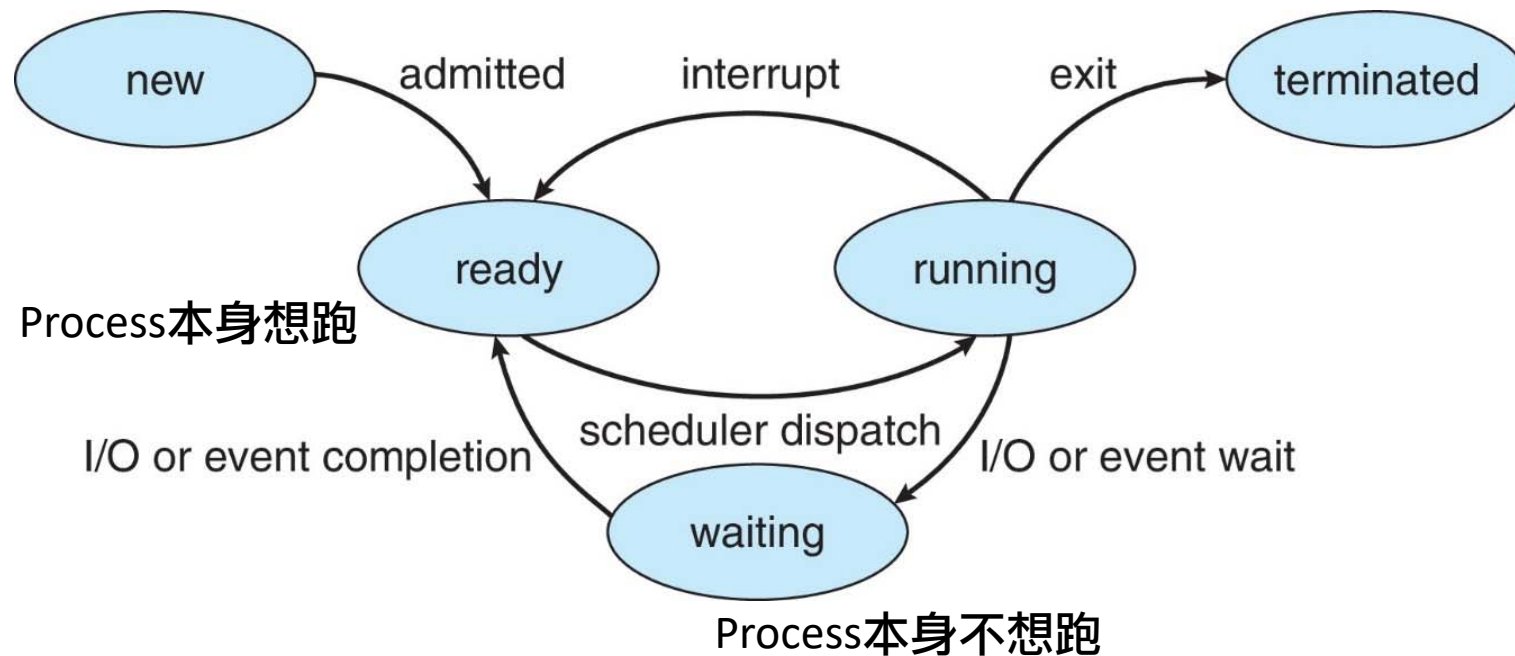
<https://gist.github.com/CMCDragonkai/10ab53654b2aa6ce55c11cfc5b2432a4>

Process state

- New
 - the process is being created
- Ready (沒在跑，但有意願跑)
 - the process is in the memory waiting to be assigned to a processor
- Running (正在跑)
 - instructions are being executed by CPU
- Waiting (沒在跑，也沒意願跑)
 - the process is waiting for events to occur
- Terminated
 - the process has finished execution

Process state

- One process is running on a processor at any instant
- Many processes may be ready or waiting



Process Control Block (PCB)

- Kernel中維護的Process資訊，每個Process都有一個PCB
 - Process state
 - Program counter Linux中稱為EIP (instruction pointer)
 - CPU registers
 - CPU scheduling information
(e.g. priority)
 - Memory-management information
(e.g. base/limit register)
 - I/O status information
 - Accounting information



See `/proc/process id`

在`/proc/<pid>/stat`中, see: <https://man7.org/linux/man-pages/man5/proc.5.html> (search:/proc/[pid]/stat)

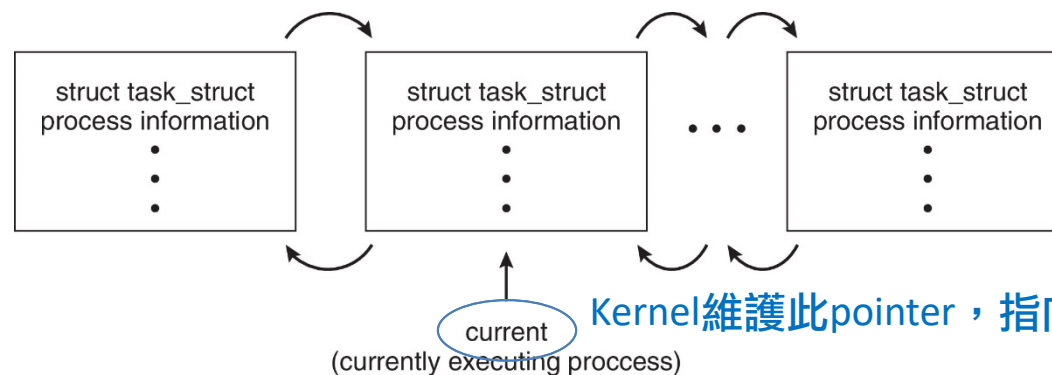
PCB in Linux

Represented by the structure `task_struct`

```

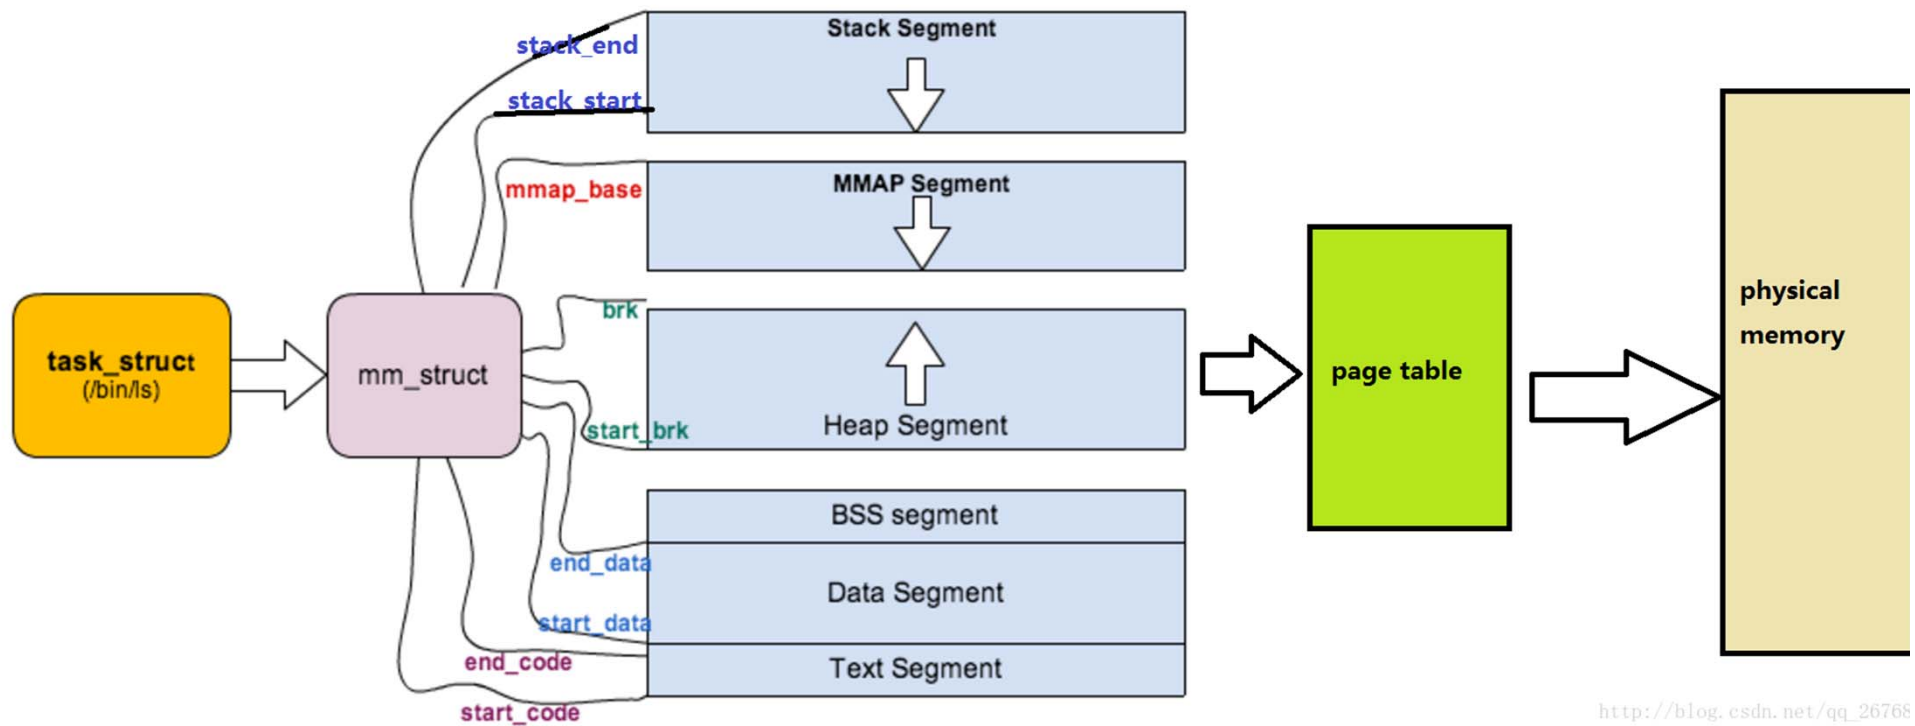
pid t_pid;                /* process identifier */
long state;               /* state of the process */
unsigned int time_slice   /* scheduling information */
struct task_struct *parent; /* this process' s parent */
struct list_head children; /* this process' s children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;     /* address space of this process */
                          (#L876)  指向Process實際存在的記憶體位址

```



Kernel維護此pointer，指向正在執行的process

<https://github.com/torvalds/linux/blob/master/include/linux/sched.h#L743>

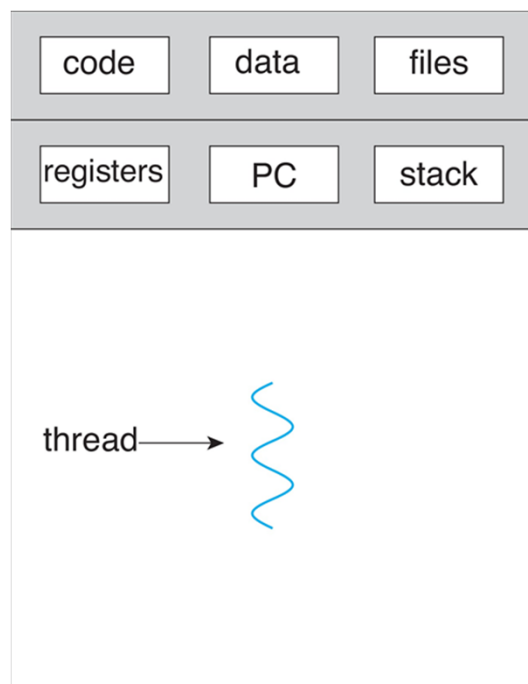


http://blog.csdn.net/qq_26768

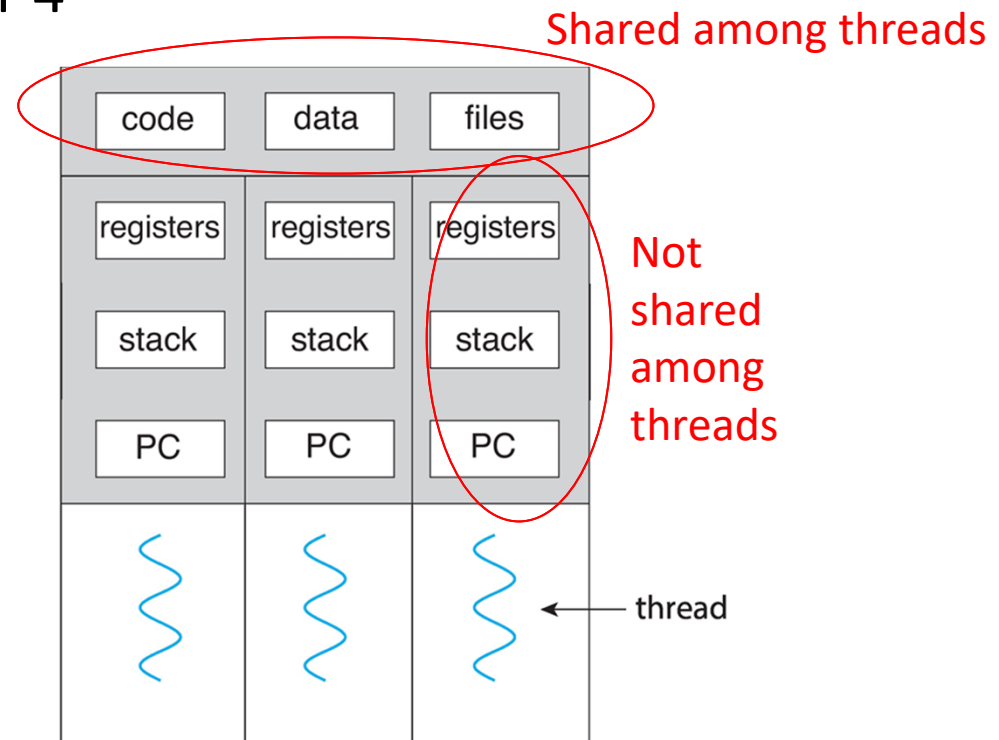
https://github.com/torvalds/linux/blob/master/include/linux/mm_types.h

Thread

- A process may have multiple PC and register states
 - Each represent a single thread of execution
 - Cover more on chapter 4



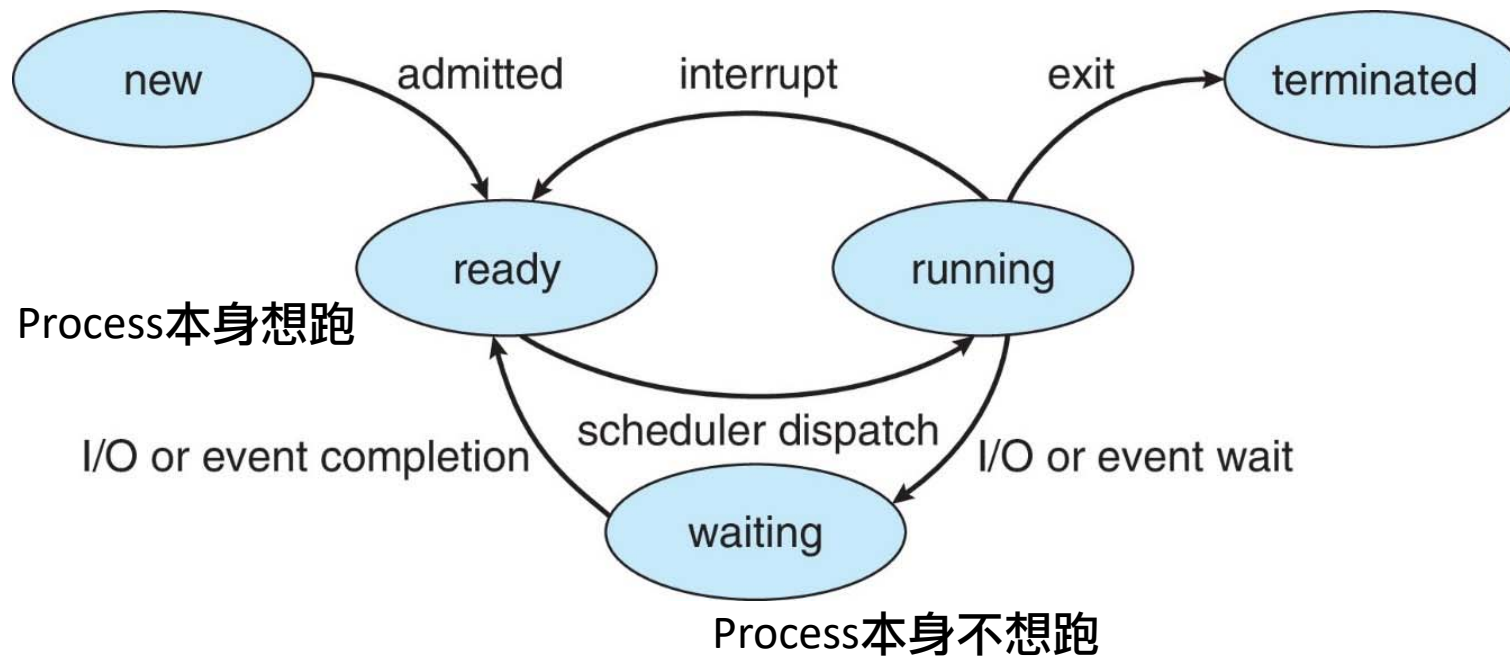
single-threaded process



multithreaded process

Process Scheduling

- Multiprogramming
 - CPU runs process at all times to maximize CPU utilization
 - Degree of multiprogramming: how many process in memory
- Time sharing
 - Switch CPU frequently such that users can interact with each program while it is running
- I/O bound vs. CPU bound
 - Most of the process running time spent on I/O or computation

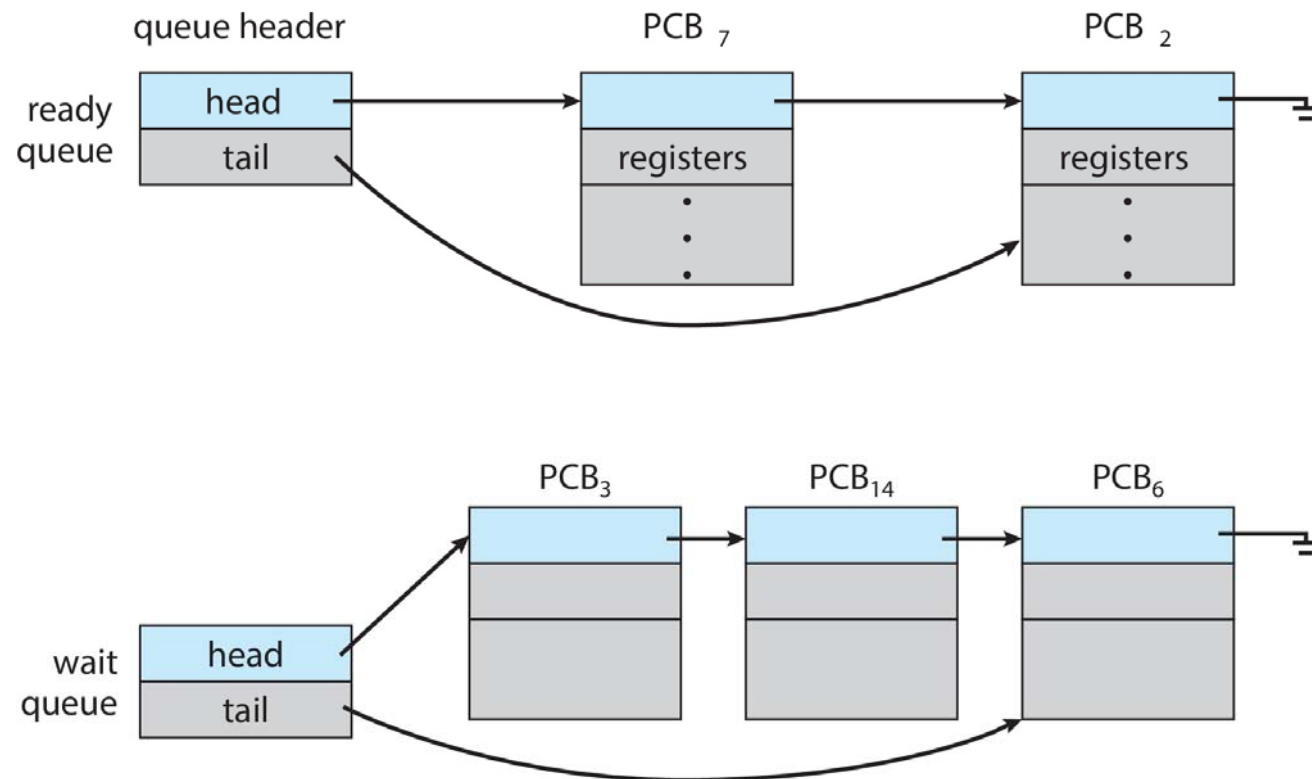


I/O bound的process可能花很多時間在這裡

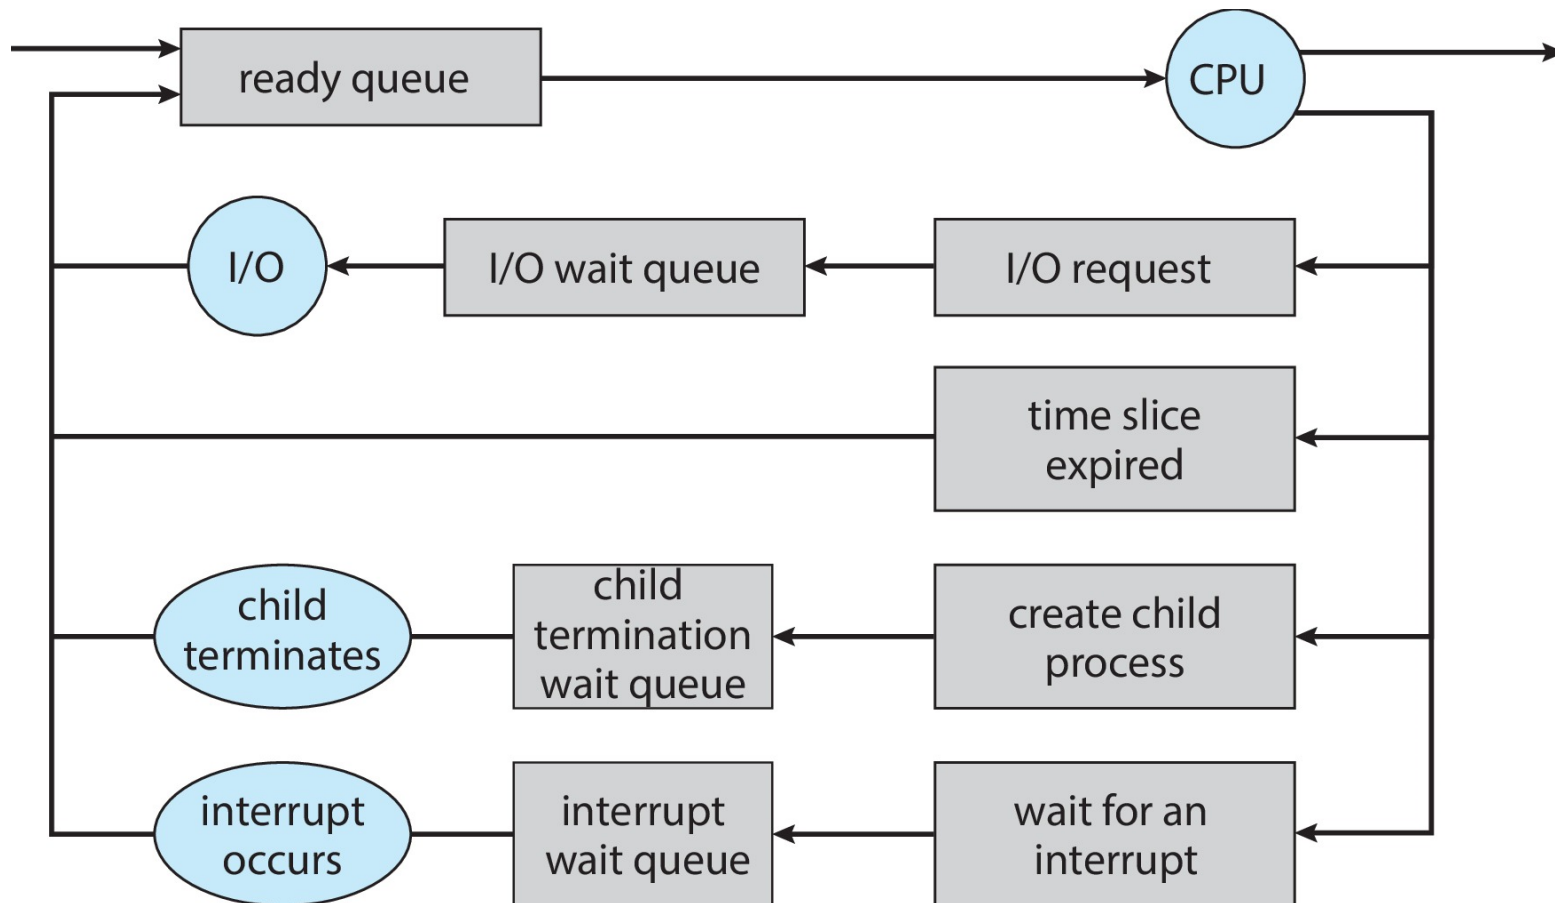
Process Scheduling Queues

- Processes migrate between the various queues (i.e. switch among states)
 - Job queue (New State) – set of all processes in the system
 - Ready queue (Ready State) – set of all processes residing in main memory, ready and waiting to execute
 - Device queue (Wait State)– set of processes waiting for an I/O device

Ready and Wait Queues



Process Scheduling Diagram



Note:此圖顯示了task移出CPU的四種情況

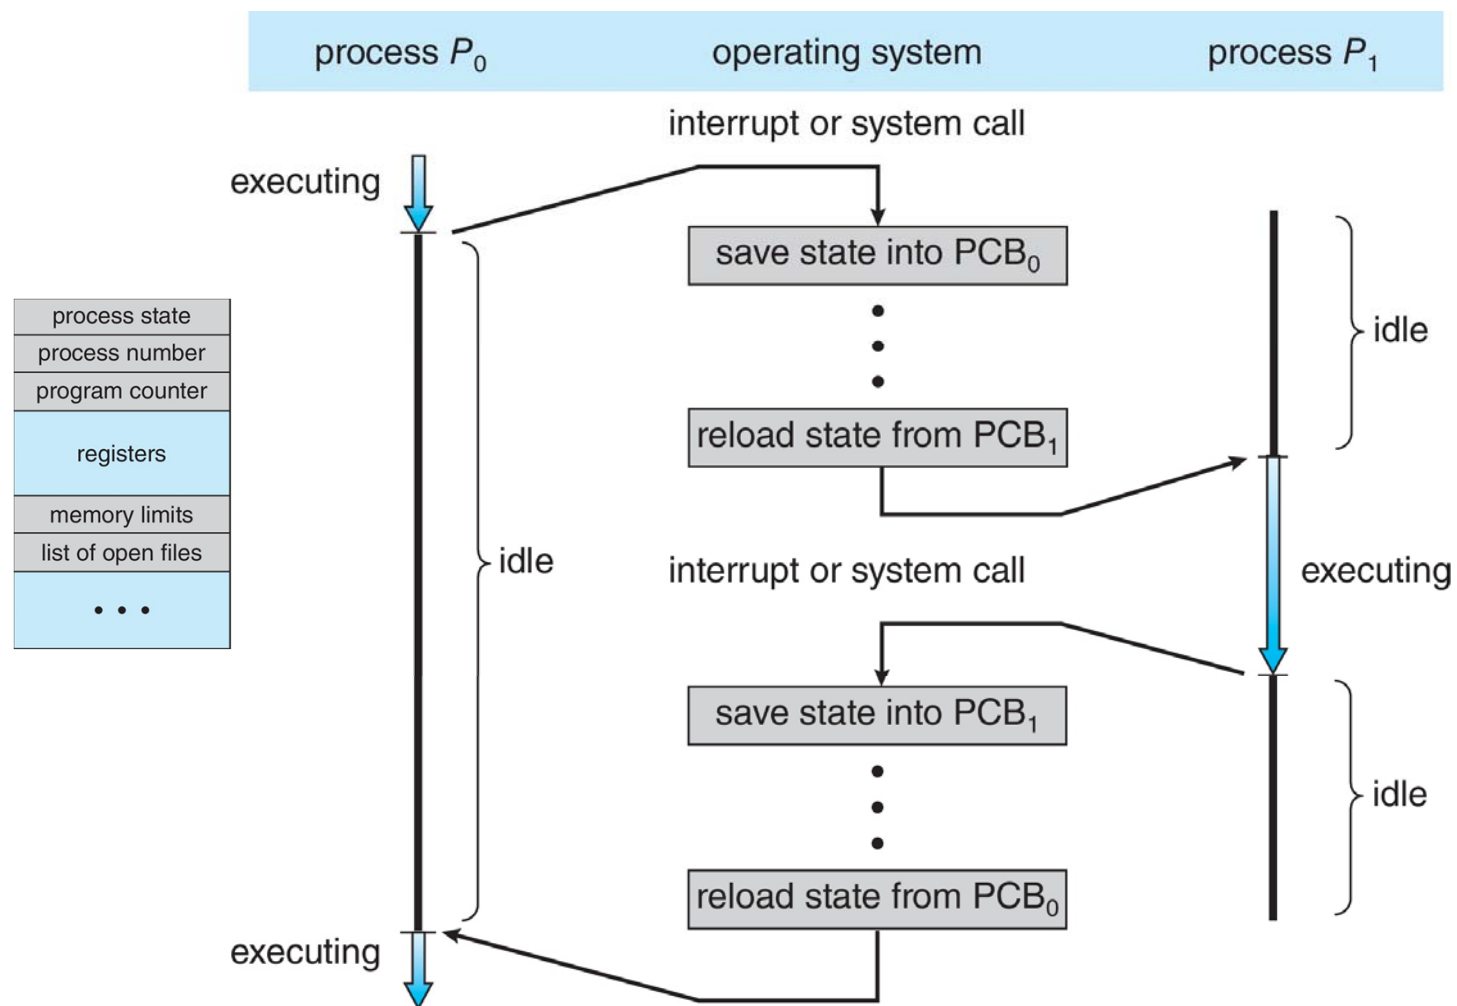
CPU Scheduling

- The role of the CPU scheduler
 - Select among the processes that are in the ready queue and allocate a CPU core to one of them
 - More on chapter 5
- Swapping
 - an intermediate form of scheduling
 - remove a process from memory → reduce the degree of multiprogramming
 - More on chapter 9

Context Switch

- Context Switch
 - Saves the state of the old process
 - Loads the saved state for the new process
 - Context-switch time is purely overhead
- Switch time (about 1~1000 ms) depends on
 - Memory speed
 - Number of registers
 - Existence of special instructions
 - a single instruction to save/load all registers
- Hardware support
 - Prepare multiple sets of registers
 - A context switch only changes the pointer to the register

Context Switch

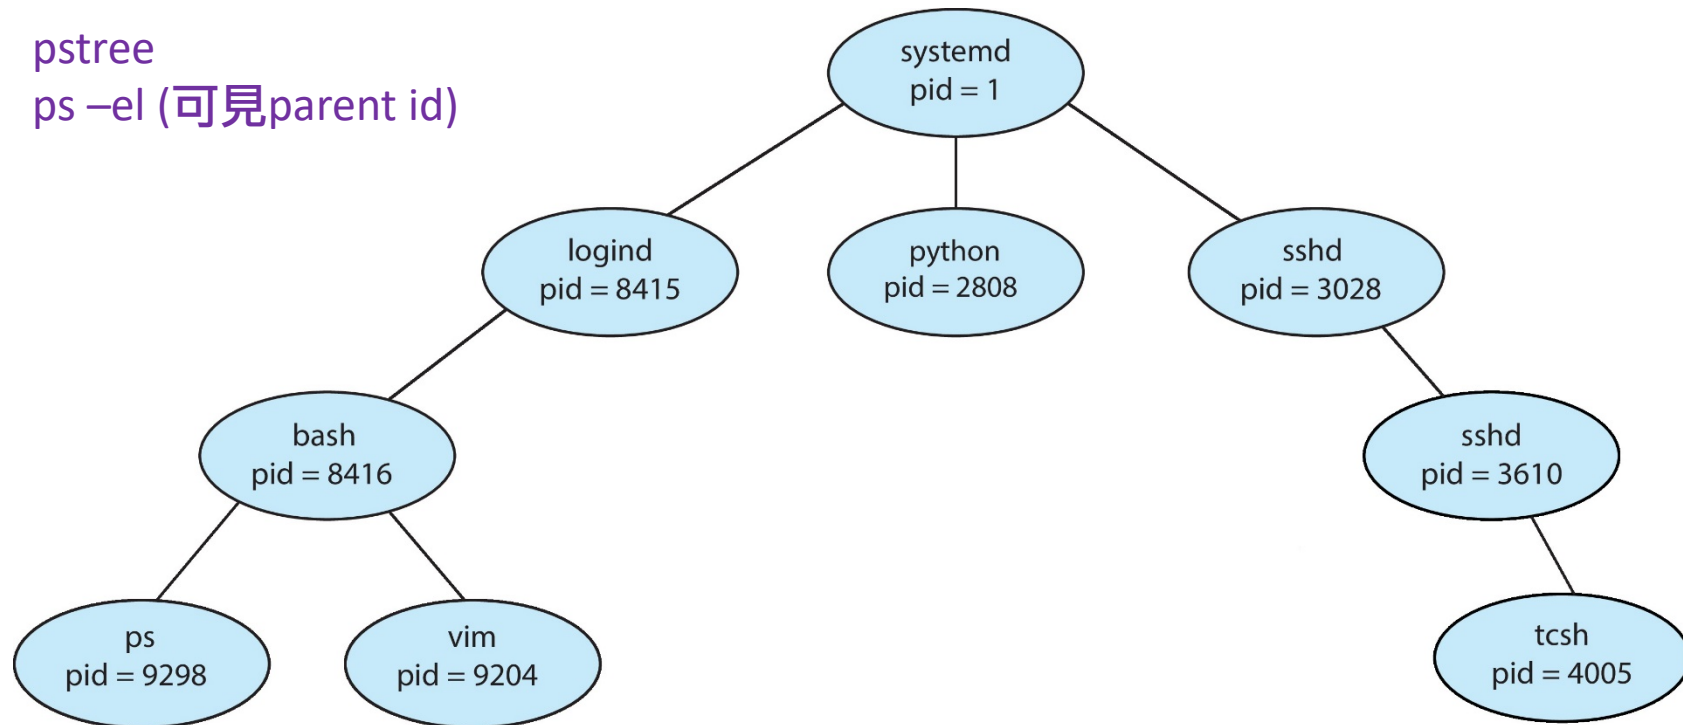


Process Tree

- Each process is identified by a unique process identifier (pid)
- Each process can create new child processes

pstree

ps -el (可見parent id)

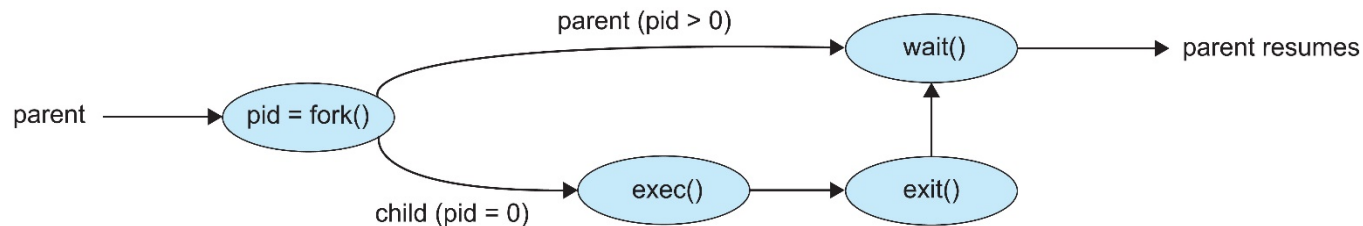


UNIX/Linux Process Creation

- fork system call
 - Create a new (child) process
 - The new process duplicates the address space of its parent
 - Child & Parent execute concurrently after fork
 - Child: return value of fork is 0
 - Parent: return value of fork is PID of the child process
- (child) execvp system call
 - Load a new binary file into memory – destroying the old code
- (parent) wait system call
 - The parent waits for one of its child processes to complete

Process Creation

- An UNIX example
 - `fork()` system call creates new process
 - `exec`**l**`p()` system call used after a `fork()` to replace the process' memory space with a new program
 - Parent process calls `wait()` for the child to terminate



```

int main()
{
    pid_t child_pid;
    child_pid = fork();

    if (child_pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed\n");
        return 1;
    }
    else if (child_pid == 0) { /* child process */
        printf("I am the child %d\n", child_pid);
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        printf("I am the parent of %d\n", child_pid);
        wait(NULL);

        printf("Child Complete\n");
    }

    return 0;
}

```

```

int main()
{
    pid_t child_pid;
    child_pid = fork();

    if (child_pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed\n");
        return 1;
    }
    else if (child_pid == 0) { /* child process */
        printf("I am the child %d\n", child_pid);
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        printf("I am the parent of %d\n", child_pid);
        wait(NULL);

        printf("Child Complete\n");
    }

    return 0;
}

```

Parent process
child_pid == 1234

```

int main()
{
    pid_t child_pid;
    child_pid = fork();

    if (child_pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed\n");
        return 1;
    }
    else if (child_pid == 0) { /* child process */
        printf("I am the child %d\n", child_pid);
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        printf("I am the parent of %d\n", child_pid);
        wait(NULL);

        printf("Child Complete\n");
    }

    return 0;
}

```

Child process
child_pid == 0


```
int main()
```

```
{
```

```
    pid_t child_pid;  
    child_pid = fork();
```

```
    Parent process  
    child_pid == 1234
```

```
    if (child_pid < 0) { /* error occurred */  
        fprintf(stderr, "Fork Failed\n");  
        return 1;  
    }
```

```
    else if (child_pid == 0) { /* child process */  
        printf("I am the child %d\n", child_pid);  
        execlp("/bin/ls", "ls", NULL);  
    }
```

```
    else { /* parent process */
```

```
        /* parent will wait for the child to complete */  
        printf("I am the parent of %d\n", child_pid);  
        wait(NULL);
```

```
        printf("Child Complete\n");
```

```
    }
```

```
    return 0;
```

```
}
```

```
int main()
```

```
{
```

```
    pid_t child_pid;  
    child_pid = fork();
```

Child process
child_pid == 0

```
    if (child_pid < 0) { /* error occurred */  
        fprintf(stderr, "Fork Failed\n");  
        return 1;  
    }
```

```
    else if (child_pid == 0) { /* child process */  
        printf("I am the child %d\n", child_pid);  
        execlp("/bin/ls", "ls", NULL);
```

Load a new binary file
into memory –
destroying the old code

問題: 會不會輸出test?

```
    }  
    else { /* parent process */  
        /* parent will wait for the child to complete */  
        printf("I am the parent of %d\n", child_pid);  
        wait(NULL);  
  
        printf("Child Complete\n");  
    }
```

```
    return 0;
```

```
}
```

```
int main()
{
```

```
    pid_t child_pid;
    child_pid = fork();
```

```
    Parent process
    child_pid == 1234
```

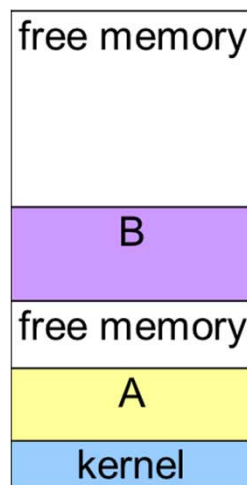
```
        if (child_pid < 0) { /* error occurred */
            fprintf(stderr, "Fork Failed\n");
            return 1;
        }
        else if (child_pid == 0) { /* child process */
            printf("I am the child %d\n", child_pid);
            execlp("/bin/ls", "ls", NULL);
        }
        else { /* parent process */
            /* parent will wait for the child to complete */
            printf("I am the parent of %d\n", child_pid);
            wait(NULL);
            printf("Child Complete\n");
        }
        return 0;
    }
```

Parent and Child Process

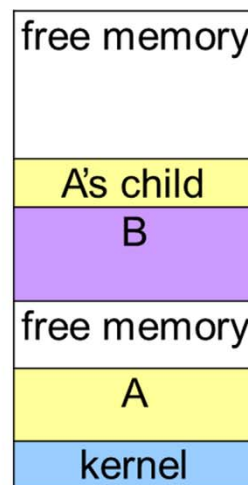
- Resource sharing: 3 modes
 - Parent and child processes **share all** resources
 - Child process shares **subset** of parent's resources
 - Parent and child **share no** resources
- Execution: 2 modes
 - 非同步 Parent and children execute concurrently
 - 同步 Parent waits until children terminate
- Address space: 2 modes
 - Child duplicate the parent
 - Child has a program loaded into it (execvp)

UNIX/Linux Process Creation

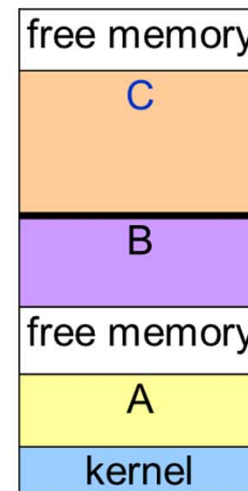
- Memory space of fork():
 - Old implementation: A's child is an exact copy of parent
 - Current implementation: use **copy-on-write** technique to store differences in A's child address space



Originally



After A does
an **fork** fork時只會為
child創建新的
stack seg



After the child
does an **execp**

Process Termination

- Terminate when the last statement is executed or `exit()`
 - All resources of the process, including physical & virtual memory, open files, I/O buffers, are deallocated
- Parent may terminate execution of children processes by specifying its PID (`abort`)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
- Cascading termination (some systems enforce this)
 - killing (exiting) parent → killing (exiting) all its children

Process Termination

- Cascading termination
 - If a process terminates, then all its children must also be terminated
- Wait for child process to terminate

```
pid = wait(&status);
```

- Zombie (child早死)
 - Child process ends before the parent calls wait()
 - Child process 資源被回收，但PCB仍在
- Orphan (parent早死)
 - Parent terminated without calling wait(), leaving child alone
 - 換parent: Linux **systemd** can be the parent of all process
 - call wait() periodically to collect the orphan child
 - Also allows other process to be the parent of the orphan child

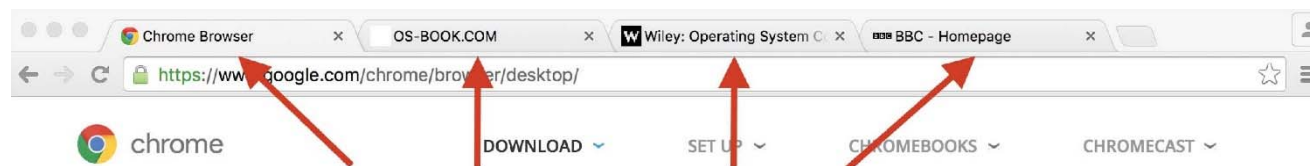
Ps. 在parent calls wait或parent結束後，pcb就會被回收

Android Process Hierarchy

- In android, the system maintains an importance hierarchy of processes
 - Terminate based on the importance of processes
- Importance hierarchy
 - Foreground process
 - Visible process
 - Service process (apparent to the user; ex: music)
 - Background process
 - Empty process

Chrome Browser

- Browsers with single process
 - If one web site causes trouble, entire browser can hang or crash
- Chrome is multi-process; 3 different types of processes:
 - **Browser process** manages user interface, disk and network I/O
 - **Renderer process** renders web pages, deals with HTML, JavaScript. A new renderer created for each website opened
 - Runs in sandbox restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in process** for each type of plug-in



Each tab represents a separate process.

Inter-process Communication

- IPC
 - Exchanging data among multiple processes
 - Shared memory; message passing
- Cooperating process
 - Share data with other processes
- Purposes
 - information sharing
 - computation speedup (viable only in multi-core CPUs)
 - modularity

Consumer & Producer Problem

- Producer process produces information that is consumed by a consumer process

- A buffer as the shared memory

- The buffer is a circular array with size **BUFFER_SIZE**

- first free: in

- first available: out

- empty: $in == out$

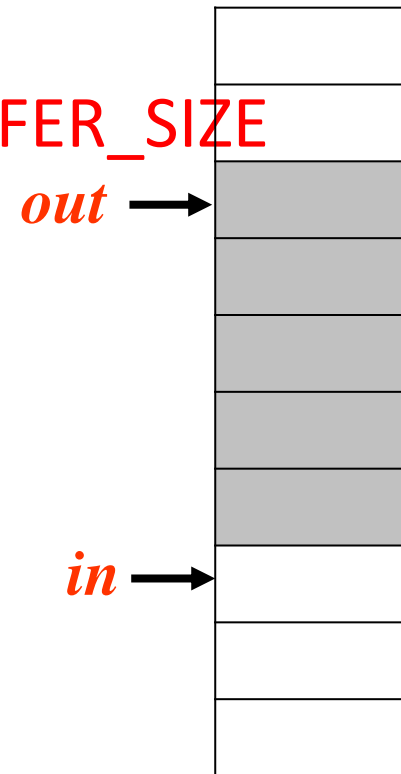
- full:

$(in+1) \% \text{BUFFER_SIZE} == out$

```
#define BUFFER_SIZE 10

typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

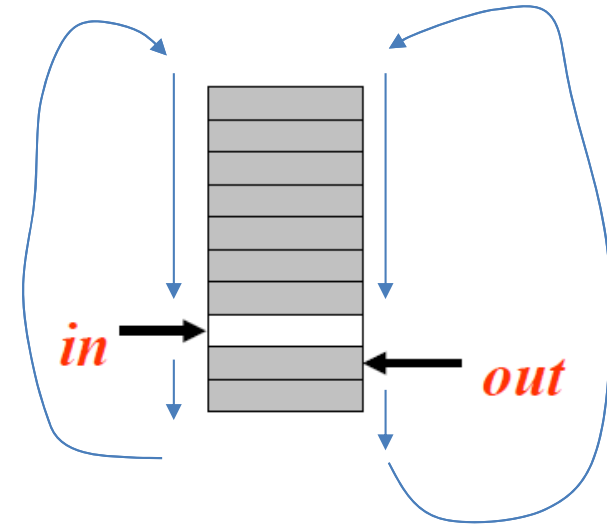


Shared-Memory Solution

```
/*producer*/  
while (1) {  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; //wait if buffer is full  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

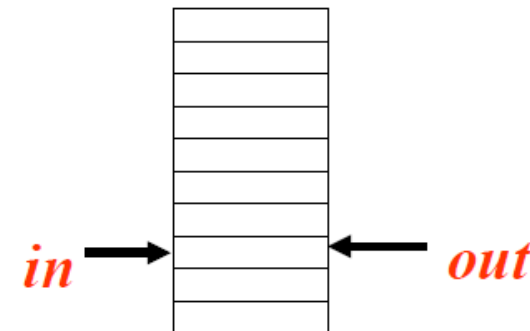
the buffer is full

"in" only modified by producer



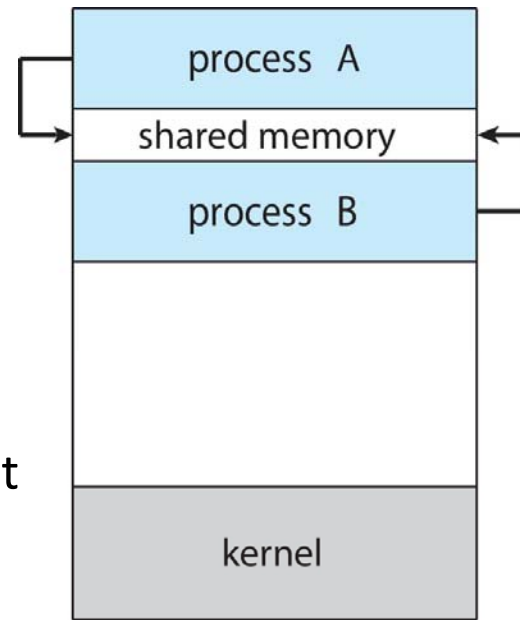
```
/*consumer*/  
while (1) {  
    while (in == out); //wait if buffer is empty  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
}
```

"out" only modified by consumer



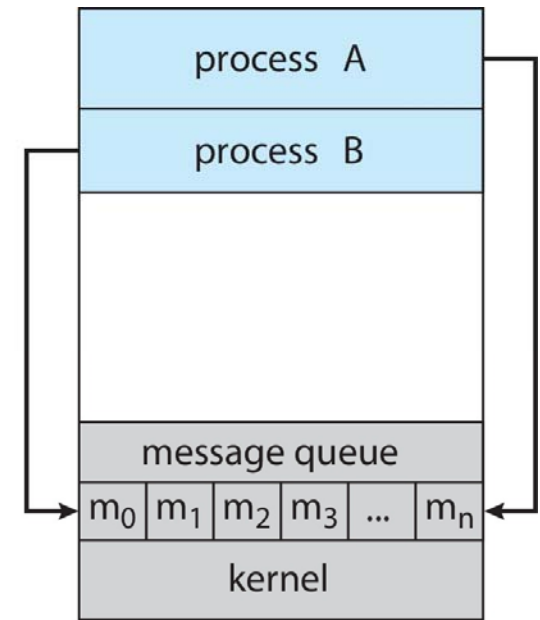
Communication Methods

- Shared memory:
 - Require more careful user synchronization
 - Implemented by memory access: faster speed
 - 不需要OS支援
- Message passing:
 - No sync issues: more efficient for small data
 - Use send/recv (sys call) to transmit messages
 - slower than shared memory
 - 需要OS支援



(a)

Shared memory.



(b)

Message passing

Shared Memory

- Processes are responsible for
 - Establishing a region of shared memory
 - The hosting process creates a shared-memory region resides in the address space of the process
 - The communicating processes must agree to remove memory access constraint from OS
 - Determining the form of the data and the location
 - Ensuring data are not written simultaneously by processes
 - i.e. critical section problem
 - OS不介入，通訊二方要自己實現

Examples of IPC Systems - POSIX

See <https://chmodcommand.com/>

- POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

int

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Use `mmap()` to memory-map a file pointer to the shared memory object
- Reading and writing to shared memory is done by using the pointer returned by `mmap()`.

	Owner Rights (u)	Group Rights (g)	Others Rights (o)
Read (4)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
Write (2)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
Execute (1)	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>

<https://github.com/greggagne/osc10e/blob/master/ch3/shm-posix-producer.c>

IPC POSIX Producer

```
int main()
{
/* the size (in bytes) of shared memory object */
const int SIZE = 4096;
/* name of the shared memory object */
const char *name = "OS";
/* strings written to shared memory */
const char *message_0 = "Hello";
const char *message_1 = "World!";

/* shared memory file descriptor */
int shm_fd;
/* pointer to shared memory object */
void *ptr;
```

```
/* create the shared memory object */
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

/* configure the size of the shared memory object */
ftruncate(shm_fd, SIZE);

/* memory map the shared memory object */
ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);
```

```
/* write to the shared memory object */
sprintf(ptr, "%s", message_0);
ptr += strlen(message_0);
sprintf(ptr, "%s", message_1);
ptr += strlen(message_1);
```

使用mmap回傳的指標來寫入共享記憶體

```
return 0;
```

```
}
```

<https://github.com/greggagne/osc10e/blob/master/ch3/shm-posix-producer.c>

IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

此敘述執行後，shared memory立即失效

<https://github.com/greggagne/osc10e/blob/master/ch3/shm-posix-consumer.c>

Message-Passing System

- Enable processes to communicate without sharing the same address space
 - Communicating via a queue (pipe) provided by OS
- IPC facility provides two operations:
 - Send(message) – message size fixed or variable (P.128中)
 - Receive(message)
- To communicate, processes need to
 - Establish a communication link (pipe)
 - Exchange a message via send/receive

Message-Passing System

- Three issues
 - Direct or indirect communication (naming)
 - Synchronous or asynchronous communication
 - Automatic or explicit buffering

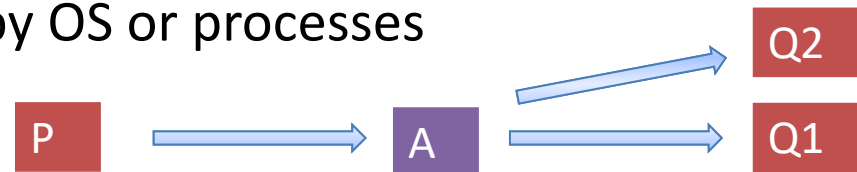
Direct communication



- Processes must name each other explicitly:
 - Send (P, message) – send a message to proc P
 - Receive (Q, message) – receive a message from process Q
- Properties of communication link
 - Links are established automatically among processes
 - One-to-One relationship between links and processes
- Drawback
 - The communicating parties are hard coded

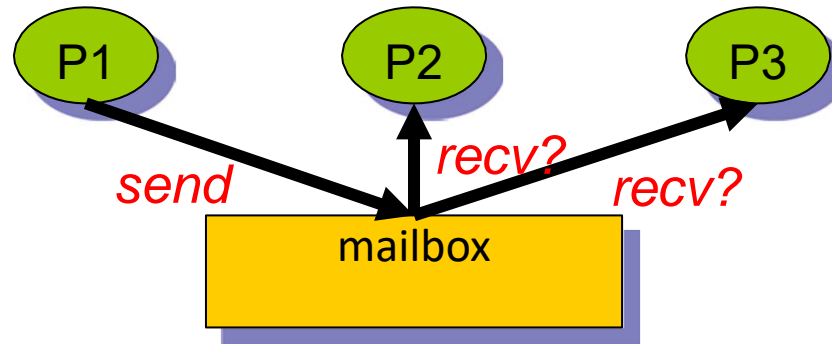
Indirect communication

- Messages are exchanged via mailboxes (ports)
 - Each mailbox has a unique ID
 - Processes can communicate if they share a mailbox
 - Send (A, message) – send a message to mailbox A
 - Receive (A, message) – receive a message from mailbox A
- Properties of communication link
 - Link established only if processes share a common mailbox
 - Many-to-Many relationship between links and processes
 - Link may be unidirectional or bi-directional
 - Mailbox can be owned either by OS or processes



Indirect communication

- Which will receive the message sent by P1?



- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select the receiver

Synchronization

- Different send/recv sync combinations are possible
- Synchronous messaging passing (blocking)
 - Blocking send: sender is blocked until the message is received by receiver or by the mailbox
 - Blocking receive: receiver is blocked until the message is available
- Asynchronous messaging passing (non-blocking)
 - Non-blocking send: sender sends the message and resumes operation
 - Non-blocking receive: receiver receives a valid message or a null

Queue Buffering

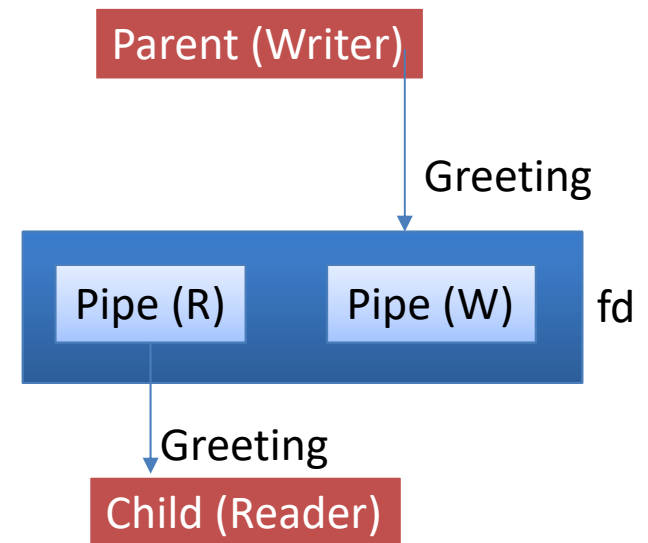
- Buffer
 - Queue of messages attached to the link
- Implementation alternatives
 - Zero capacity: blocking send/receive
 - Bounded capacity: if full, sender will be blocked
 - Unbounded capacity: sender never blocks

Pipes

- A conduit (導管) allowing two processes to communicate
- Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?
 - Can the pipes be used over a network?
- **Ordinary pipes (parent-child)**
 - A parent process creates a pipe
 - Uses the pipe to communicate with a child process
- **Named pipes**
 - can be accessed without a parent-child relationship

Ordinary Pipes

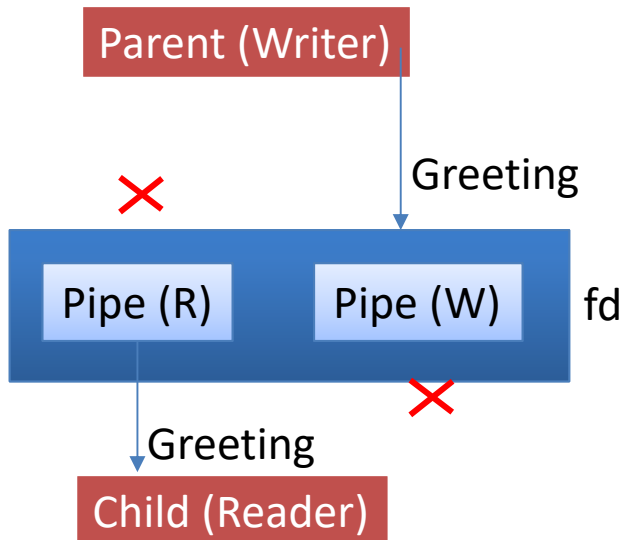
- Feature
 - Ordinary pipes are unidirectional
 - Require parent-child relationship between communicating processes
- Producer writes to one end
 - the **write-end** of the pipe
- Consumer reads from the other end
 - the **read-end** of the pipe



```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
```

```
#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1
```

```
int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;
```



```
/* create the pipe */
if (pipe(fd) == -1) {
    fprintf(stderr, "Pipe failed");
    return 1;
}
```

```
/* fork a child process */
pid = fork();
```

Parent

```
if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    return 1;
}
```

```
if (pid > 0) { /* parent process */
    /* close the unused end of the pipe */
    close(fd[READ_END]);
```

```
/* write to the pipe */
write(fd[WRITE_END], write_msg, strlen(write_msg)+1);
```

```
/* close the write end of the pipe */
close(fd[WRITE_END]);
```

```
}
else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);
```

```
/* read from the pipe */
read(fd[READ_END], read_msg, BUFFER_SIZE);
printf("read %s", read_msg);
```

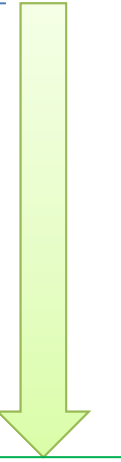
```
/* close the read end of the pipe */
close(fd[READ_END]);
```

```
}
```

```
return 0;
```

```
}
```

Child



Named Pipes

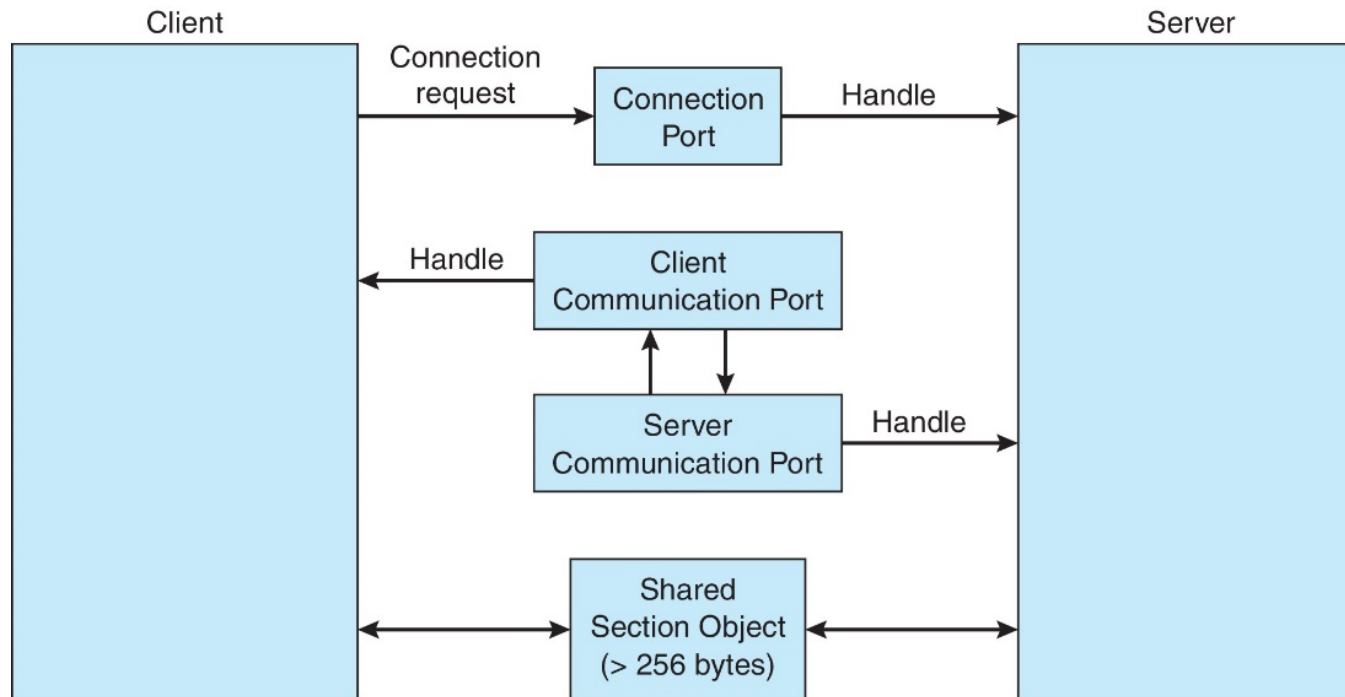
- More powerful than ordinary pipes
 - Communication is bidirectional
 - No parent-child relationship is necessary
 - Several processes can use the named pipe
- Provided on both UNIX and Windows systems
 - UNIX only supports half-duplex
 - Windows supports full-duplex

<https://github.com/3p3r/node-named-pipe>
mkfifo

Windows ALPC

- ALPC
 - Advanced Local Procedure Call
- Steps
 - Server publishes a connection port publicly
 - Client sends a request to the connection port
 - Server setup a private channel (contains 2 communication ports)
 - Client obtain the handle of the communication port from server
 - Messaging
 - Small messages: (≤ 256 bytes) directly via communication ports
 - Big messages: section object (shared memory)
 - Huge messages: direct memory access via specific API

Windows ALPC



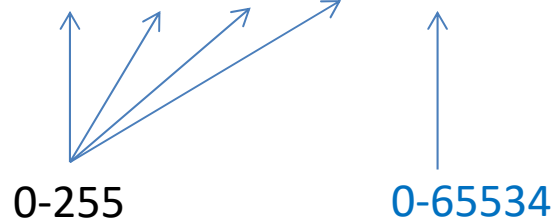
Socket

- 網路通訊程式
 - 位於二台不同電腦的程式互相傳送訊息
- Internet上每台電腦至少有一個固定位址

- IP + Port

- 例如

- 140.119.168.10: 80



每個IP都有65535個port(埠)

http://www.nccu.edu.tw

網路名稱轉譯

www.nccu.edu.tw → 140.119.168.10

http:// → 80

IP網路通訊程式基礎

- 不同Port代表不同服務
- Well-known port
 - 80: HTTP
 - 443: HTTPS
 - 21: FTP
 - 22: SSH
 - 23: Telnet



IP網路通訊程式基礎

- 特殊位址

- 127.0.0.1 → localhost

- 指向自己

- 同一台電腦不同port間也可以通訊

```
127.0.0.1: 12345 →  
127.0.0.1: 8080
```

- 192.168.

- 虛擬位址 (只在LAN中有效)

- 239.

- 群播 (multicast)

Socket Server

```
let net = require('net');  
  
let server = net.createServer(function(s) {  
  s.setEncoding('utf8');  
  s.on('data', function(data) {  
    console.log(s.remoteAddress);  
    console.log(s.remotePort);  
    console.log(data);  
  });  
});  
  
server.listen(1337);
```

→ 設定編碼為字串

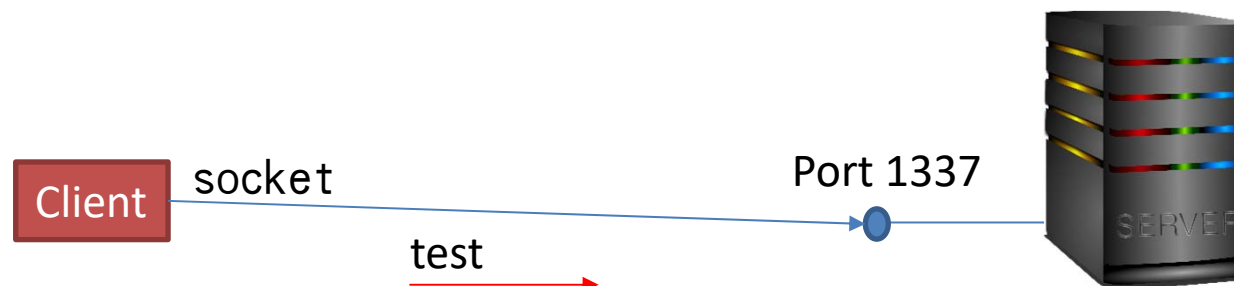
Port 1337



若沒有設定編碼為字串，則傳入的data會是Buffer物件
若是如此，則要使用data.toString()，資料才能正常顯示。

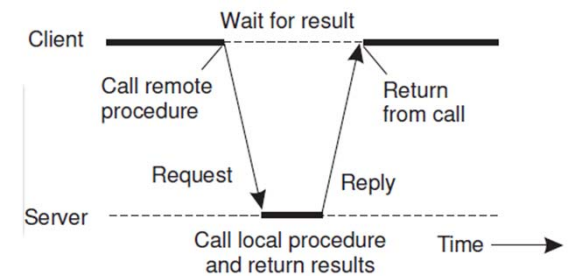
Socket Client

```
let net = require('net');  
  
let client = new net.Socket();  
  
client.connect(1337, '127.0.0.1', function () {  
  console.log('connected');  
  client.write('test', () => console.log('written'));  
});
```

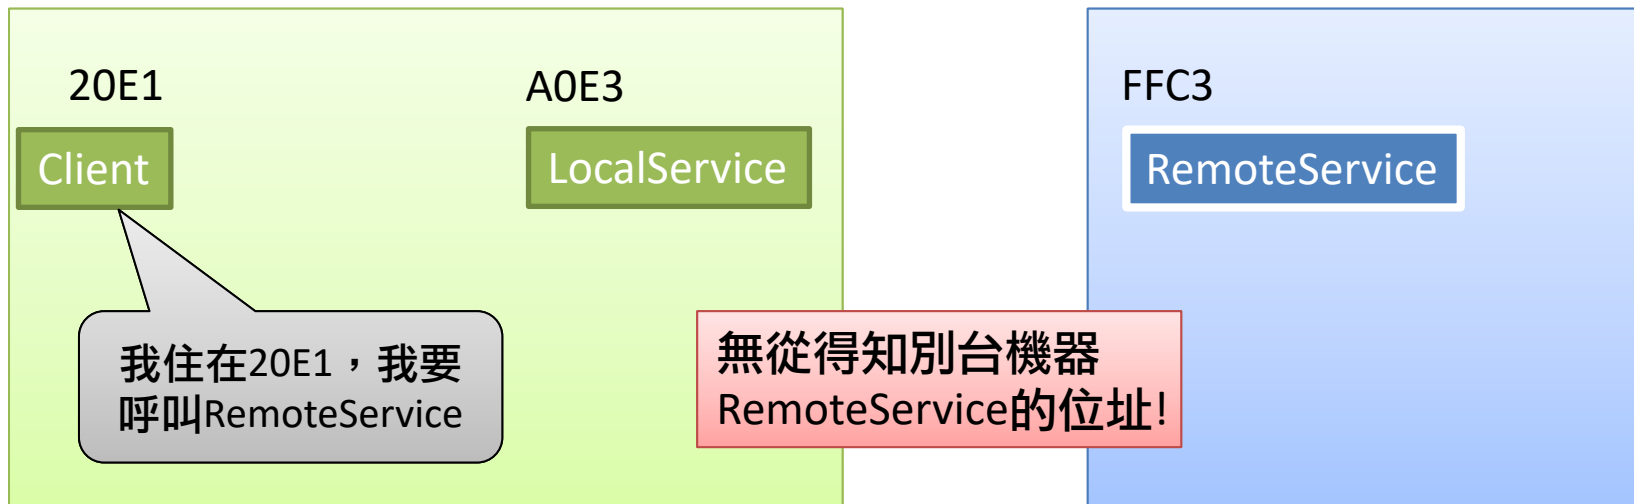
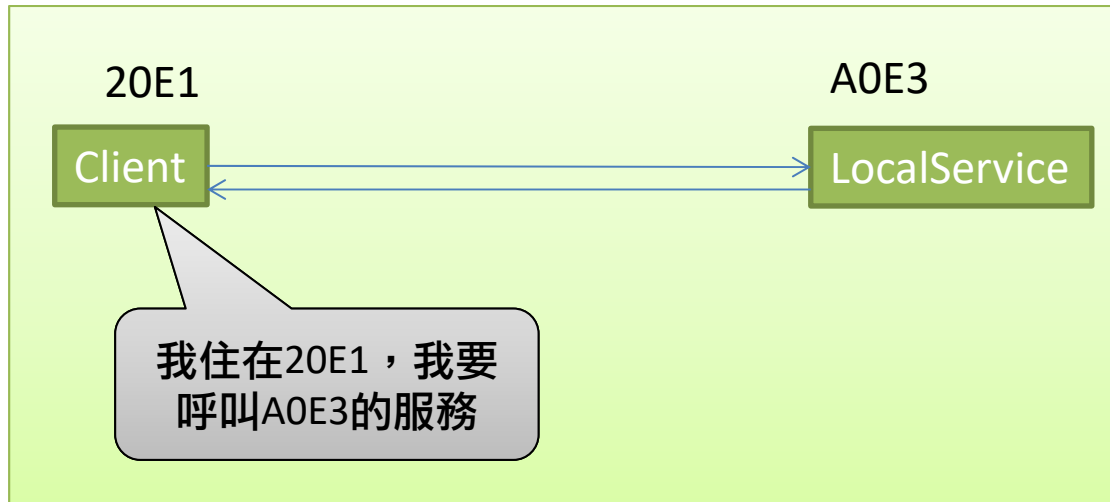


Remote Procedure Call (RPC)

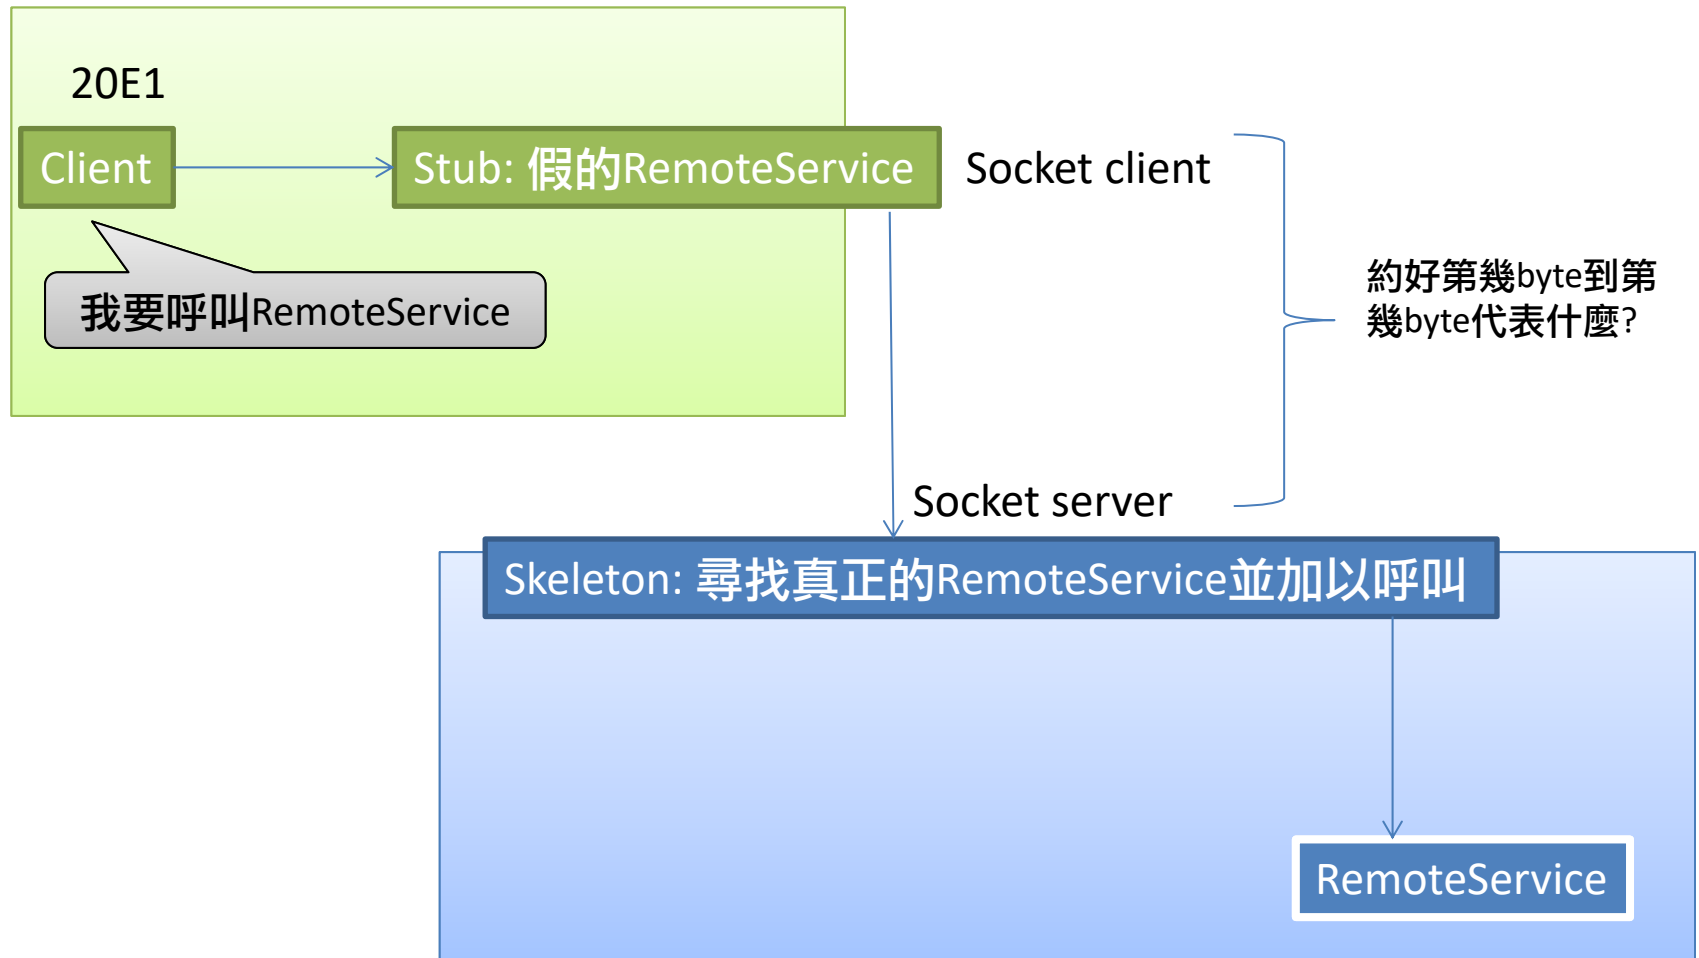
- Motivation
 - A type of the direct communication mechanism
 - Calling procedure on a remote host “as if” calling a local procedure (Saif and Greaves, 2001)
 - Warning: unaware of “remoting” is considered harmful
- Approach
 - Communication between caller & callee can be hidden by using a proxied procedure-call mechanism



Remote Procedure Calls



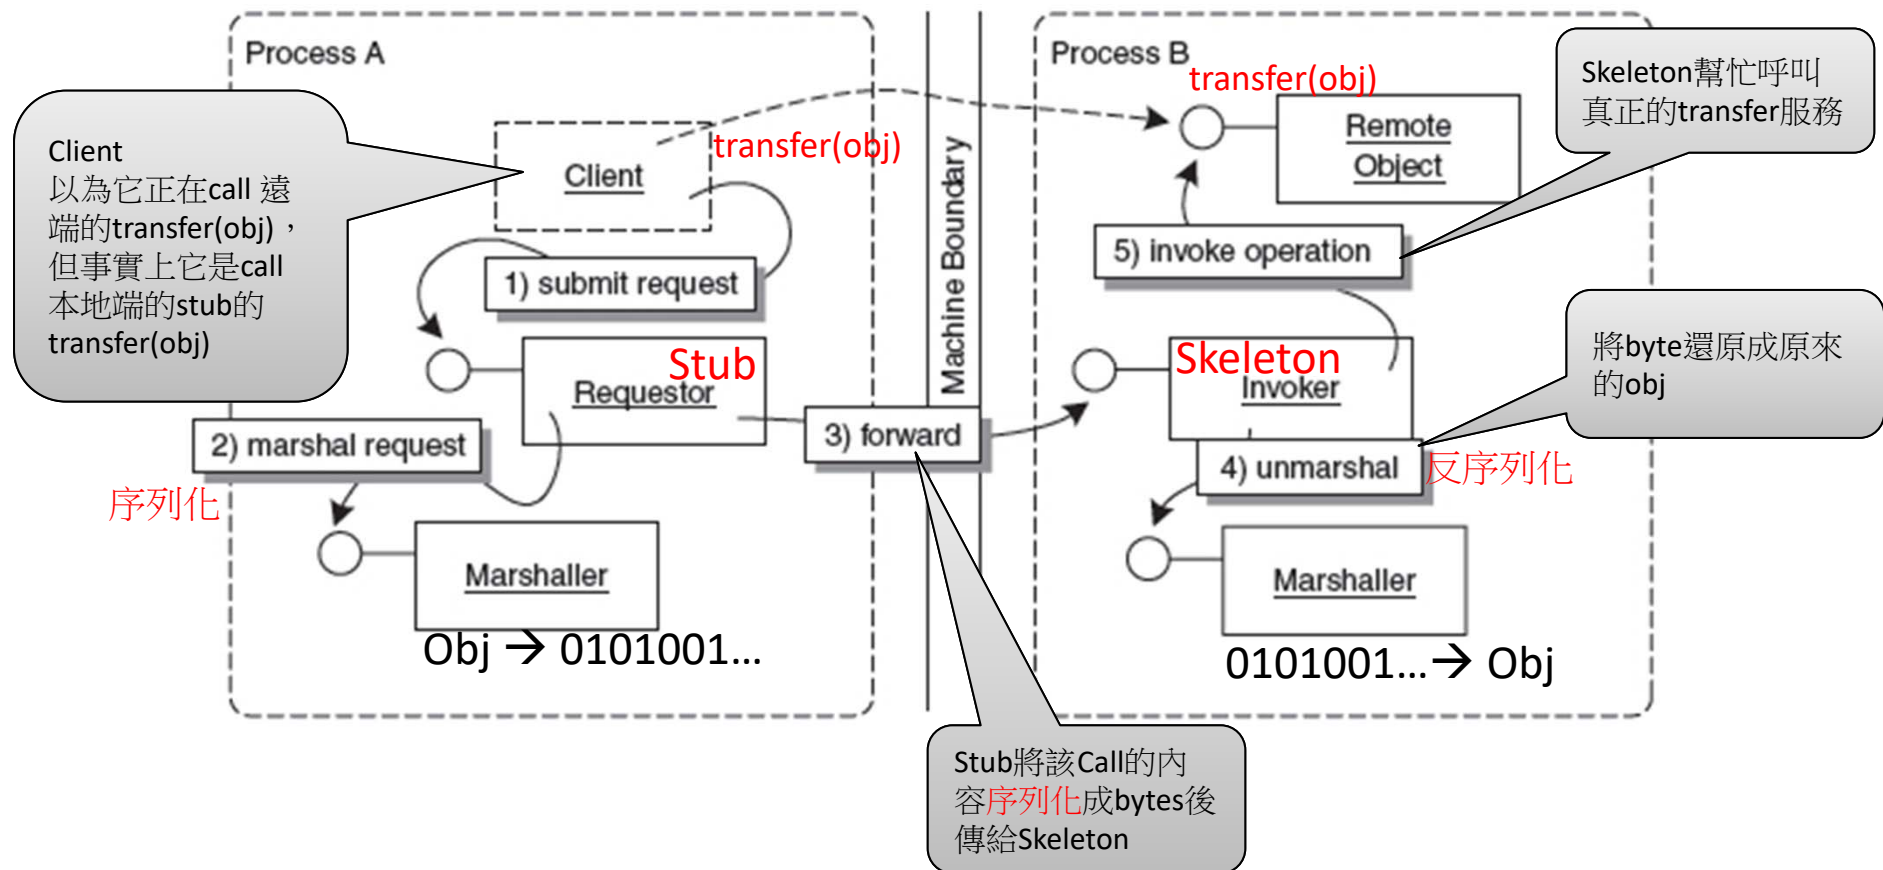
Remote Procedure Calls



Serialization (Marshalling) 序列化

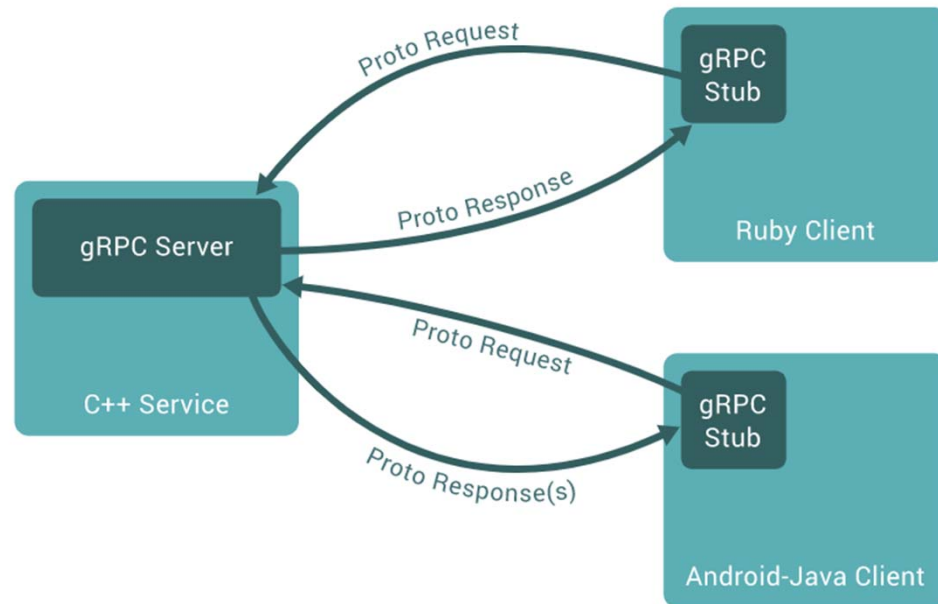
- 參數如何傳送?
 - Client and server machines may have different data representations
 - E.g., byte ordering
 - Serialization序列化: transforming a value into a sequence of bytes
 - Client and server have to agree on the same encoding standard
 - 議題
 - How are basic data values represented (integers, floats, characters)
 - How are complex data values represented (arrays, unions)

呼叫遠端函式(物件)的架構



Case: gRPC

- A modern open source high performance RPC framework that can run in any environment
 - Efficiently connect services in and across data centers
 - Technology stack
 - HTTP 2
 - Protocol Buffers



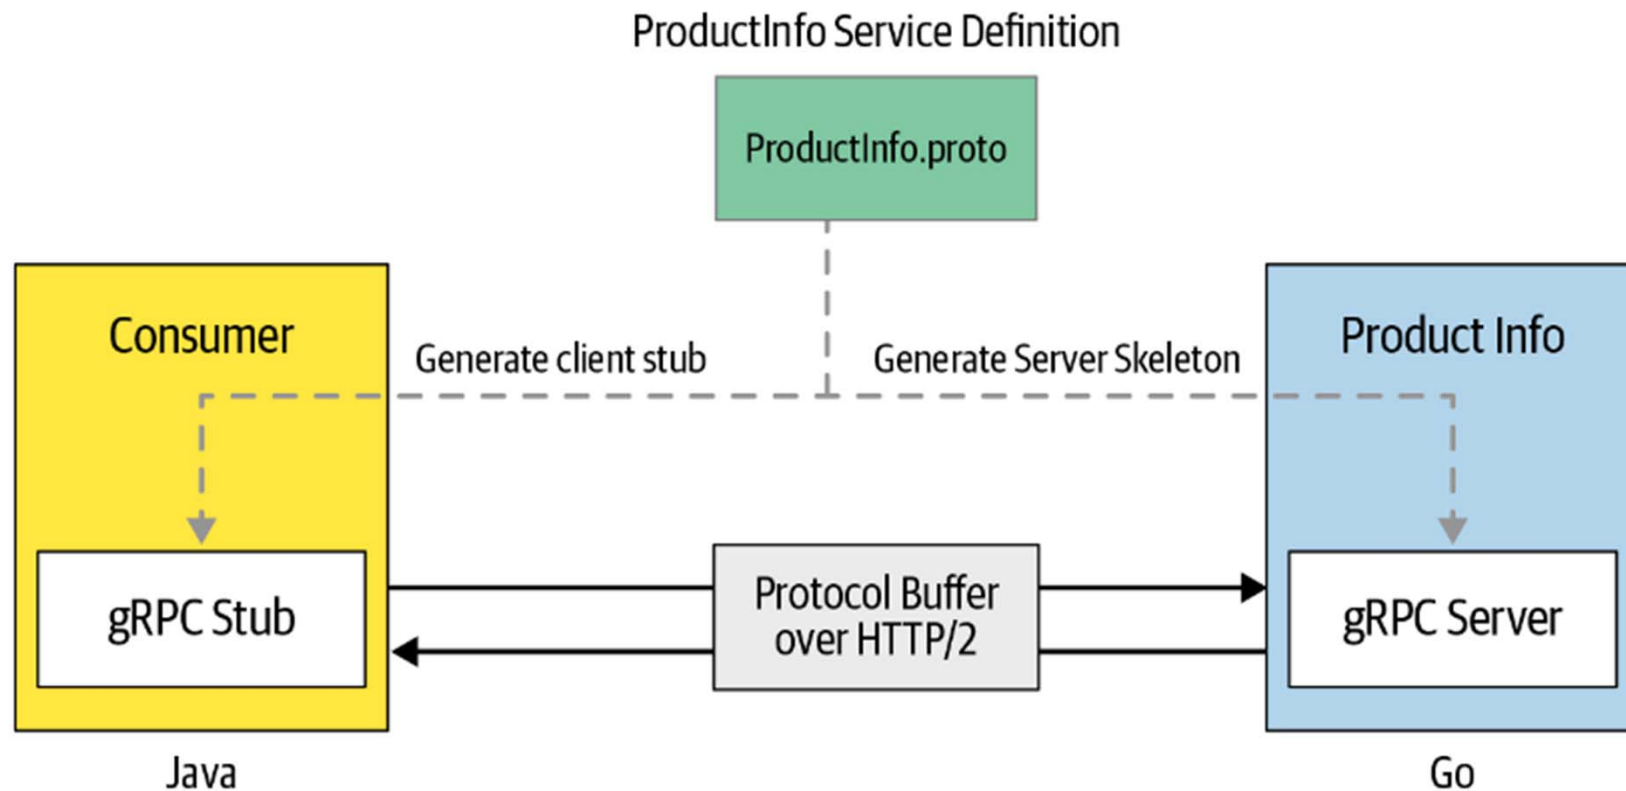
gRPC in the Realworld

- gRPC has been widely adopted for building microservices and cloud native applications
- Netflix
 - Initially using in-house RESTful solution on HTTP/1.1
 - With the adoption of gRPC, Netflix has seen a massive boost in developer productivity
 - Creating a client, which could take up to two to three weeks, takes a matter of minutes with gRPC.

gRPC in the Realworld

- Dropbox
 - Dropbox runs hundreds of polyglot microservices, which exchange millions of requests per second
 - Initial solution
 - A homegrown RPC framework with a custom protocol for manual serialization
 - Apache Thrift
 - Legacy HTTP/1.1-based RPC framework + protobuf
 - New solution: Courier
 - A gRPC-based RPC framework
 - A customized solution to meet specific requirements like authentication, authorization, service discovery, service statistics, event logging, and tracing tools

A Server and a Client based on gRPC



Service Definition

```
// The greeting service definition.
service Greeter {
  // Sends a greeting
  rpc SayHello (HelloRequest) returns (HelloReply) {}
}

// The request message containing the user's name.
message HelloRequest {
  string name = 1;
}

// The response message containing the greetings
message HelloReply {
  string message = 1;
}
```

```
function sayHello(call, callback) {  
  callback(null, {message: 'Hello ' + call.request.name});  
}  
  
function main() {  
  var server = new grpc.Server();  
  server.addService(hello_proto.Greeter.service,  
    {sayHello: sayHello});  
  server.bind('0.0.0.0:50051', grpc.ServerCredentials.createInsecure());  
  server.start();  
}
```

```
function main() {  
  
    var client = new hello_proto.Greeter('localhost:50051',  
                                         grpc.credentials.createInsecure());  
    client.sayHello({name: 'you'}, function(err, response) {  
        console.log('Greeting:', response.message);  
    });  
  
}
```

課後閱讀 (會考)

- P115 手持裝置的Multitasking
- P135 Mach Messaging Passing: Task Self port的用途
- P136 什麼是Bootstrap server
- P136 Mach Message Passing如何傳送大檔?
- P138 ALPC是什麼? 作用為何?
- P139 ALPC為什麼需要section object?
- P151 Android中提供什麼支援來開發RPC?

Q & A