

SQL-The Relational Database Standard-III

政治大學

資訊科學系

沈錕坤

Summary of SQL Queries

SELECT <attribute list>

FROM <table list>

[WHERE <condition>]

[GROUP BY <grouping attribute(s)>]

[HAVING <group condition>]

[ORDER BY <attribute list>]

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

Dnumber	Dlocation
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

WORKS_ON

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

Essn	Dependent_name	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

Update in SQL

Updates in SQL

- ◆ three SQL commands to modify the database
 - INSERT
 - DELETE
 - UPDATE

Updates in SQL (cont.)

- ◆ **E.g.** Insert a tuple for a new EMPLOYEE for whom we only know the FNAME, LNAME, and SSN attributes.

```
INSERT INTO EMPLOYEE (FNAME, LNAME, SSN)  
VALUES ('Richard', 'Marini', '653298653')
```

- ◆ The constraints specified in the DDL commands are automatically enforced by the DBMS when updates are applied to the database

U3A

- ◆ **Insertion of multiple tuples** resulting from a query into a relation
<E.g.> to create a temporary table that has the name, number of employees, and total salaries for each department.

```
U3A:  CREATE TABLE  DEPTS_INFO  
      (DEPT_NAME  VARCHAR(10),  
      NO_OF_EMPS  INTEGER,  
      TOTAL_SAL   INTEGER);
```

INSERT INTO

```
DEPTS_INFO (DEPT_NAME, NO_OF_EMPS, TOTAL_SAL)  
SELECT  DNAME, COUNT (*), SUM (SALARY)  
FROM    DEPARTMENT, EMPLOYEE  
WHERE   DNUMBER = DNO  
GROUP BY DNAME ;
```

DELETE

- ◆ Removes tuples from a relation
- ◆ Includes a WHERE-clause to select the tuples to be deleted
- ◆ The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause
- ◆ Tuples are deleted from only one table at a time (unless CASCADE is specified on a referential integrity constraint)
- ◆ A missing WHERE-clause specifies that all tuples in the relation are to be deleted; the table then becomes an empty table
- ◆ Referential integrity should be enforced

U4A/U4B/U4C/U4D

**U4A: DELETE FROM EMPLOYEE
WHERE LNAME = 'Brown'**

**U4B: DELETE FROM EMPLOYEE
WHERE SSN = '123456789'**

**U4C: DELETE FROM EMPLOYEE
WHERE DNO IN
(SELECT DNUMBER
FROM DEPARTMENT
WHERE DNAME = 'Research')**

U4D: DELETE FROM EMPLOYEE *delete all records*

UPDATE

- ◆ Used to modify attribute values of one or more selected tuples
- ◆ A WHERE-clause selects the tuples to be modified
- ◆ An additional SET-clause specifies the attributes to be modified and their new values
- ◆ Each command modifies tuples in the **same relation**
- ◆ **Referential integrity** should be enforced

U5/U6

U5: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively

```
UPDATE PROJECT
SET      PLOCATION = 'Bellaire', DNUM = 5
WHERE    PNUMBER = 10
```

U6: Give all employees in the 'Research' department a 10% raise in salary.

```
UPDATE EMPLOYEE
SET      SALARY = SALARY * 1.1
WHERE DNO IN
        ( SELECT DNUMBER
          FROM    DEPARTMENT
          WHERE   DNAME = 'Research' )
```

View in SQL

Relational Views in SQL

- A view is a single **virtual table** that is derived from other tables –
 - (1) **base tables** or
 - (2) previously defined **views**
- ◆ A view does **not necessarily exist** in **physical form**, which limits the possible update operations that can be applied to views
- ◆ There are no limitations on querying a view
- ◆ **CREATE VIEW** command specify a view by specifying a (virtual) table name and a defining query
- ◆ The view attribute names can be inherited from the attribute names of the tables in the defining query

Relational Views in SQL (cont.)

- ◆ One advantages of a view is to **simplify** the **specification** of queries.
- ◆ Views can also be used as a **security** and **authorization** mechanism
- ◆ DBMS responsible for keeping the view up-to-date if the base tables on which the view is defined are modified
- ◆ It is the responsibility of **DBMS**, not the user, to make sure that the view is **up to date**
- ◆ View is **not realized** at the time of **view definition**, but at the time we specify a query on the view

V1

♦ **CREATE VIEW WORKS_ON1 AS**
SELECT FNAME, LNAME, PNAME, HOURS
FROM EMPLOYEE, PROJECT, WORKS_ON
WHERE SSN=ESSN **AND** PNO=PNUMBER ;

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

PROJECT

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

WORKS_ON

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

QUERIES ON VIEWS

- ◆ Retrieve the last name and first name of all employees who work on 'ProjectX'.

```
QV1: SELECT  PNAME, FNAME, LNAME  
        FROM    WORKS_ON1  
        WHERE   PNAME='ProjectX' ;
```

```
* CREATE VIEW    WORKS_ON1 AS  
SELECT          FNAME, LNAME, PNAME, HOURS  
FROM            EMPLOYEE, PROJECT, WORKS_ON  
WHERE           SSN=ESSN AND PNO=PNUMBER ;
```


V2

CREATE VIEW DEPT_INFO (DEPT_NAME, NO_OF_EMPS, TOTAL_SAL)
AS

specify column names

SELECT DNAME, COUNT(*), SUM(SALARY)
FROM DEPARTMENT, EMPLOYEE
WHERE DNUMBER=DNO
GROUP BY DNAME ;

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

Drop View

A view is removed using the **DROP VIEW** command

V1A: DROP VIEW WORKS_ON1 ;

V2A: DROP VIEW DEPT_INFO ;

View 是 Virtual Table , DBMS 如何實作 View 的功能 ?



View Implementation

- ◆ Two main approaches
 - Query modification
 - Modifying view query into a query on the underlying base tables
 - Disadvantage: inefficient for views of complex queries
 - View materialization
 - Physically creating a temporary view table
 - Incremental update view when base tables are updated
 - If view is not queried for a period of time, system automatically remove the physical view table

View Implementation (cont.)

- ◆ Query modification: modifying view query into a query on the underlying base tables

create
view

```
CREATE VIEW WORKS_ON1 AS
SELECT FNAME, LNAME, PNAME, HOURS
FROM EMPLOYEE, PROJECT, WORKS_ON
WHERE SSN = ESSN AND PNO = PNUMBER ;
```

query
view

```
SELECT PNAME, FNAME, LNAME
FROM WORKS_ON1
WHERE PNAME = 'ProjectX' ;
```



```
SELECT PNAME, FNAME, LNAME
FROM EMPLOYEE, PROJECT, WORKS_ON
WHERE SSN = ESSN AND PNO = PNUMBER
AND PNAME='ProjectX' ;
```

Updating of Views

- ◆ A view update operation may be mapped in **multiple ways** to update operations on the defining base relations - **ambiguity**
- ◆ updating views is still an active research area

Ex: To update the WORKS_ON1 view by modifying the PNAME attribute of 'John Smith' from 'ProductX' to 'ProductY'.

```
UV1:  UPDATE  WORKS_ON1  
      SET      PNAME = 'ProductY'  
      WHERE  LNAME='Smith' AND  
            FNAME='John' AND  
            PNAME='ProductX'
```

View Update 如何對應到其 Base Table 的 Update ?

John Smith 參與的計畫不變，只是改名為 Product Y

OR

John Smith 參與的計畫改變，改參與Product Y 計畫？



Updating of Views (cont.)

- ◆ Updates on the base relations to give the desired update on the view
- ◆ Two possibilities

(1) Change the name of the 'ProductX' tuple in the PROJECT relation to 'ProductY'

```
UPDATE PROJECT
SET PNAME = 'ProductY'
WHERE PNAME = 'ProductX'
```

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

PROJECT

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

WORKS_ON

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

Updating of Views (cont.)

(2) Relate 'John Smith' to the 'ProductY' PROJECT tuple in place of the 'ProductX' PROJECT tuple

UPDATE WORKS_ON

**SET PNO = (SELECT PNUMBER
FROM PROJECT
WHERE PNAME = 'ProductY')**

*the project #
of ProductY*

**WHERE ESSN = (SELECT SSN
FROM EMPLOYEE
WHERE LNAME = 'Smith'
AND FNAME = 'John')**

John Smith

**AND PNO = (SELECT PNUMBER
FROM PROJECT
WHERE PNAME = 'ProductX')**

*works on
ProductX*

Updating of Views (cont.)

- ◆ Some view updates may not make much sense

UV2: UPDATE DEPT_INFO

SET TOTAL_SAL = 100000 如何分配？

WHERE DNAME = 'Research' ;

- ◆ In general, we cannot guarantee that any view can be updated
- ◆ A view update is unambiguous only if **one update** on the base relations can accomplish the desired update effect on the view
- ◆ If a view update can be mapped to **more than one update** on the underlying base relations, we must have a certain procedure to choose the desired update

Updating of Views (cont.)

- ◆ A view with a single defining table is **updatable** if the view attributes contain the **primary key**
- ◆ Views defined on multiple tables using **joins** are generally **not updatable**
- ◆ Views defined using **grouping and aggregate** functions are **not updatable**

WITH

- ◆ 假設薪水最高的員工不只一位，列出所有這些員工的SSN。

- **CREATE VIEW** Max_Salary(Value)

- AS SELECT Max**(Salary)

- FROM** Employee;

- SELECT SSN**

- FROM** Employee, Max_Salary

- WHERE** Employee.Salary = Max_Salary.Value;

- **WITH** Max_Salary(Value)
 - AS SELECT Max**(Salary)
 - FROM** Employee

- SELECT SSN**

- FROM** Employee, Max_Salary

- WHERE** Employee.Salary = Max_Salary.Value;

Allows you to give the sub-query
block a name

Assertion

semantic integrity
constraint

Constraint: salary of an employee must not be greater than the salary of the manager of the department that the employee works for

(MySQL does not support assertion & domain)

```
CREATE ASSERTION SALARY_CONSTRAINT  
CHECK (NOT EXISTS (SELECT * FROM EMPLOYEE E,  
                    EMPLOYEE M, DEPARTMENT D  
                    WHERE E.SALARY > M.SALARY AND  
                    E.DNO=D.DNUMBER AND D.MGRSSN=M.SSN));
```

```
CREATE DOMAIN D_NUM AS INTEGER  
CHECK (D_NUM > 0 AND D_NUM < 21);
```

essentially a data type
w/ constraints that can be
used in multiple tables

Additional Features of SQL

Additional Features of SQL

- ◆ Granting, revoking of privileges to users to access certain relations e.g. `GRANT SELECT, INSERT ON 'ALICE'@'localhost'`
`REVOKE INSERT ON employees FROM 'Alice'@'localhost'`
- ◆ **Embed** SQL in programming language C, C++, COBOL, PASCAL based on cursor
- ◆ Transaction control granularity for concurrency control, recovery → *an indivisible unit of work*
- ◆ Physical database design parameters, file structures for relations, access paths as indexes
 - A* **atomic** : All changes to the data must be performed or not at all
 - C* **onsistent** : data must be consistent before & after the transaction
 - I* **solated** : no other process can change the data during the transaction
 - D* **urable** : the changes of the transaction must persist

Creating Indexes in SQL

- ◆ An SQL base relation generally corresponds to a stored file
- ◆ Statements can create and drop indexes on base relations
- ◆ These statements have been removed from SQL2 because they specify physical access paths - not logical concepts
- ◆ One or more indexing attributes are specified for each index
- ◆ **CREATE INDEX** statement is used
- ◆ Each index is given an *index name*

I1: *index name*
CREATE INDEX LNAME_INDEX
ON EMPLOYEE(LNAME);
 relation *attribute list*

Rank			DBMS	Database Model	Score		
Mar 2023	Feb 2023	Mar 2022			Mar 2023	Feb 2023	Mar 2022
1.	1.	1.	Oracle +	Relational, Multi-model ⓘ	1261.29	+13.77	+9.97
2.	2.	2.	MySQL +	Relational, Multi-model ⓘ	1182.79	-12.66	-15.45
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model ⓘ	922.01	-7.08	-11.77
4.	4.	4.	PostgreSQL +	Relational, Multi-model ⓘ	613.83	-2.67	-3.10
5.	5.	5.	MongoDB +	Document, Multi-model ⓘ	458.78	+6.02	-26.88
6.	6.	6.	Redis +	Key-value, Multi-model ⓘ	172.45	-1.39	-4.31
7.	7.	7.	IBM Db2	Relational, Multi-model ⓘ	142.92	-0.04	-19.22
8.	8.	8.	Elasticsearch	Search engine, Multi-model ⓘ	139.07	+0.47	-20.88
9.	9.	↑ 10.	SQLite +	Relational	133.82	+1.15	+1.64
10.	10.	↓ 9.	Microsoft Access	Relational	132.06	+1.03	-3.37
11.	↑ 12.	↑ 14.	Snowflake +	Relational	114.40	-1.26	+28.17
12.	↓ 11.	↓ 11.	Cassandra +	Wide column	113.79	-2.43	-8.35
13.	13.	↓ 12.	MariaDB +	Relational, Multi-model ⓘ	96.84	+0.03	-11.47
14.	14.	↓ 13.	Splunk	Search engine	87.97	+0.89	-7.39
15.	15.	↑ 16.	Amazon DynamoDB +	Multi-model ⓘ	80.77	+1.08	-1.03
16.	16.	↓ 15.	Microsoft Azure SQL Database	Relational, Multi-model ⓘ	77.44	-1.31	-7.23
17.	17.	17.	Hive	Relational	70.91	-1.21	-10.31
18.	18.	18.	Teradata	Relational, Multi-model ⓘ	63.74	+0.71	-5.11
19.	19.		Databricks	Multi-model ⓘ	60.86	+0.52	
20.	20.	↓ 19.	Neo4j +	Graph	53.51	-1.92	-6.16
21.	↑ 22.	↑ 24.	Google BigQuery +	Relational	53.44	+0.99	+6.78


(cf: <https://db-engines.com/en/ranking>)

Database Programming

- ◆ Approaches to DB programming
 - Embedding DB commands in a general-purpose programming language
 - Embedded SQL
 - Dynamic SQL : SQL statements are constructed at runtime
 - SQLJ : a standard that allows embedding SQL into JAVA
 - Using a library of DB functions (API)
 - SQL/CLI (Call Level Interface, ODBC) → Open Database Connectivity: A widely used implementation of SQL/CLI
 - JDBC : Java Database Connectivity
 - Designing a brand-new language
 - Oracle PL/SQL : Oracle's procedural extension for SQL and Oracle's relational database.

Impedance mismatch is a term used in computer science to describe the problem that arises when two systems or components that are supposed to work together have different data models, structures, or interfaces that make communication difficult or inefficient.

Impedance Mismatch

- ◆ Problem occurs because of difference between the DB model & PL model  *Programming Language*
 - Attribute data type vs. data type of PL: binding
 - Mapping between query result data structure & data structure in PL: cursor (iterator variable) is used to loop over the tuples in a query result
 - Impedance mismatch is less a problem when a special DB PL is designed.

Typical Sequence of Interaction in DB Programming


◆ Client/server model

- Client program handles the logic of a software application
- Client includes some calls to one or more database server to access or update data

- Front-end vs. Back-end


◆ Common sequence

- Client program establish or open a connection to the database server
- Once the connection is established, program interact with the DB by submitting queries, updates or other commands
- Close the connection





Client-Server

There is a request/response protocol associated with any client-server architecture:



Client-Server (cont.)





Front end vs. Back end.

Summary

- ◆ SQL
- ◆ DDL, DML, View, Indexing, ...
- ◆ Embedd SQL + Host language
- ◆ Syntax

```
SELECT      <attribute list>
FROM        <table list>
[ WHERE      <condition> ]
[ GROUP BY <grouping attribute(s)> ]
[ HAVING    <group condition> ]
[ ORDER BY <attribute list> ]
```

SQL queries run in this order

FROM + JOIN



WHERE



GROUP BY



HAVING



SELECT (window functions
happen here !)



ORDER BY



LIMIT