

# Transaction Processing Concepts

政治大學  
資訊科學系  
沈錕坤



中國信託  
Chinatrust

# 自動櫃員機交易明細單

ATM Transaction(TX) Receipt

交 易 日 09 / 11 / 18 Date	時 間 11:22 Time	機 號 04953095 Atm No.	交易類別 CW7 TX Type
交易序號 455288 TX No.	發卡行編號 Isu Bank No.	帳 號 / 卡 號 498022*****0028 Account No.	
交 易 金 額 10,000 TX Amount		轉 入 帳 號 Trans for Act No.	
帳 戶 餘 額 179,818,931 Account Balance		可 用 餘 額 Available Balance	
訊 息 代 號 OK MSG Code	手 續 費 100 Fee	發鈔張數 千元 佰元 10 00 Cash detail	
授 權 碼 002316	交易代號：CWD取款・TWD轉帳 INO餘額查詢		

(幣別: 新台幣)

# Transaction Processing System

- ◆ System with large DB & hundreds of **concurrent users** that are executing DB **transactions**
- ◆ Require high **availability** & fast **response time** for hundreds of concurrent users

# Transaction

## ◆ Transaction

- A **logical unit** of DB processing
- Each transaction includes **one or more DB access operations** (insertion, deletion, modification, retrieval operations)
- e.g. 轉帳: 從帳號X轉帳N元到帳號Y

```
read(X);  
X=X-N;  
write(X);  
read(Y);  
Y=Y+N;  
write(Y);
```

# Interleaved Processing

- ◆ Modern computer
  - time-sharing system
  - Interleaved processing

```
read(X);  
X=X-N;  
write(X);  
read(Y);  
Y=Y+N;  
write(Y);
```

Serial  
processing

```
read(X);  
X=X+M;  
write(X);
```

```
read(X);  
X=X-N;
```

```
read(X);  
X=X+M;
```

```
write(X);  
read(Y);
```

```
write(X);  
Y=Y+N;
```

```
write(Y);
```

Interleaved  
processing

# Transaction Processing

- ◆ Concurrency control
- ◆ Recovery

# Concurrency Problem: Why Concurrency Control Is needed

# Concurrency Problems

- ◆ Types of concurrency problems
  - Lost update
  - Temporary update (dirty read): uncommitted dependency
  - Incorrect summary
  - Unrepeatable read



# Lost Update Problem

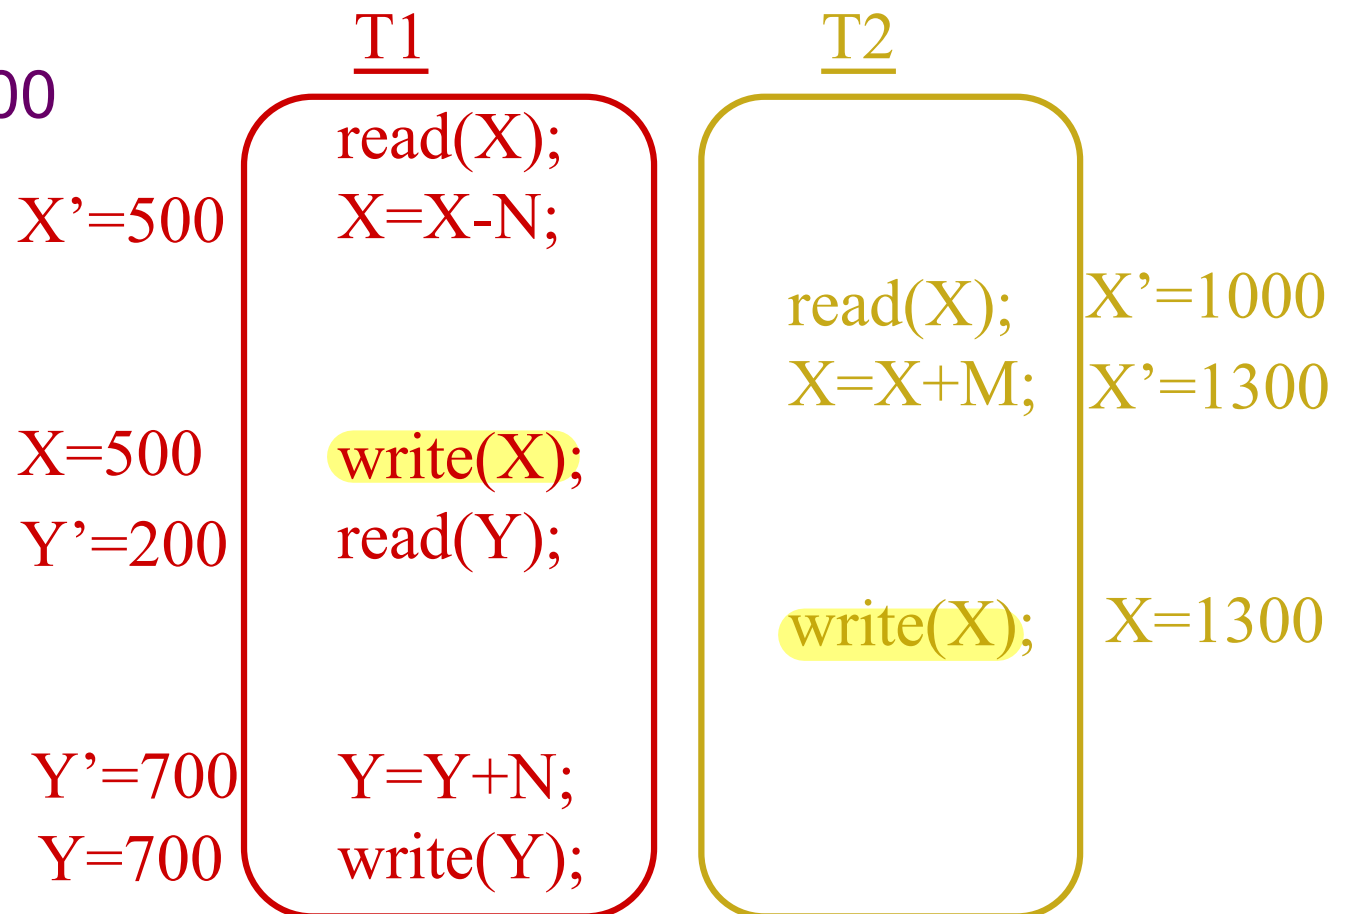
$X=1000, Y=200$

X transfers  $N=500$  to Y,

$M=300$  is transferred to X

$X=1000-500+300=800$

$Y=200+500=700$



# Temporary Update Problem

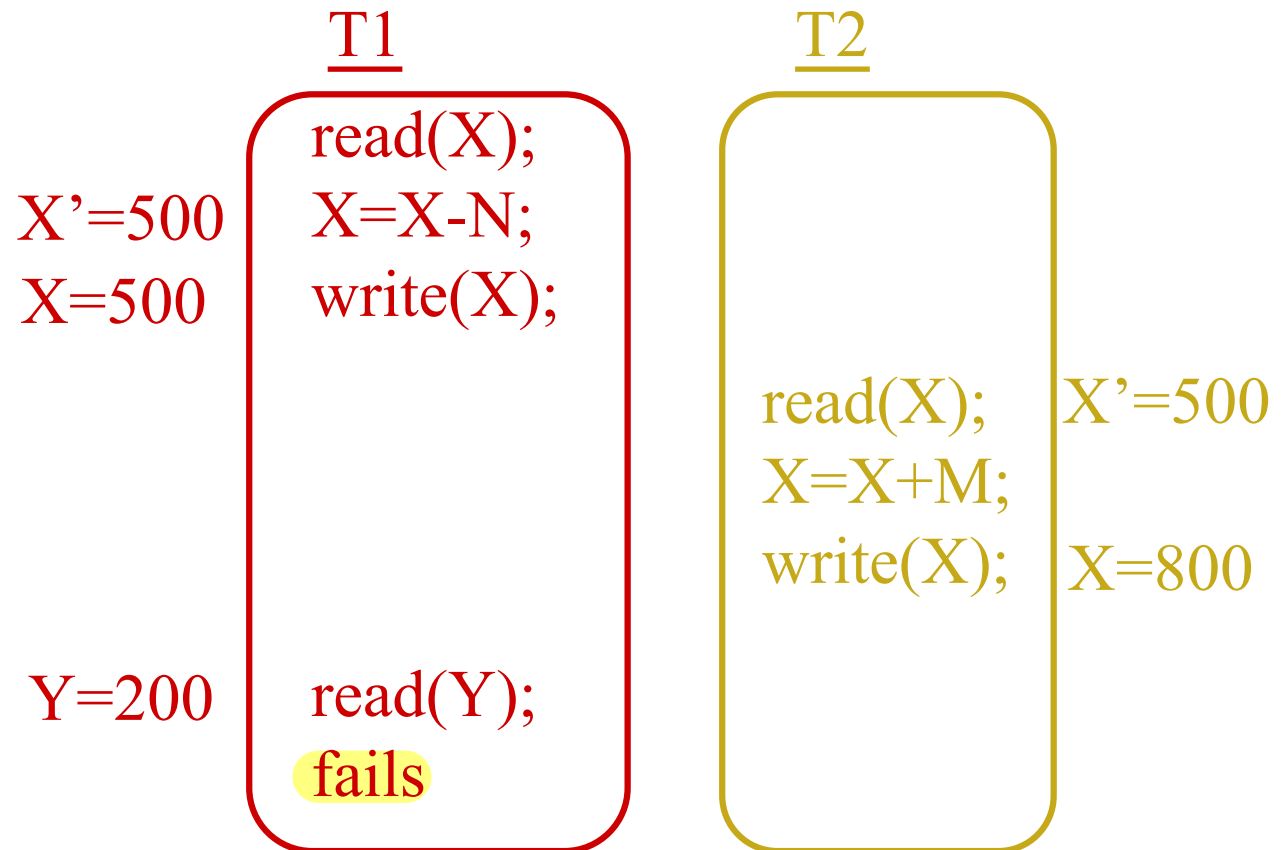
$X=1000, Y=200$

X transfers  $N=500$  to Y

$M=300$  is transferred to X

$X=1000-500+300=800$

$Y=200+500=700$

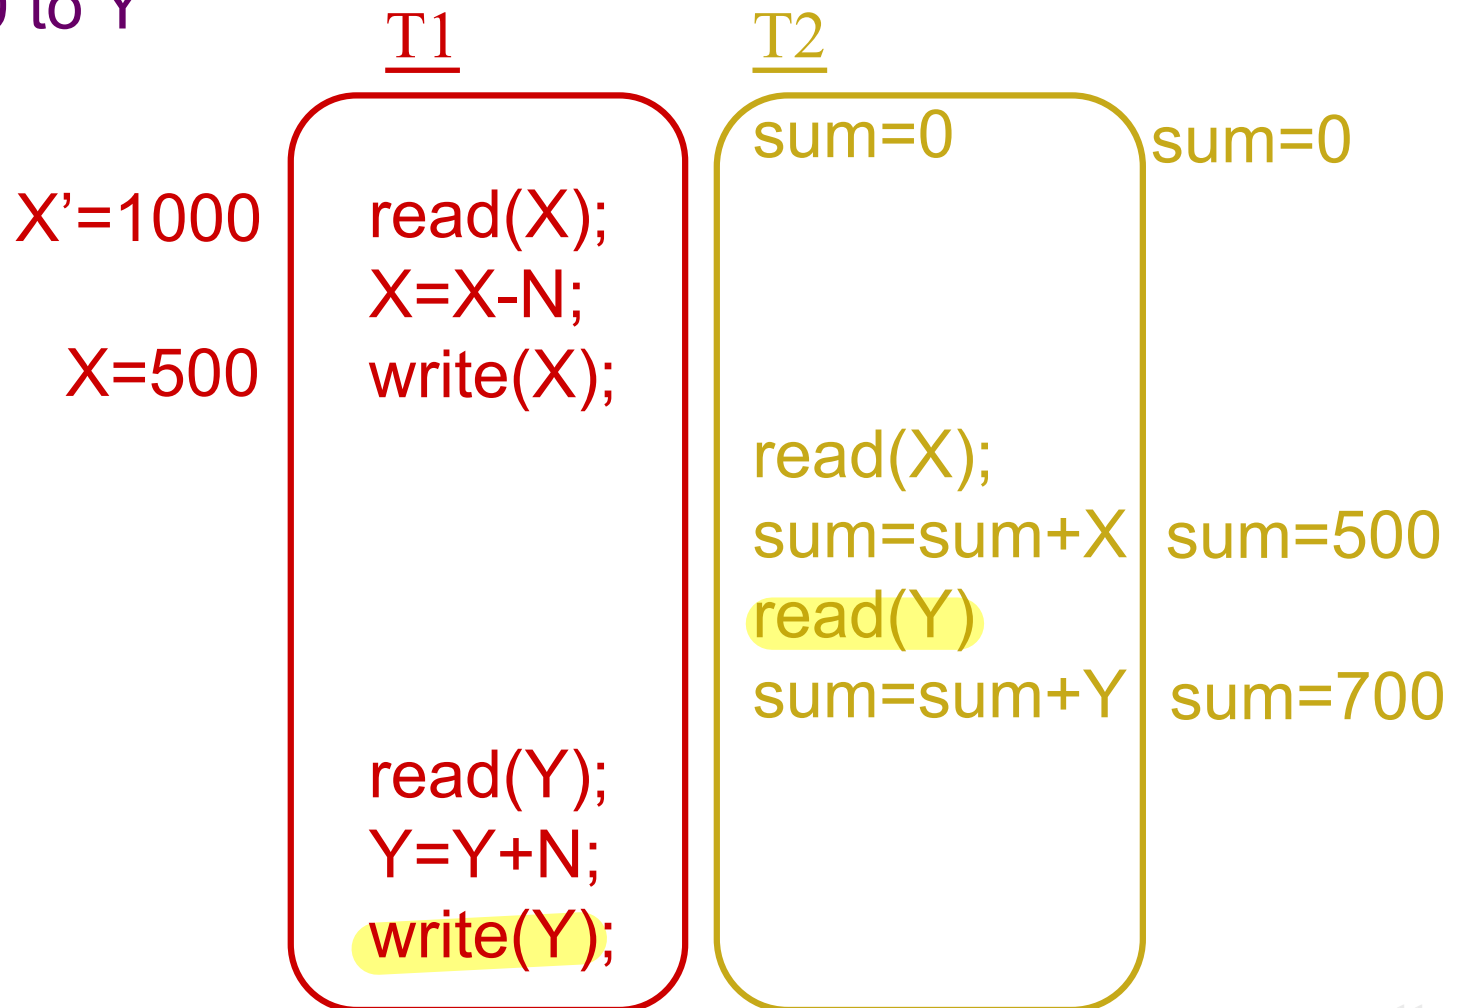


# Incorrect Summary Problem

X=1000, Y=200

X transfers N=500 to Y

X+Y=1200



# Unrepeated Read Problem

- ◆ A transaction  $T$  reads an item twice & the item is changed by another transaction  $T'$  between the two reads
- ◆  $T$  receives different values for its two reads of the same item

# Solution for Concurrency Control Problems

Critical

Section

## ◆ Locking techniques

$X=1000, Y=200$

X transfers  $N=500$  to Y,  
 $M=300$  is transferred to Y

$X=1000-500+300=800$

$Y=200+500=700$

$X'=1000$

$X'=500$

$X=500$

$Y'=200$

$Y'=700$

$Y=500$

交易 1

writelock(X)

read(X);

$X=X-N$ ;

write(X)

unlock(X)

writelock(X)

read(Y);

$Y=Y+N$ ;

write(Y);

unlock(X)

交易 二

writelock(X)

read(X);

$X=X+M$ ;

write(X);

unlock(X)

$X'=500$

$X'=800$

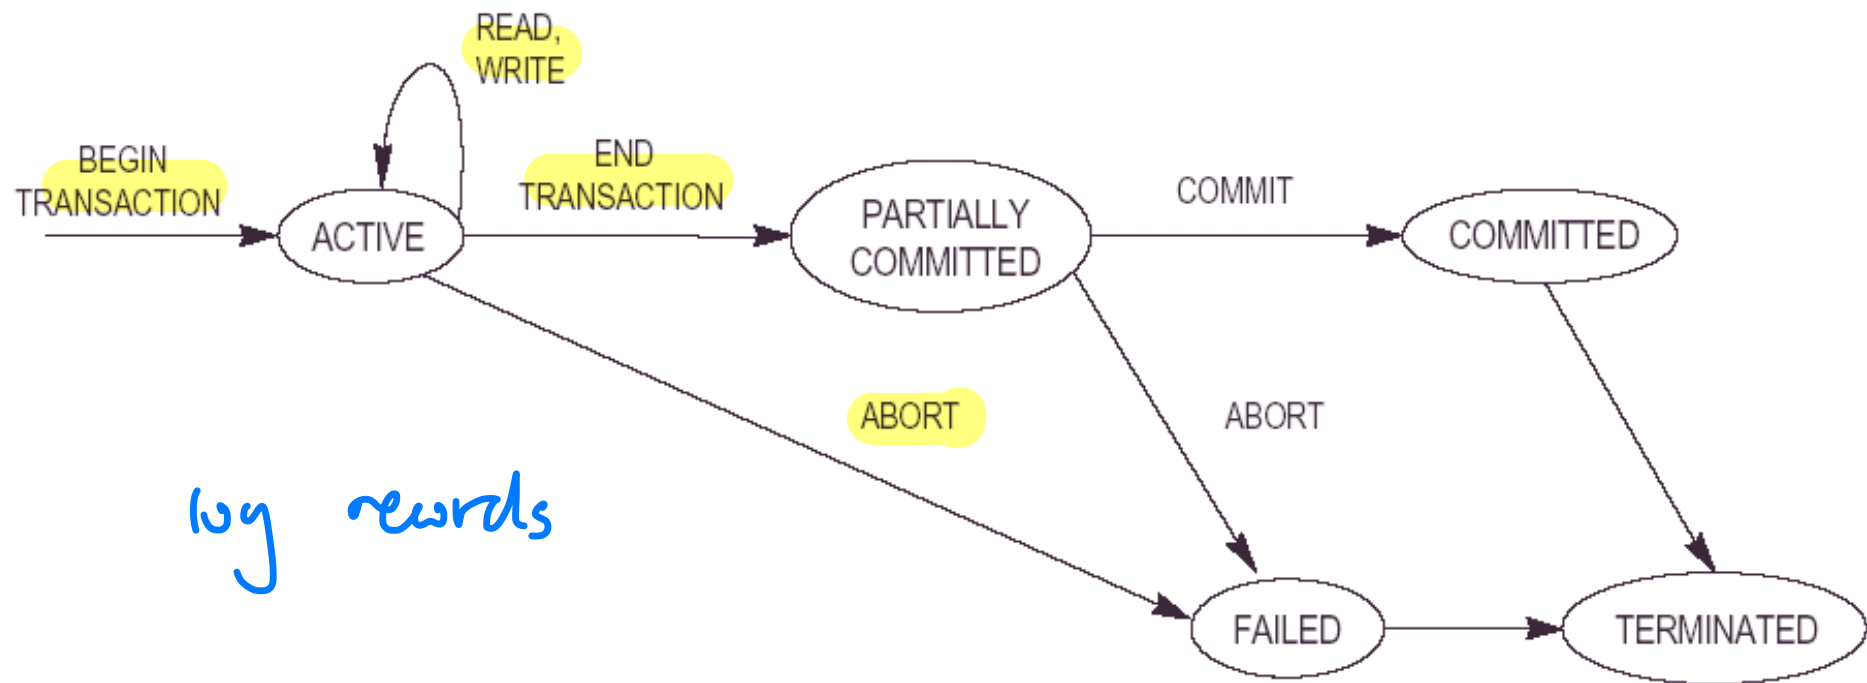
$X=800$

# Recovery Problem: Why Recovery Is needed

# Types of Failures

- ◆ Computer failure (system crash)
- ◆ Disk failure
- ◆ Transaction or system error (overflow, interrupt)
- ◆ Local errors or exception conditions detected by the transaction (e.g. data not found)
- ◆ Concurrency control enforcement (deadlock)
- ◆ Physical problems

# Transaction States





# System Log

- ◆ To be recovered from failures that affect transactions, system maintains a **log**
- ◆ Log keep track of all transaction operations that affect the values of DB items
- ◆ Log is kept on disk
- ◆ Log is **not affected** by any type of **failure** except for disk failure or catastrophe
- ◆ Log is **periodically backed up** to archival storage  
(e.g. magnetic tape)

# System Log (cont.)

- ◆ Log records
  - start transaction
  - transaction has changed value of DB item X
  - transaction has read the value of DB item X
  - transaction has completed successfully & committed to the DB
  - transaction has been aborted
- ◆ Undo, Redo
- ◆ Force-writing
  - before a transaction reaches its commit point, any portion of the log that has not been written to the disk yet must now be written to the disk

# ACID Properties of Transactions

## ◆ Atomicity

- A transaction is an atomic unit of processing
- It's either performed in its entirety or not performed at all "all or nothing"
- Responsibility of transaction **recovery** subsystem

## ◆ Consistency preservation

- The complete execution takes the DB from one consistent state to another
- Responsibility of **DB programmer or integrity constraint** of DBMS

# ACID Properties of Transactions (cont.)

## ◆ Isolation

- A transaction should appear as though it's being executed in isolation from other transactions
- Responsibility of **concurrency control** subsystem

## ◆ Durability (permanency)

- Changes applied to DB must not be lost because of any failure
- Responsibility of transaction **recovery** subsystem

# Schedules & Recoverability

# Interleaved Processing

- ◆ Modern computer
  - time-sharing system
  - Interleaved processing

```
read(X);  
X=X-N;  
write(X);  
read(Y);  
Y=Y+N;  
write(Y);
```

Serial  
processing

```
read(X);  
X=X+M;  
write(X);
```

```
read(X);  
X=X-N;  
read(X);  
X=X+M;
```

Interleaved  
processing

```
write(X);  
read(Y);  
write(X);  
Y=Y+N;  
write(Y);
```

# Schedule of Transactions

- ◆ A schedule  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is
  - An ordering of the operations of the transactions
  - Subject to the constraint that: for each transaction  $T_i$ , the operations of  $T_i$  in  $S$  must appear in the same order in which they occur in  $T_i$
  - However, operations from other transactions  $T_j$  can be interleaved with the operations of  $T_i$  in  $S$
- Partial ordering
  - a **partial order** on a set: an arrangement such that, for **some** pairs of elements, one precedes the other.

# Schedules of Transactions (cont.)

- ◆ For the purpose of recovery & concurrency control, only
  - read\_item (r)
  - write\_item (w)
  - commit (c)
  - abort (a)operations are considered



# Conflict Operations

- ◆ Two operations in a schedule are **conflict** if
  - They belong to **different transactions**
  - They **access the same item x**
  - **At least one** of the operations is a **write\_item(x)**
- ◆ E.g.
  - S<sub>a</sub>: r1(x); r2(x); w1(x); r1(y); w2(x); w1(y)
  - r1(x) & w2(x) conflict
  - r2(x) & w1(x) conflict
  - w2(x) & w1(y) do not conflict (different items)
  - r1(x) & w1(x) do not conflict (same transactions)

# Complete schedule

- ◆ A schedule  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is said to be a **complete schedule** if
  - All operations in the transactions must appear in the complete schedule including a **commit** or **abort** operations as the **last operations** for each transaction in the schedule
  - For any pair of operations from the **same transaction**  $T_i$ , their order of appearance in  $S$  is the same as their order of appearance in  $T_i$
  - For any **two conflicting** operations, one of the two must occur before the other in the schedule


# Recoverable Schedule


- ◆ Recoverable schedule: once a transaction is committed, it should never be necessary to roll back.
- ◆ A schedule  $S$  is <sup>not recoverable</sup> recoverable if
  - <sup>One</sup> No transaction  $T$  in  $S$  commits until <sup>before</sup> all transactions  $T'$  (that have written an item that  $T$  reads) have committed
- ◆ E.g.
  - $r1(x); \underline{w1(x)}; r2(x); r1(y); w2(x); \underline{c2}; a1$ ; not recoverable  $T=2, T'=\{1\}$
  - $r1(x); r2(x); w1(x); r1(y); w2(x); \underline{c2}; w1(y); c1$ ; recoverable
  - $r1(x); \underline{w1(x)}; r2(x); r1(y); w2(x); w1(y); \underline{c1}; c2$ ; recoverable
  - $r1(x); w1(x); r2(x); r1(y); w2(x); w1(y); a1; a2$ ;

# Serializability of Schedules

# Serial Schedules


- ◆ A schedule  $S$  is **serial**
  - if every transaction  $T$  in  $S$ , all the operations of  $T$  are executed **consecutively** in the schedule
  - Otherwise, **non-serial**
- ◆ In a serial schedule
  - **only one** transaction at a time is **active**
  - The commit (or abort) of the active transaction initiates execution of the next transaction





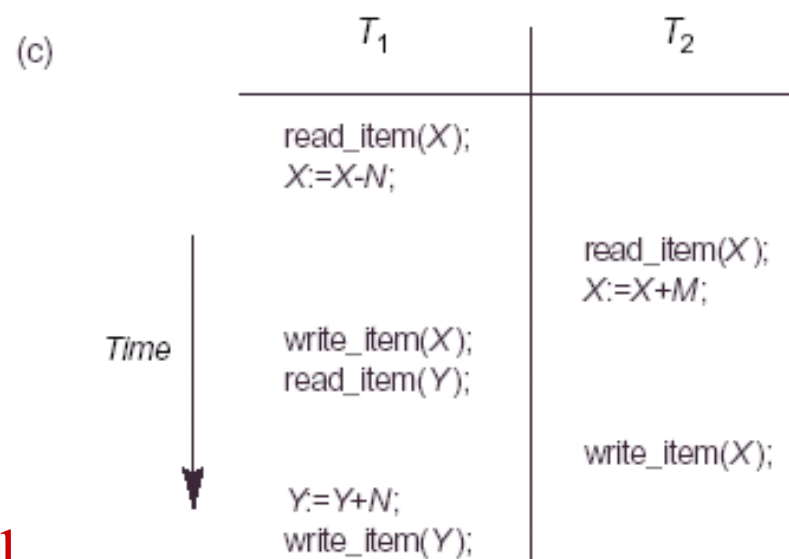
Schedule A

serial



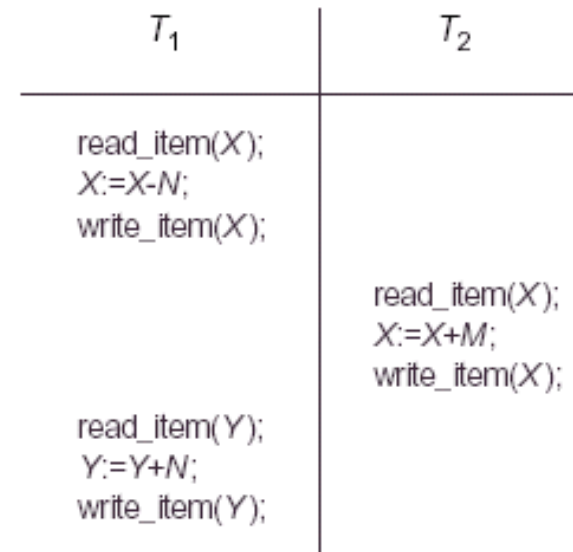
Schedule B

serial



Schedule C

not  
serial



Schedule D

not  
serial

# Serializable

- ◆ A schedule  $S$  of  $n$  transactions is **serializable** if it is equivalent to some **serial schedule** of the same  $n$  transactions
- ◆ If a non-serial schedule  $S$  is serializable
  - > it is correct

# Conflict Equivalent

- ◆ Two schedules are said to be **conflict equivalent** if the **order** of any two **conflicting operations** is the **same** in both schedules
- ◆ If two conflicting operations are applied in different orders in two schedules, the effect can be different

≡ one can be transformed to another by  
swapping consecutive non-conflicting  
operations




# Conflict Operations

- ◆ Two operations in a schedule are **conflict** if
  - They belong to **different transactions**
  - They **access the same item X**
  - **At least one** of the operations is a **write\_item(x)**
- ◆ E.g.
  - $S_a$ :  $r1(x); r2(x); w1(x); r1(y); w2(x); w1(y)$
  - $r1(x)$  &  $w2(x)$  conflict
  - $r2(x)$  &  $w1(x)$  conflict
  - $w2(x)$  &  $w1(y)$  do not conflict (different items)
  - $r1(x)$  &  $w1(x)$  do not conflict (same transactions)

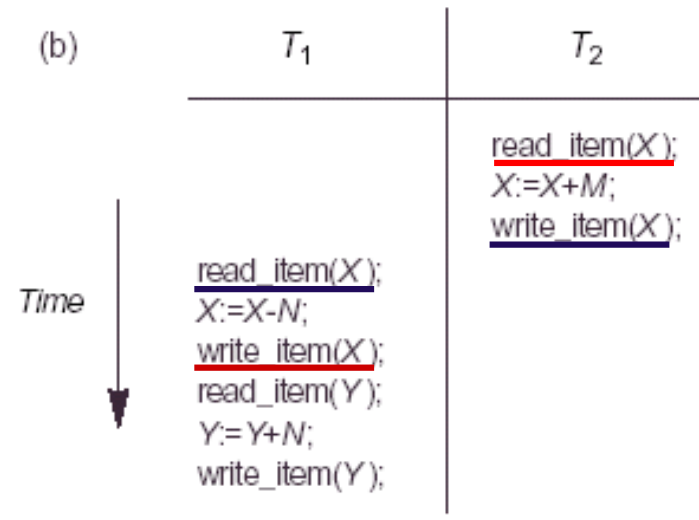
# Conflict Serializable

- ◆ A schedule  $S$  is **conflict serializable** if it is **conflict equivalent** to some **serial schedule**  $S'$


★ Conflict serializability is a conservative test for serializability. (may give some false negatives)




**(serial) schedule A**



**(serial) schedule B**



**(non-serializable) schedule C**



**(serializable) schedule D**

# Testing for Conflict Serializability

- ◆ There exists a simple algorithm to test for conflict serializability
- ◆ However
  - most concurrency control methods do not actually test for serializability
  - Protocols or rules are developed to guarantee that a schedule is serializability (e.g. two phase-locking)

# Testing for Conflict Serializability (cont.)

- ♦ Algorithm to test for conflict serializability
  - precedence graph
    - A directed graph  $G = (N, E)$
    - $N = \{T_1, T_2, \dots, T_n\}$ , one node for each transaction  $T_i$  in the schedule
    - $E = \{T_j \rightarrow T_k \mid \text{one of operations in } T_j \text{ appears in the schedule before some conflicting operation in } T_k\}$

## ♦ Algorithm

1. construct each node in  $G$  for each transaction  $T_i$
2. create an edge  $(T_i \rightarrow T_j)$  in  $G$  if  $w_i(x)$  appears before  $r_j(x)$  in  $S$
3. create an edge  $(T_i \rightarrow T_j)$  in  $G$  if  $r_i(x)$  appears before  $w_j(x)$  in  $S$
4. create an edge  $(T_i \rightarrow T_j)$  in  $G$  if  $w_i(x)$  appears before  $w_j(x)$  in  $S$
5. The schedule  $S$  is (conflict) **serializable** if and only if

$G$  has no cycles


conflicting operations




## Testing for Conflict Serializability (cont.)

(acyclic)

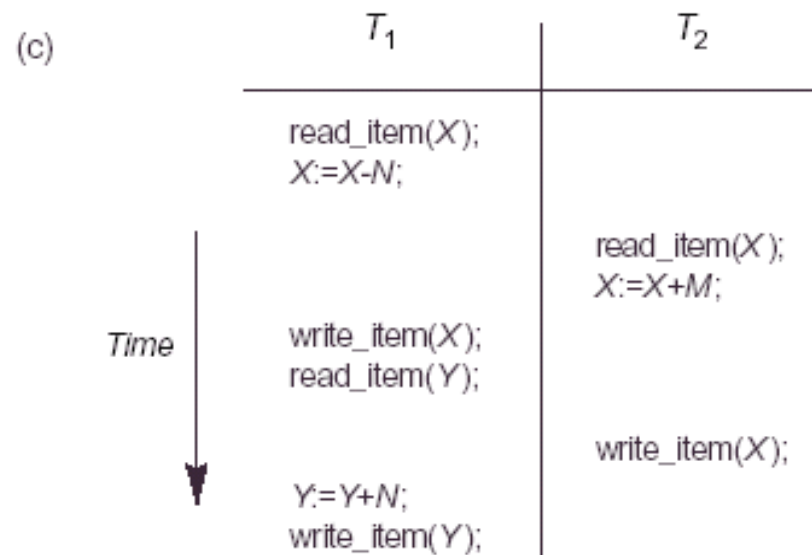
- ◆ If there is no cycle in the precedence graph, we can create an equivalent serial schedule  $S'$  that is equivalent to  $S$




**(serial) schedule A**




**(serial) schedule B**



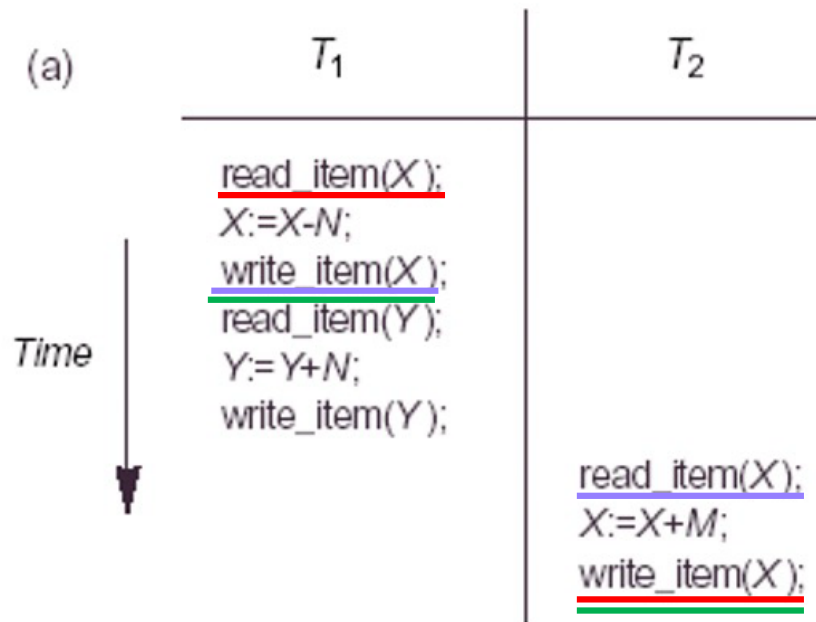
**(non-serializable) schedule C**




**(serializable) schedule D**



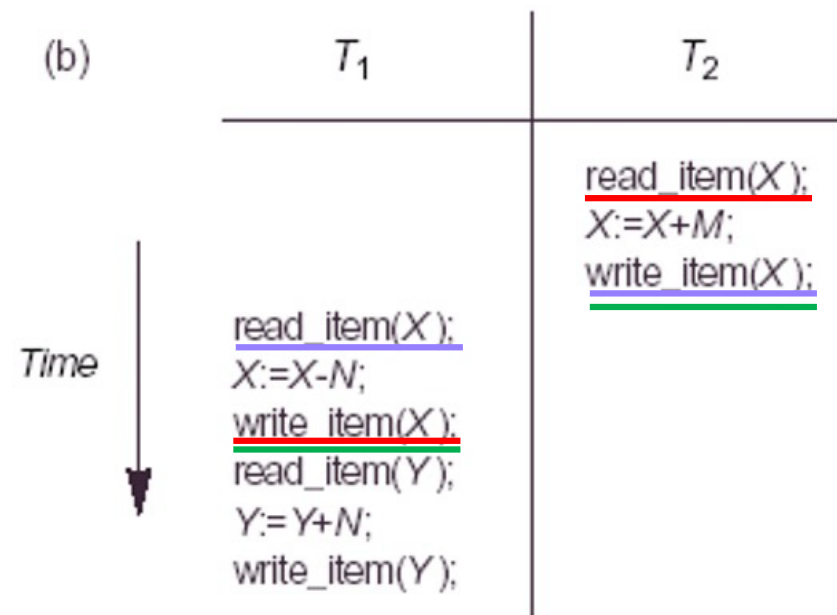
(serial) **schedule A**



Schedule A




(serial) **schedule B**




Schedule B

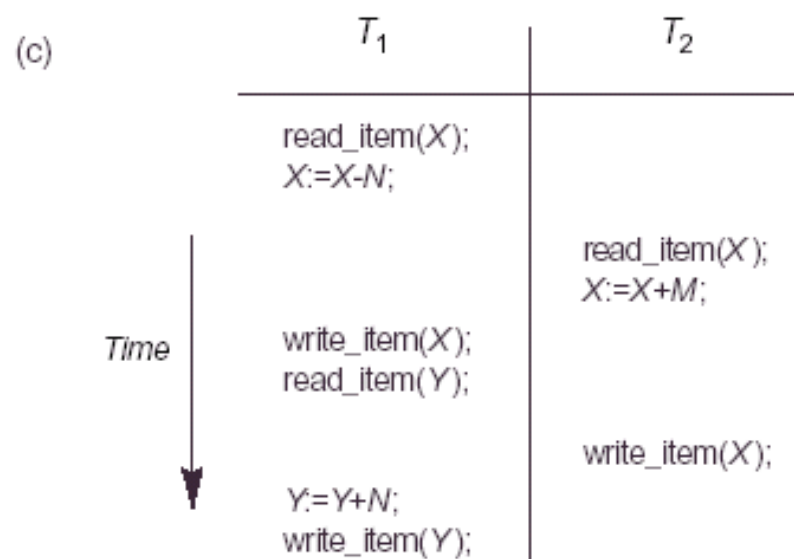





**(serial) schedule A**



**(serial) schedule B**




**(non-serializable) schedule C**




**(serializable) schedule D**


(c)

**(non-serializable)****schedule C**


(d)

**(serializable)****schedule D**

(c)



Schedule C



Schedule D


(a)

transaction $T_1$	transaction $T_2$	transaction $T_3$
<code>read_item (X);</code> <code>write_item (X);</code> <code>read_item (Y);</code> <code>write_item (Y);</code>	<code>read_item (Z);</code> <code>read_item (Y);</code> <code>write_item (Y);</code> <code>read_item (X);</code> <code>write_item (X);</code>	<code>read_item (Y);</code> <code>read_item (Z);</code> <code>write_item (Y);</code> <code>write_item (Z);</code>

(b)

	transaction $T_1$	transaction $T_2$	transaction $T_3$
		<code>read_item (Z);</code> <code>read_item (Y);</code> <code>write_item (Y);</code>	<code>read_item (Y);</code> <code>read_item (Z);</code>
Time ↓	<code>read_item (X);</code> <code>write_item (X);</code>		<code>write_item (Y);</code> <code>write_item (Z);</code>
	<code>read_item (Y);</code> <code>write_item (Y);</code>	<code>read_item (X);</code>  <code>write_item (X);</code>	

Schedule E



**Schedule E**


(a)

transaction $T_1$	transaction $T_2$	transaction $T_3$
<code>read_item (X);</code> <code>write_item (X);</code> <code>read_item (Y);</code> <code>write_item (Y);</code>	<code>read_item (Z);</code> <code>read_item (Y);</code> <code>write_item (Y);</code> <code>read_item (X);</code> <code>write_item (X);</code>	<code>read_item (Y);</code> <code>read_item (Z);</code> <code>write_item (Y);</code> <code>write_item (Z);</code>

(c)

	transaction $T_1$	transaction $T_2$	transaction $T_3$
			<code>read_item (Y);</code> <code>read_item (Z);</code>
Time ↓	<code>read_item (X);</code> <code>write_item (X);</code>	<code>read_item (Z);</code>	<code>write_item (Y);</code> <code>write_item (Z);</code>
	<code>read_item (Y);</code> <code>write_item (Y);</code>	<code>read_item (Y);</code> <code>write_item (Y);</code> <code>read_item (X);</code> <code>write_item (X);</code>	

Schedule F



**Schedule F**

# Uses of Serializability

- ◆ Serial schedule
  - inefficient processing (no interleaving)
- ◆ Serializable schedule
  - In practice, interleaving is determined by O.S.scheduler
  - It is difficult to determine how operations of a schedule will be interleaved beforehand to ensure serializability
  - Most practical DBMS use methods that ensure serializability, without having to test the schedules

# Uses of Serializability (cont.) \*

- ◆ Protocols used in practical DBMS
  - Two-phase locking (2PL)
    - Locking data items to prevent concurrent transactions from interleaving with one another
    - Enforcing additional condition that guarantees serializability
    - Used in most DBMS
  - Timestamp ordering
    - Each transaction is assigned a unique timestamp
    - Protocol ensures that any conflicting operations are executed in the order of the transaction timestamps
  - Multiversion protocol
    - Maintain multiple versions of data items
  - Optimistic protocol
    - Check for possible serializability violations after transactions terminate but before they are permitted to commit

# Concurrency Control Technique: Two Phase Locking

# Locks

## ◆ Lock

- a **variable** associated with a data item that describes the status of the item with respect to possible operations that can be applied to it.
- there is **one lock** for each data item in the DB.
- Locks are used as a means of synchronizing the access by concurrent transactions to the database items.



# Binary Locks

## ◆ A binary lock

- have two states or values: locked and unlocked (1 and 0).
- If the value of the lock on X is 1, item X cannot be accessed by a database operation that requests the item.
- If the value of the lock on X is 0, the item can be accessed when requested, and the lock value is changed to 1.
- a binary lock enforces **mutual exclusion** on the data item
- lock(X): current state of the lock associated with item X

## ◆ Two operations

- lock\_item
- unlock\_item.

```

lock_item(X):
B:  if LOCK(X) = 0          (* item is unlocked *)
      then LOCK(X) ← 1      (* lock the item *)
    else
      begin
      wait (until LOCK(X) = 0
            and the lock manager wakes up the transaction);
      go to B
      end;
unlock_item(X):
    LOCK(X) ← 0;             (* unlock the item *)
    if any transactions are waiting
      then wakeup one of the waiting transactions;

```

\* lock\_item and unlock\_item operations must be implemented as **indivisible units** (known as **critical sections** in operating systems)

# Read/Write Locks

- ◆ 3 locking operations
  - `read_lock(X)`
  - `write_lock(X)`
  - `unlock(X)`.
- ◆ `Lock(X)`, has three possible states
  - read-locked
  - write-locked
  - unlocked.
- ◆ **read**-locked = **share**-locked, because other transactions are allowed to read the item,
- ◆ **write**-locked = **exclusive**-locked, because a single transaction exclusively holds the lock on the item.

# Read/Write Locks (cont.)

1. A transaction  $T$  must issue the operation  $\text{read\_lock}(X)$  or  $\text{write\_lock}(X)$  before any  $\text{read\_item}(X)$  operation is performed in  $T$ .
2. A transaction  $T$  must issue the operation  $\text{write\_lock}(X)$  before any  $\text{write\_item}(X)$  operation is performed in  $T$ .
3. A transaction  $T$  must issue the operation  $\text{unlock}(X)$  after all  $\text{read\_item}(X)$  and  $\text{write\_item}(X)$  operations are completed in  $T$ .
4. A transaction  $T$  will not issue a  $\text{read\_lock}(X)$  operation if it already holds a read (shared) lock or a write (exclusive) lock on item  $X$ . (This rule may be relaxed.)
5. A transaction  $T$  will not issue a  $\text{write\_lock}(X)$  operation if it already holds a read (shared) lock or write (exclusive) lock on item  $X$ . (This rule may also be relaxed.)
6. A transaction  $T$  will not issue an  $\text{unlock}(X)$  operation unless it already holds a read (shared) lock or a write (exclusive) lock on item  $X$ .

**read\_lock(X):**

```
B:  if LOCK(X) = "unlocked"
    then begin LOCK(X) ← "read-locked";
        no_of_reads(X) ← 1
    end
else if LOCK(X) = "read-locked"
    then no_of_reads(X) ← no_of_reads(X) + 1
else begin
    wait (until LOCK(X) = "unlocked"
        and the lock manager wakes up the transaction);
    go to B
end;
```

**write\_lock(X):**

```
B:  if LOCK(X) = "unlocked"
    then LOCK(X) ← "write-locked"
else begin
    wait (until LOCK(X) = "unlocked"
        and the lock manager wakes up the transaction);
    go to B
end;
```

**unlock (X):**

```
if LOCK(X) = "write-locked"
    then begin LOCK(X) ← "unlocked";
        wakeup one of the waiting transactions, if any
    end
else if LOCK(X) = "read-locked"
    then begin
        no_of_reads(X) ← no_of_reads(X) - 1;
        if no_of_reads(X) = 0
            then begin LOCK(X) = "unlocked";
                wakeup one of the waiting transactions, if any
            end
        end
    end;
```

# Lock Conversion

- ◆ A transaction that already holds a lock on item  $X$  is allowed under certain conditions to **convert** the lock from one locked state to another.
- ◆ For example, it is possible for a transaction  $T$  to issue a **read\_lock**( $X$ ) and then later to **upgrade** the lock by issuing a **write\_lock**( $X$ ) operation.
- ◆ If  $T$  is the only transaction holding a read lock on  $X$  at the time it issues the **write\_lock**( $X$ ) operation, the lock can be upgraded; otherwise, the transaction must wait.
- ◆ It is also possible for a transaction  $T$  to issue a **write\_lock**( $X$ ) and then later to **downgrade** the lock by issuing a **read\_lock**( $X$ ) operation.

# Two Phase Locking

- ◆ A transaction is said to follow the two-phase locking protocol if all locking operations (read\_lock, write\_lock) precede the first unlock operation in the transaction.
- ◆ Two phases
  1. Expanding or growing phase: new locks on items can be acquired but none can be released
  2. Shrinking phase: existing locks can be released but no new locks can be acquired.
- ◆ It can be proved that, if every transaction in a schedule follows the two-phase locking protocol, the schedule is guaranteed to be serializable, obviating the need to test for serializability of schedules.

X=20, Y=30

$T_1$	$T_2$
<pre> read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X); </pre>	<pre> read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y); </pre>

T1, T2:  $X=50, Y=80$

T2, T1:  $Y=50, X=70$

(c)

$$Y' = 30$$

Time

$X' = 20$

$X' = 50$

**X = 50**

$T_1$	$T_2$
<pre> read_lock(Y); read_item(Y); unlock(Y);  write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X); </pre>	<pre> read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y); </pre>

$X' = 20$

$$Y' = 30$$
$$Y = 50$$



$T_1$	$T_2$
read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); $X := X + Y$ ; write_item(X); unlock(X);	read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); $Y := X + Y$ ; write_item(Y); unlock(Y);

$T_1'$	$T_2'$
read_lock(Y); read_item(Y); write_lock(X); unlock(Y) read_item(X); $X := X + Y$ ; write_item(X); unlock(X);	read_lock(X); read_item(X); write_lock(Y); unlock(X) read_item(Y); $Y := X + Y$ ; write_item(Y); unlock(Y);

# Transaction Support in SQL

EXEC SQL SET Transaction

Read Write

Diagnostics size 5

Isolation Level Serializable;

EXEC SQL Insert into Employee ...;

EXEC SQL Update Employee...;

EXEC SQL Commit;

Goto The\_End;

Undo; EXEC SQL Rollback

The\_End

# Conclusions

- ◆ Concurrent transactions
- ◆ Concurrency Problems
- ◆ Recovery Problems
- ◆ Transaction State
- ◆ ACID Properties
- ◆ Recoverable schedules
- ◆ Serializable schedules
- ◆ 2 Phase Locking