

Course Topics

- **Intelligent agents (AIMA Ch. 2)**
- Search (AIMA Ch. 3, 4, 6, 5)
- Reasoning (AIMA Ch. 7-9, 12-15)
- Machine learning (AIMA Ch. 19-20)
- Deep learning (AIMA Ch. 22)
- Natural language processing (AIMA Ch. 24)
- Generative AI (Optional)

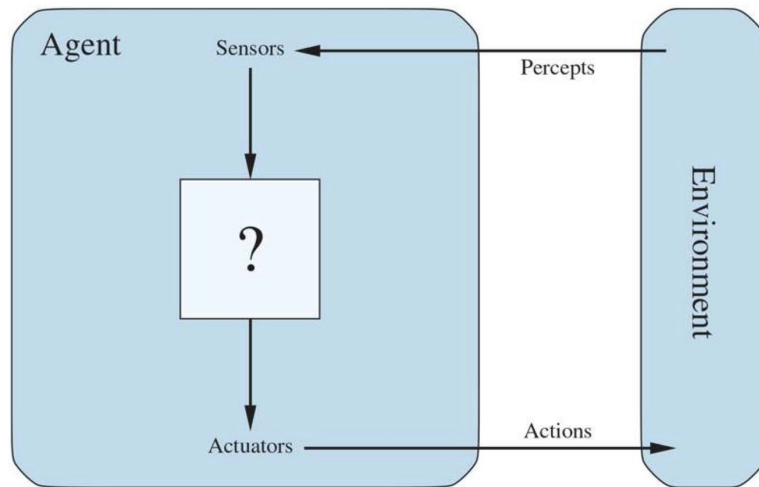
Intelligent Agent

L.-Y. Wei

Spring 2024

Agent


- An **agent** is anything that can be viewed as perceiving its **environment** through **sensors** and acting upon that environment through **actuators**
 - Robotic agent
 - e.g., camera/robotic arm
 - Software agent
 - e.g., keyboard/writing files
 - ...




- An agent's behavior is described by the **agent function** that maps any given percept sequence to an action $f(\text{percept sequence}) = \text{action}$

Example: Vacuum World

- Percepts
 - Location: A or B
 - Status: clean or dirty
- Actions
 - Move right
 - Move left
 - Suck (clean the square that it occupies)
 - NoOp



Percept sequence	Action
$[A, \textit{Clean}]$	\textit{Right}
$[A, \textit{Dirty}]$	\textit{Suck}
$[B, \textit{Dirty}]$	\textit{Suck}
$[A, \textit{Clean}], [A, \textit{Dirty}]$	\textit{Suck}
\vdots	\vdots
$[A, \textit{Clean}], [A, \textit{Clean}], [A, \textit{Clean}]$	\textit{Right}
$[A, \textit{Clean}], [A, \textit{Clean}], [A, \textit{Dirty}]$	\textit{Suck}
\vdots	\vdots



What is the **right** function?

function REFLEX-VACUUM-AGENT($[location, status]$) **returns** an action

if $status = \textit{Dirty}$ **then return** \textit{Suck}
else if $location = A$ **then return** \textit{Right}
else if $location = B$ **then return** \textit{Left}

Rational Agent

- A **rational agent** should select an action that is **expected** to **maximize** its **performance measure**, given the percept sequence to date

* to date = until now

To design a rational agent,
we must specify the **task environment**

Task Environment

- PEAS
 - **P**erformance
 - **E**nvironment
 - **A**ctuators
 - **S**ensors

Example - An Automated Taxi Driver

Agent Type	Performance Measure	Environment	Actuators	Sensors
Taxi driver	Safe, fast, legal, comfortable trip, maximize profits, minimize impact on other road users	Roads, other traffic, police, pedestrians, customers, weather	Steering, accelerator, brake, signal, horn, display, speech	Cameras, radar, speedometer, GPS, engine sensors, accelerometer, microphones, touchscreen

Task Environment Types

- Fully observable vs. partially observable
 - Fully observable
 - If an agents' sensors give it access to the complete state of the environment at each point in time
 - Partially observable
 - e.g., a vacuum agent with only a local dirt sensor cannot tell whether there is dirt in other squares

Task Environment Types (cont.)

- Deterministic vs. nondeterministic
 - Deterministic
 - The next state of the environment is completely determined by the current state and the action
 - Nondeterministic
 - e.g., there's a chance of rain tomorrow
(the possibilities are listed without being quantified)
 - **Stochastic**
 - A model of environment explicitly deals with probabilities

Task Environment Types (cont.)

- Episodic vs. sequential

- Episodic

- In each episode, the agent receives a percept and then performs a single action
 - The next episode does not depend on the actions taken in previous episodes
 - e.g., many classification tasks

- Sequential

- The current decision could affect all future decisions



Task Environment Types (cont.)

- Static vs. dynamic

- Dynamic

- If the environment can change while an agent is deliberating, the environment is dynamic for that agent

- **Semidynamic**

- The environment does not change with time
 - The agent's performance score changes with the time

(a)

(v) engage in long & careful consideration



Task Environment Types (cont.)

- Discrete vs. continuous
 - Discrete
 - e.g., the chess environment has a finite number of distinct states
 - Continuous
 - e.g., taxi driving is a continuous-time problem
 - A continuous-state
 - Actions are continuous (steering angles, etc.)
- Single-agent vs. multiagent

Examples of Task Environment Types

	Observable?	Agents?	Deterministic?	Episodic?	Static?	Discrete?
Crossword puzzle	Fully	Single	Deterministic	Sequential	Static	Discrete
Go with a clock	Fully	Multi	Deterministic	Sequential	Semi	Discrete
Poker	Partially	Multi	Stochastic	Sequential	Static	Discrete
Taxi driving	Partially	Multi	Stochastic	Sequential	Dynamic	Continuous

Examples of Task En

Eric Wallace*
UC BerkeleyNicholas Tomlin*
UC BerkeleyAlbert Xu*
UC BerkeleyKevin Yang*
UC BerkeleyEshaan Pathak*
UC BerkeleyMatthew L. Ginsberg
Matthew Ginsberg, LLCDan Klein
UC Berkeley

{ericwallace, nicholas_tomlin, albertxu3, klein}@berkeley.edu

Abstract

We present the Berkeley Crossword Solver, a state-of-the-art approach for automatically solving crossword puzzles. Our system works by generating answer candidates for each crossword clue using neural question answering models and then combines loopy belief propagation with local search to find full puzzle solutions. Compared to existing approaches, our system improves exact puzzle accuracy from 71% to 82% on crosswords from *The New York Times* and obtains 99.9% letter accuracy on themeless puzzles. Additionally, in 2021, a hybrid of our system and the existing Dr.Fill system outperformed all human competitors for the first time at the American Crossword Puzzle Tournament. To facilitate research on question answering and crossword solving, we analyze our system's remaining errors and release a dataset of over six million question-answer pairs.



Figure 1: A partially-solved example crossword puzzle from the 2021 American Crossword Puzzle Tournament, where our system scored higher than all 1033 human solvers. The highlighted fill KUNGFU answers the wordplay clue: *Something done for kicks?*

	Observable?	Agent					
Crossword puzzle	Fully	Single					
Go with a clock	Fully	Multi					
Poker	Partially	Multi	Stochastic	Sequential	Static	Discrete	
Taxi driving	Partially	Multi	Stochastic	Sequential	Dynamic	Continuous	

Real world: partially, multi-agent, nondeterministic, sequential, dynamic, continuous

Agent Structure

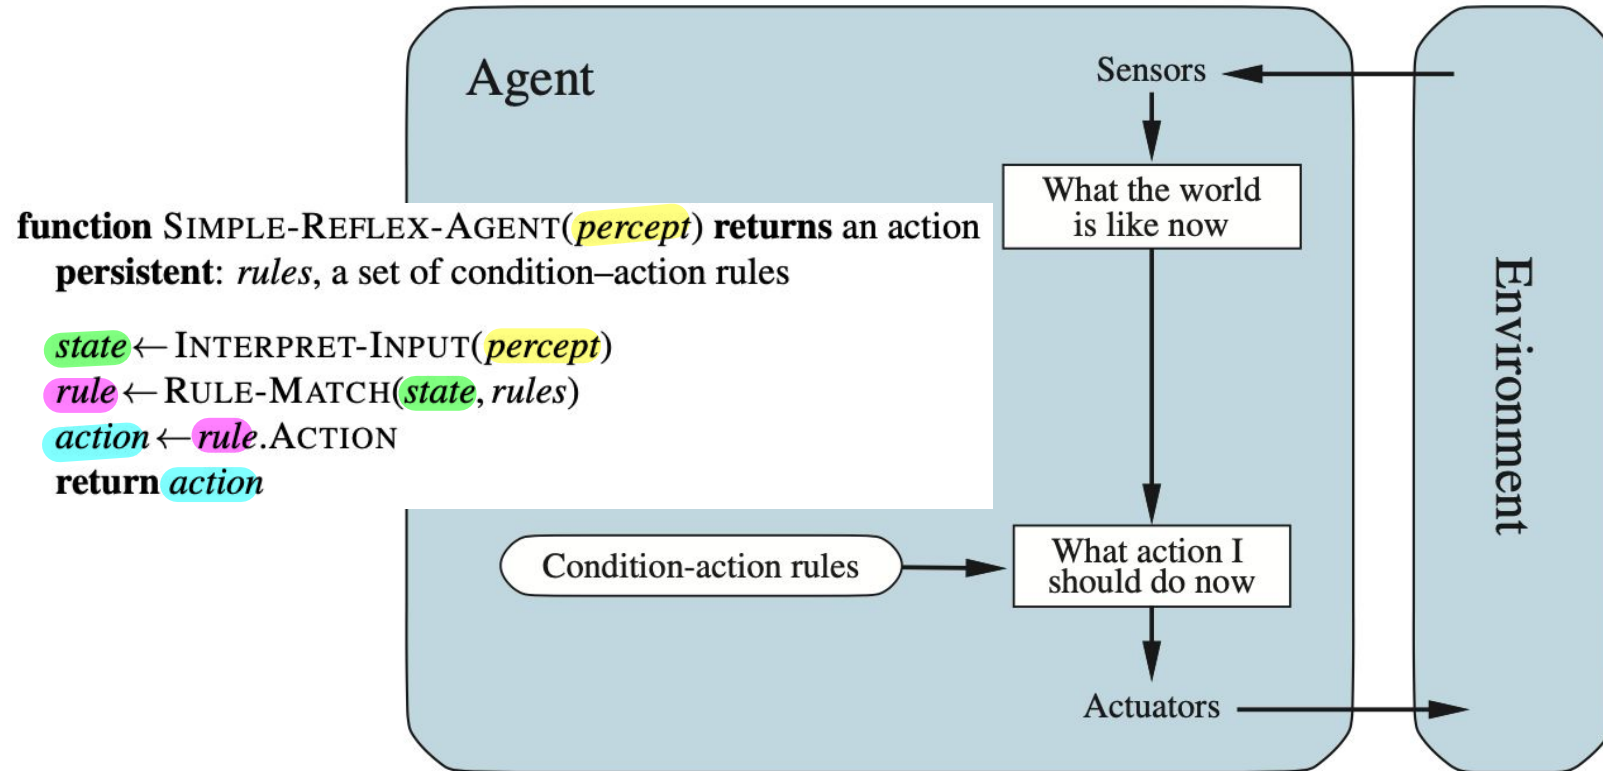
- Agent = architecture + program
 - Agent architecture includes physical sensors and actuators
 - **Agent program** runs on the physical architecture to implement the agent function (the job of AI)

Agent Types

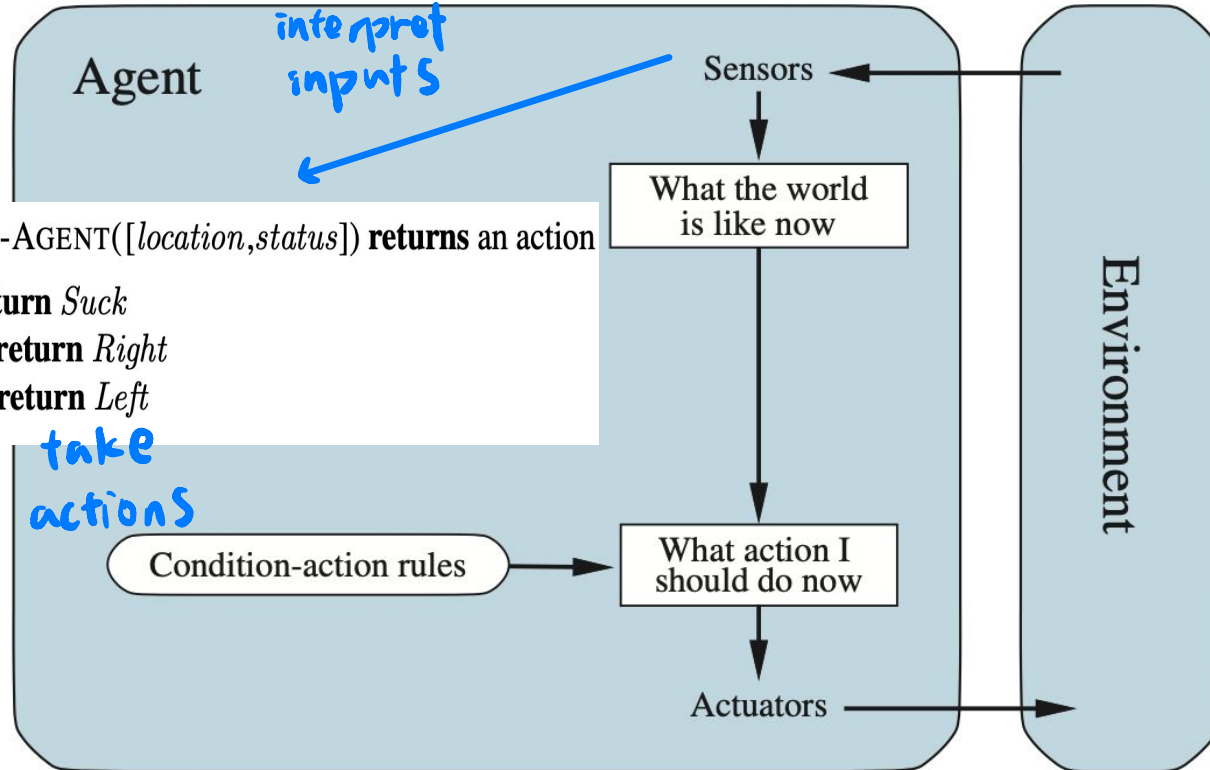
- Simple reflex agents
- Model-based reflex agents
- Goal-based agents
- Utility-based agents



Simple Reflex Agent



Simple Reflex Agent



function REFLEX-VACUUM-AGENT([location,status]) **returns** an action

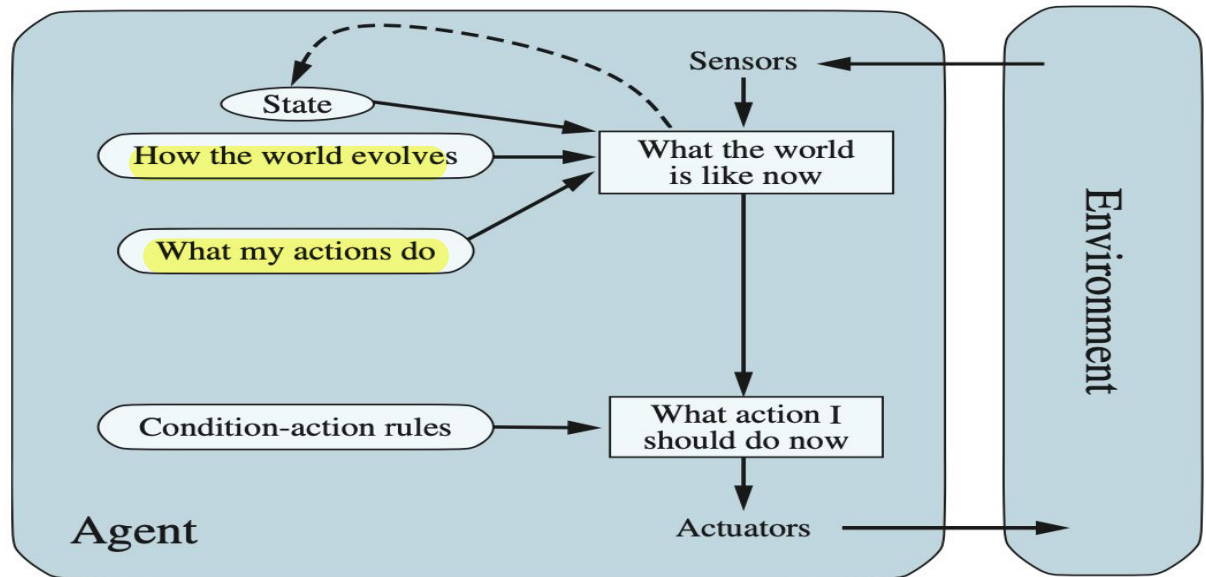
if *status* = *Dirty* **then return** *Suck*
else if *location* = *A* **then return** *Right*
else if *location* = *B* **then return** *Left*



Model-based Reflex Agent

- Keep track the part of the world it cannot see now using an **internal model** (Partial observability)
- Then choose an action in the same way as the reflex agent

Knowledge:
Transition model of the world
Sensor model



function MODEL-BASED-REFLEX-AGENT(*percept*) **returns** an action

persistent: *state*, the agent's current conception of the world state

internal model { *transition_model*, a description of how the next state depends on the current state and action
sensor_model, a description of how the current world state is reflected in the agent's percepts
rules, a set of condition–action rules
action, the most recent action, initially none

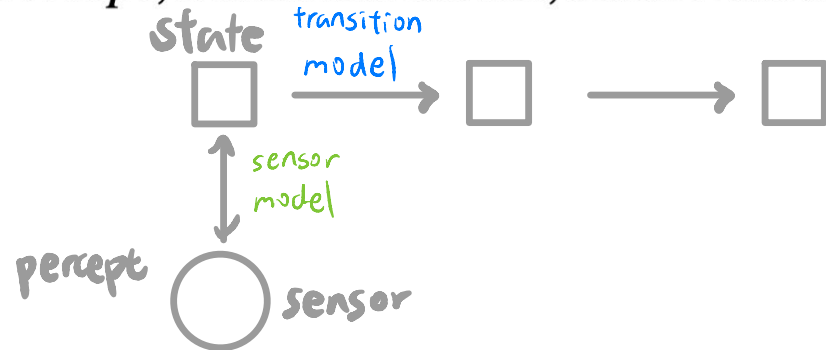
state ← UPDATE-STATE(*state*, *action*, *percept*, *transition_model*, *sensor_model*)

rule ← RULE-MATCH(*state*, *rules*)

action ← *rule*.ACTION

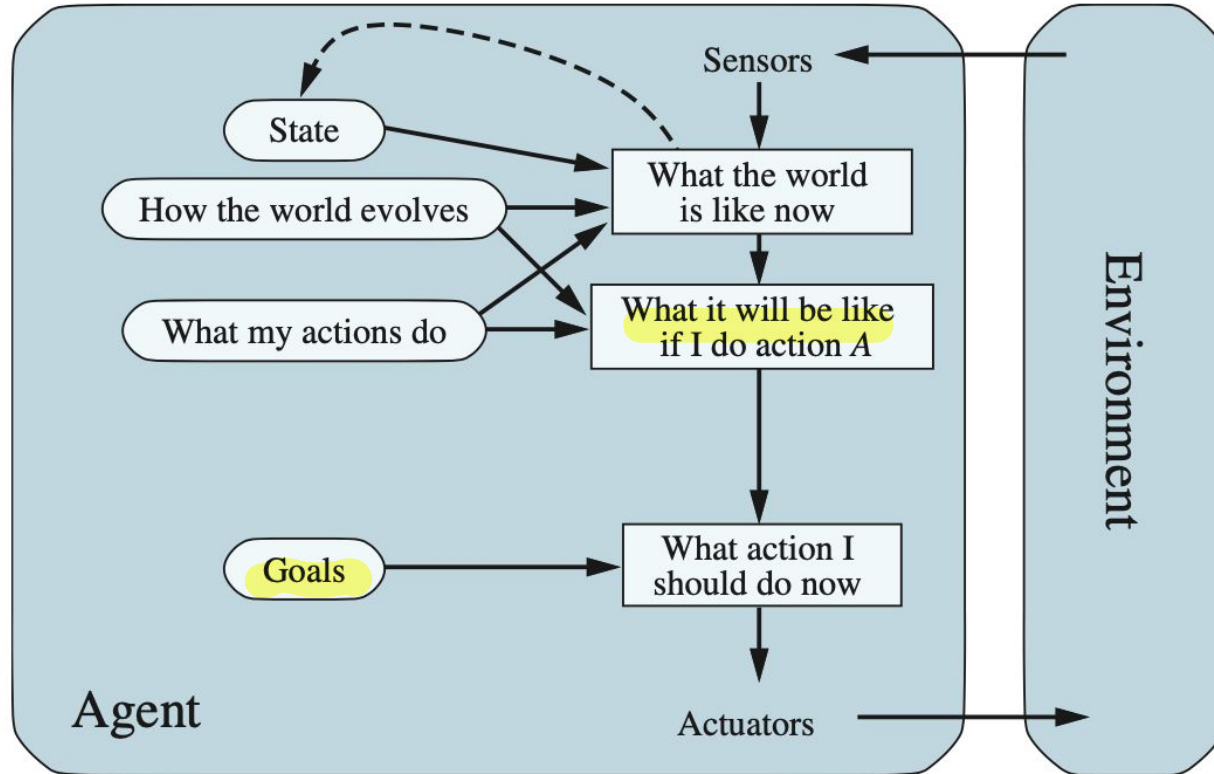
return *action*

choose an action the same way
as a reflex agent



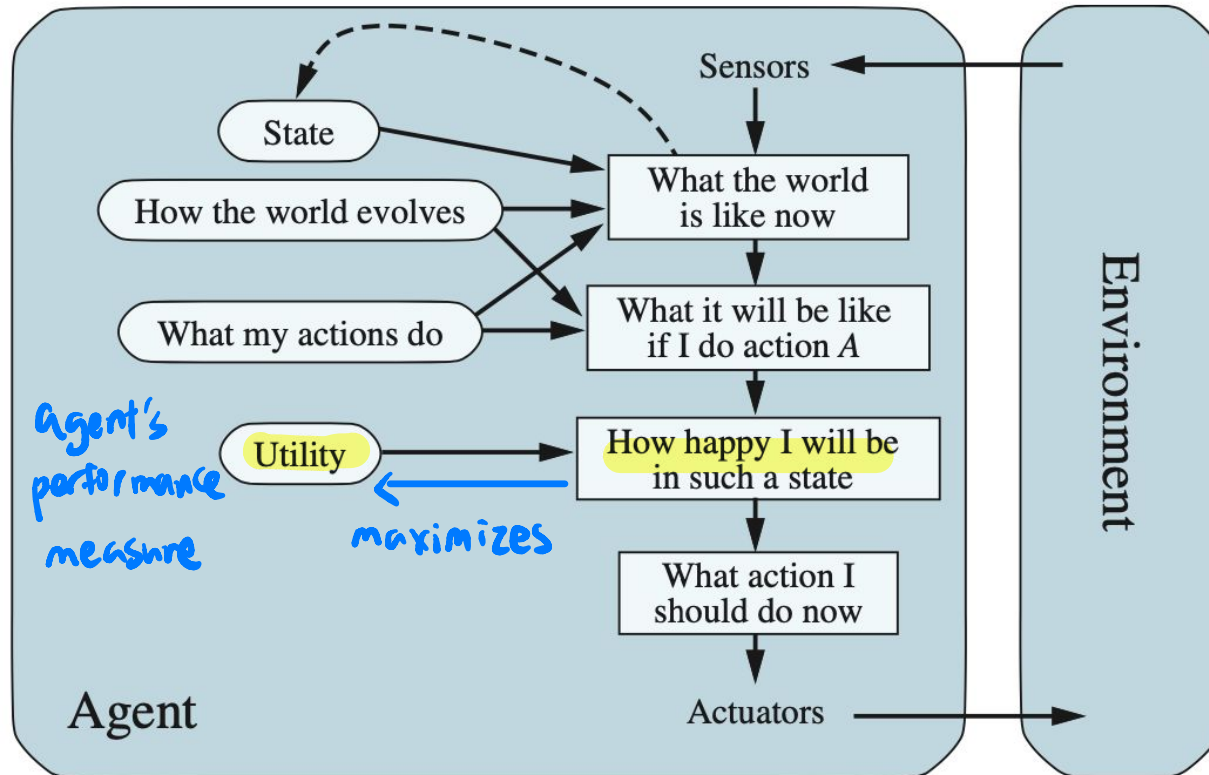


Goal-based Agent





Utility-based Agent



Agent Types

- Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents
- ⇒ Learning agents



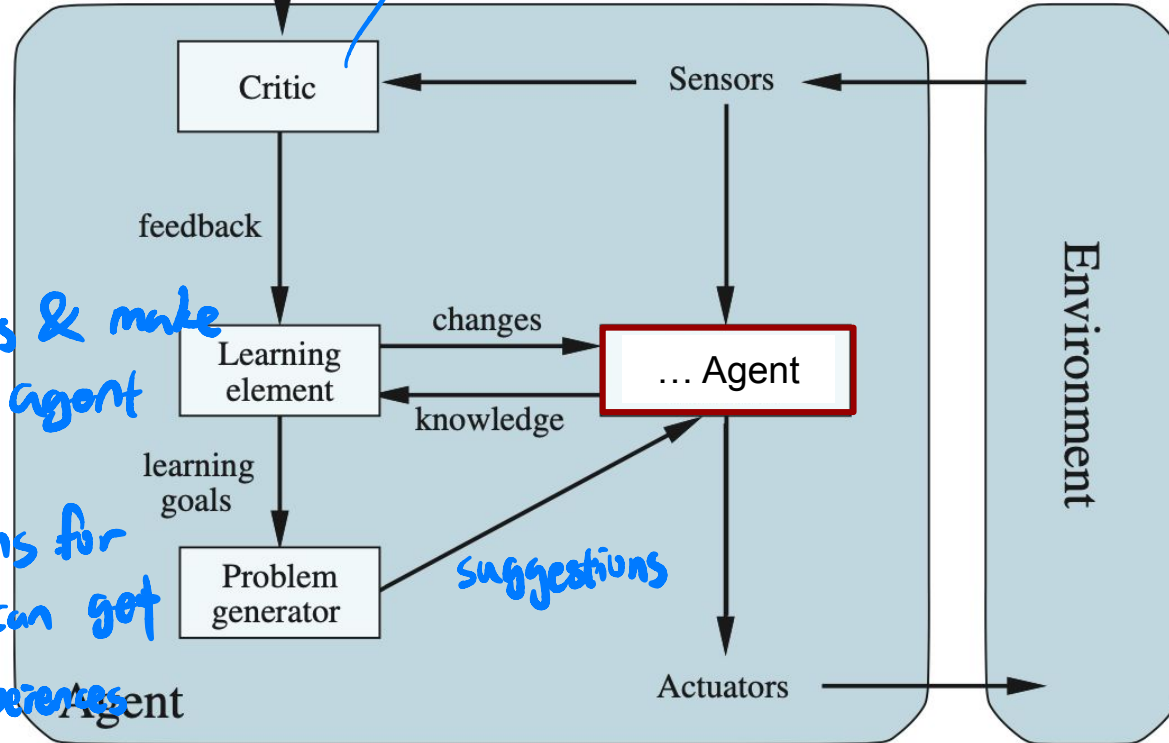
Generality



Learning Agent

Finds out how well the agent is performing & provides feedback based on fixed standards

Performance standard (be fixed)



receive feedbacks & make changes to the agent

Suggesting actions for the agent so it can get some fresh experiences