



Competitive Environments

Games vs. Search Problems

- Unpredictable opponent
 - Solution is a **strategy** specifying a move for every possible opponent reply
- Time limits

Search Topics

- Search problems (Ch. 3)
- Uninformed search (Ch. 3)
- Informed search (Ch. 3)
- Local search (Ch. 4)
- **Adversarial search (Ch. 6)**
- Constraint Satisfaction Problems (CSPs) (Ch. 5)

Adversarial Search

L.-Y. Wei

Spring 2024

Game Types


- The game most commonly studied are

- Deterministic
- Two-player
- Turn-taking
- Perfect information
 - Fully observable

(Imperfect information, e.g., poker and bridge)

- Zero-sum games 零和賽局
 - No “win-win” outcome
(i.e., what is good for one player is just as bad for the other)


Example: Tic-Tac-Toe



Two-Player Zero-Sum Games

- Initial state, S_0
- TO-MOVE(s)
 - Returns which **player** to move in state s
- ACTIONS(s)
 - Returns the set of legal moves in state s
- RESULT(s,a), transition model
 - Defines the state resulting from taking action a in state s
- IS-TERMINAL(s), terminal test
 - Checks if state s is a terminal state
- **UTILITY**(s,p), utility function/objective function/
 - Defines the final numeric value to player p when the game ends in terminal state s

Example: Initial State




Example: TO-MOVE(s)

$$\text{TO-MOVE}\left(\begin{array}{|c|c|c|}\hline & & \\ \hline & & \\ \hline & & \\ \hline\end{array}\right) = X$$

$$\text{TO-MOVE}\left(\begin{array}{|c|c|c|}\hline & & \\ \hline & X & \\ \hline & & \\ \hline\end{array}\right) = O$$

換誰？

Example: ACTIONS(s)



Example: RESULT(s,a)

RESULT(

O	X	
X	X	O
O		X

,

		O

) =

O	X	O
X	X	O
O		X

Example: IS-TERMINAL(s)

IS-TERMINAL(

O	X	
X	X	O
O		X

) = false

IS-TERMINAL(

X	O	X
	X	
X	O	O

) = true

Example: UTILITY(s, p)

UTILITY(

X	O	X
	X	
X	O	O

, X) =

UTILITY(

X	O	X
O	O	X
X	X	O

, X) =

UTILITY(

X	O	X
	O	X
	O	

, X) =

Example: UTILITY(s,p)

UTILITY(

X	O	X
	X	
X	O	O

, X) = 1

X wins

UTILITY(

X	O	X
O	O	X
X	X	O

, X) = 0

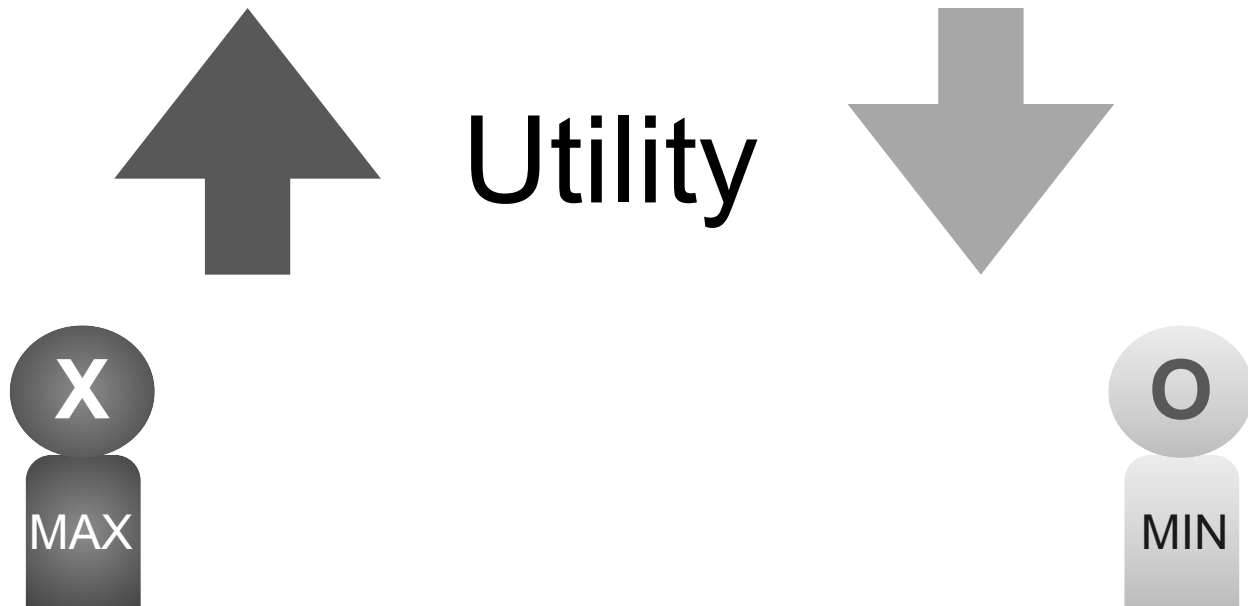
draw

UTILITY(

X	O	X
	O	X
	O	

, X) = -1


O wins



Al



Example: Game Tree for Tic-Tac-Toe




MAX wants to find a sequence of actions leading to win.

Optimal Decisions in Games?


Assume both players play optimal.

- When it is MAX's turn to move,
MAX prefers to move to a state of maximum value
- MIN prefers a state of minimum value
(i.e., minimum value for MAX and thus maximum value for MIN)

Example: Game Tree for Tic-Tac-Toe



Example: Game Tree for Tic-Tac-Toe



A General Algorithm for Perfect-Information Games

- **Minimax Search**

- Given a game tree, the **minimax value** of each state in the tree is the utility for MAX of being in that state, i.e.,

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s, \text{MAX}) & \text{if IS-TERMINAL}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if TO-MOVE}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if TO-MOVE}(s) = \text{MIN} \end{cases}$$

- Assume both players play optimally from the state to the end of the game
- Use the **depth-first exploration** of the game tree to derive the optimal strategy by working out the minimax value of each state in the tree


Example: Minimax Search for Two-Ply Game Trees

MAX



Example: Minimax Search for Two-Ply Game Trees

MAX




Example: Minimax Search for Two-Ply Game Trees


MAX

MIN


Terminal




Example: Minimax Search for Two-Ply Game Trees




Example: Minimax Search for Two-Ply Game Trees




Example: Minimax Search for Two-Ply Game Trees



Example: Minimax Search for Two-Ply Game Trees



Example: Minimax Search for Two-Ply Game Trees




Example: Minimax Search for Two-Ply Game Trees

MAX

MIN

Terminal




Example: Minimax Search for Two-Ply Game Trees


MAX

MIN

Terminal




Example: Minimax Search for Two-Ply Game Trees



Example: Minimax Search for Two-Ply Game Trees

△ MAX

▽ MIN




Minimax Search Algorithm

```
function MINIMAX-SEARCH(game, state) returns an action  
  player  $\leftarrow$  game.TO-MOVE(state)  
  value, move  $\leftarrow$  MAX-VALUE(game, state)  
  return move
```

```
function MAX-VALUE(game, state) returns a (utility, move) pair  
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null  
   $v$ , move  $\leftarrow -\infty$   
  for each  $a$  in game.ACTIONS(state) do  
     $v2$ ,  $a2$   $\leftarrow$  MIN-VALUE(game, game.RESULT(state,  $a$ ))  
    if  $v2 > v$  then  
       $v$ , move  $\leftarrow v2$ ,  $a$   
  return  $v$ , move
```

```
function MIN-VALUE(game, state) returns a (utility, move) pair  
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null  
   $v$ , move  $\leftarrow +\infty$   
  for each  $a$  in game.ACTIONS(state) do  
     $v2$ ,  $a2$   $\leftarrow$  MAX-VALUE(game, game.RESULT(state,  $a$ ))  
    if  $v2 < v$  then  
       $v$ , move  $\leftarrow v2$ ,  $a$   
  return  $v$ , move
```



Properties of Minimax

Note. b is the branching factor and m is the depth limit

- Complete?
 - Yes, if the tree is finite
- Optimal cost?
 - Yes
- Time complexity?
 - $O(b^m)$
- Space complexity?
 - $O(bm)$ (depth-first exploration)

) DFS

Example of Three Players


to move

A

B

C

A



Example of Three Players

to move

A

B


C

A

(1, 2, 6)

A B C

Note. Vectors give the utility of the state from each player's viewpoint



Example of Three Players


to move

A

B

C

A



$\begin{matrix} A & B & C \end{matrix}$

Note. Vectors give the utility of the state from each player's viewpoint

Example of Three Players


to move

A

B

C

A



A B C

Note. Vectors give the utility of the state from each player's viewpoint

Example of Three Players


to move

A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint

Example of Three Players


to move

A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint


Example of Three Players

to move
A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint


Example of Three Players

to move
A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint

Example of Three Players

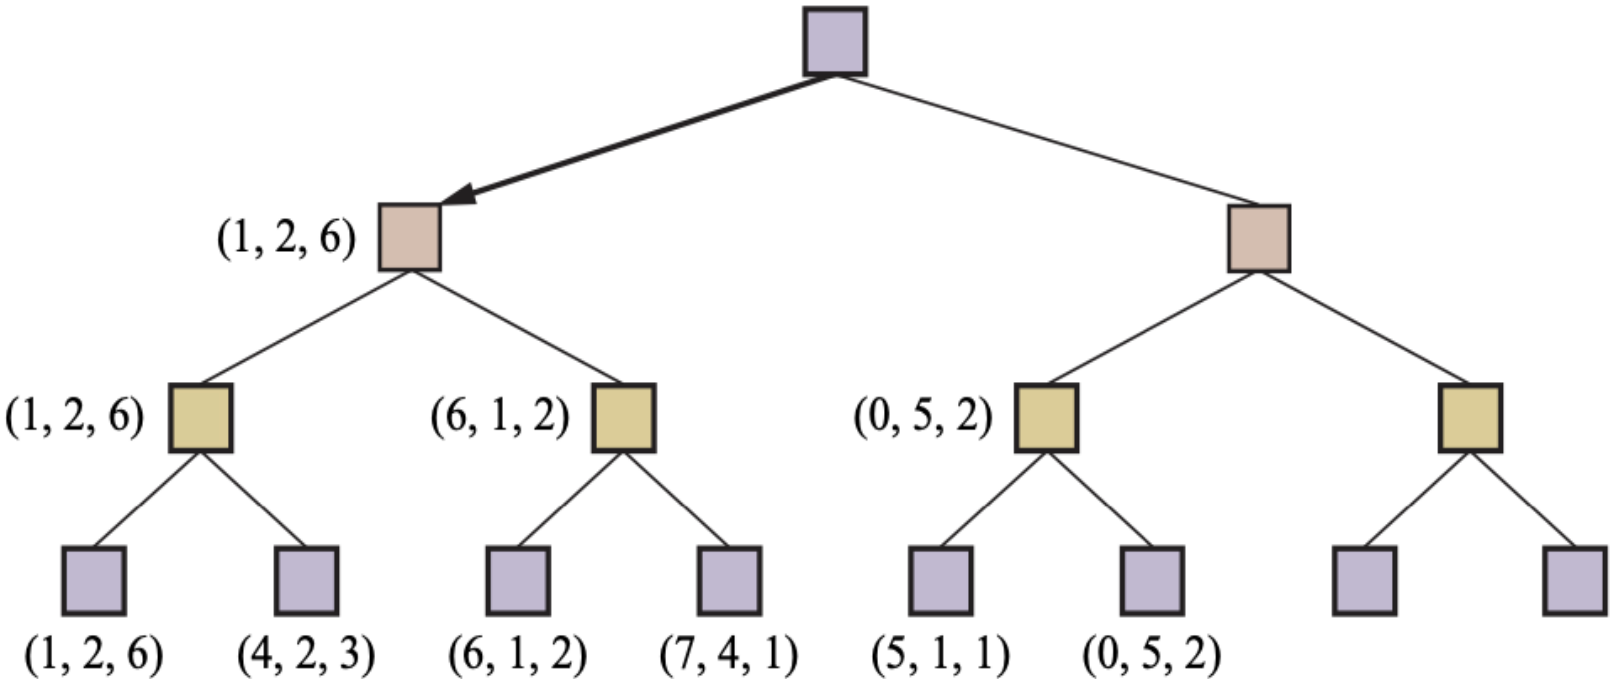
to move

A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint

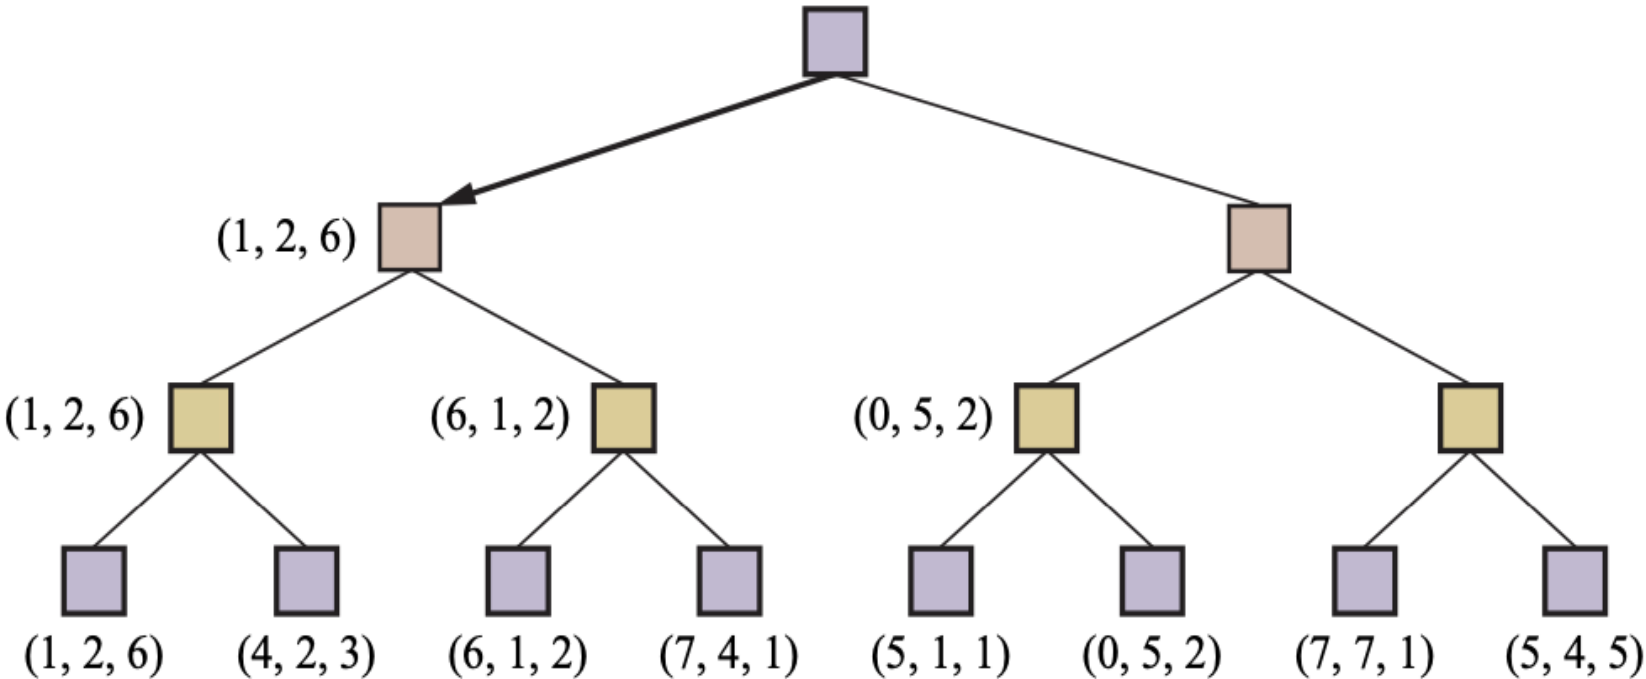
Example of Three Players

to move
A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint

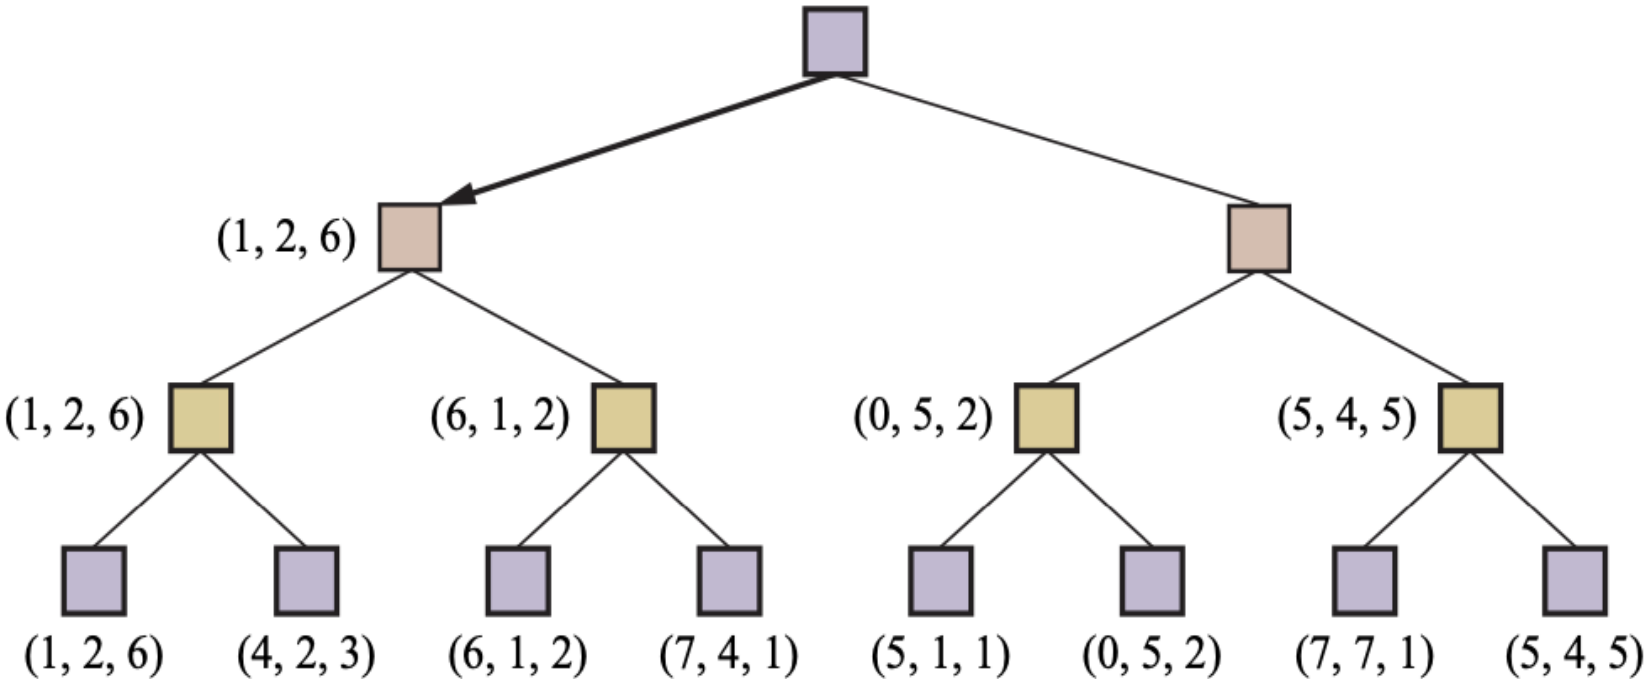
Example of Three Players

to move
A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint

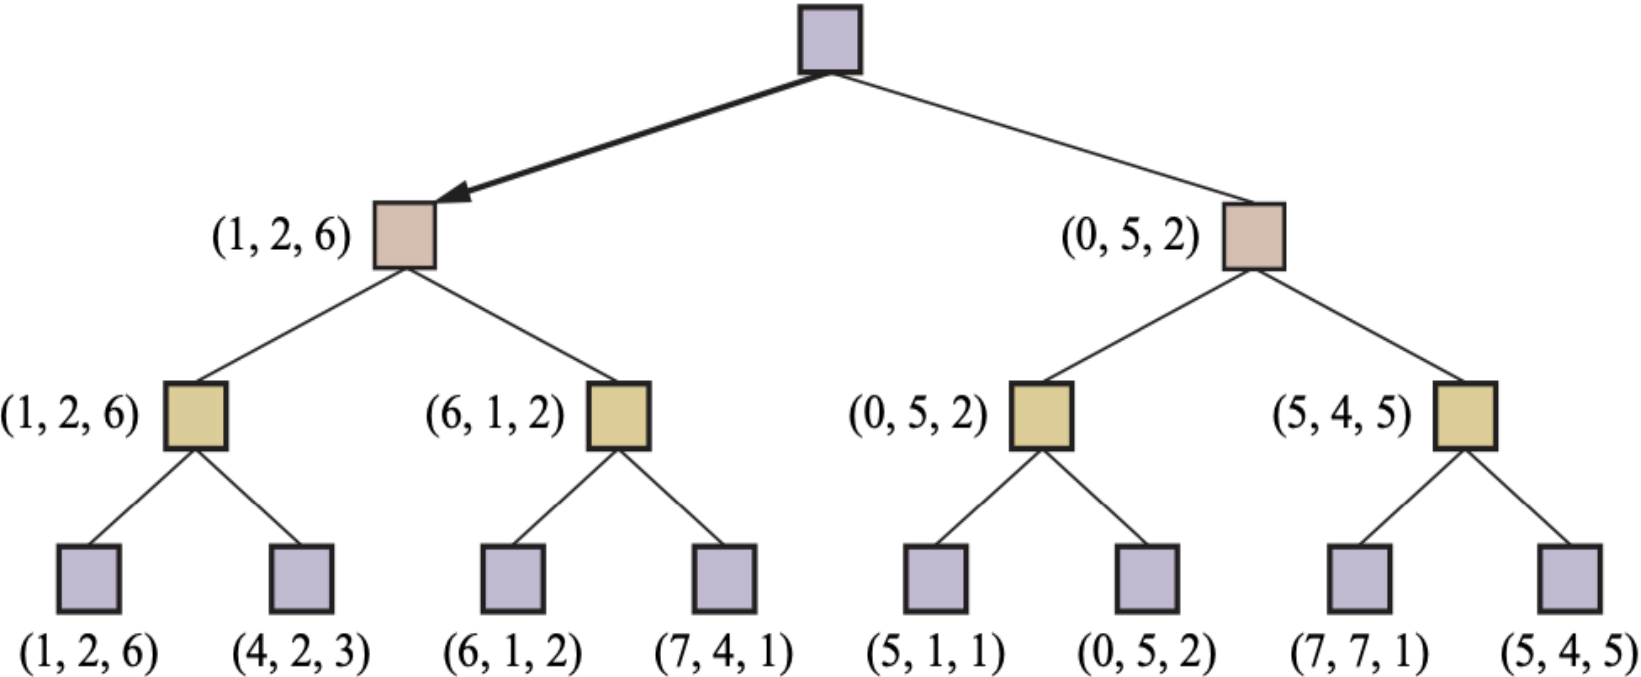
Example of Three Players

to move
A

B

C

A



Note. Vectors give the utility of the state from each player's viewpoint

Example of Three Players

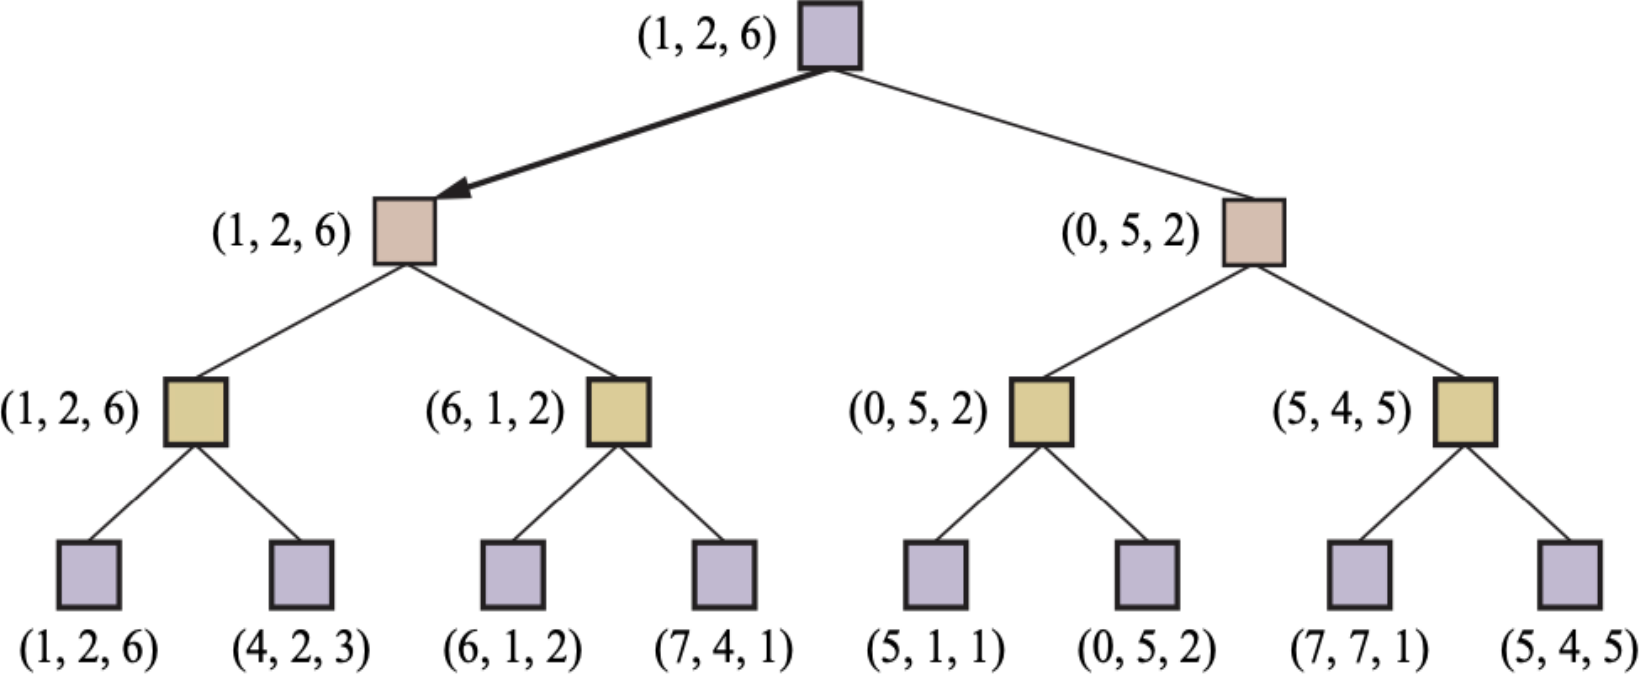
to move

A

B

C

A




Note. Vectors give the utility of the state from each player's viewpoint

Performance?

Observation: Minimax Search

MAX

MIN



Expand All Nodes


Improvement ?

Expand ~~All~~ Nodes
Partial

Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX


MIN



Pruning Technique?

MAX

MIN



Alpha-Beta Pruning

- α = the value of the best choice (i.e., highest-value) we have found so far at any choice point along the path for **MAX**
(Think: α = “at least”)
- β = the value of the best choice (i.e., lowest-value) we have found so far at any choice point along the path for **MIN**
(Think: β = “at most”)


Player

Opponent

•
•
•


Player

Opponent




Alpha-Beta Pruning

- α = the value of the best choice (i.e., highest-value) we have found so far at any choice point along the path for **MAX**
(Think: α = “at least”)
- β = the value of the best choice (i.e., lowest-value) we have found so far at any choice point along the path for **MIN**
(Think: β = “at most”)



Alpha-Beta Pruning

- α = the value of the best choice (i.e., highest-value) we have found so far at any choice point along the path for **MAX**
(Think: α = “at least”)
- β = the value of the best choice (i.e., lowest-value) we have found so far at any choice point along the path for **MIN**
(Think: β = “at most”)




Example: Pruning




Example: Pruning

$$[-\infty, +\infty] \triangle A$$


Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning




Example: Pruning



Example: Pruning



Example: Pruning



Alpha-Beta Search Algorithm

function ALPHA-BETA-SEARCH(*game*, *state*) **returns** an action
 player \leftarrow *game*.TO-MOVE(*state*)
 value, *move* \leftarrow MAX-VALUE(*game*, *state*, $-\infty$, $+\infty$)
 return *move*

Alpha-Beta Search Algorithm

```
function MAX-VALUE(game, state,  $\alpha$ ,  $\beta$ ) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
   $v \leftarrow -\infty$ 
  for each a in game.ACTIONS(state) do
     $v2, a2 \leftarrow$  MIN-VALUE(game, game.RESULT(state, a),  $\alpha$ ,  $\beta$ )
    if  $v2 > v$  then
       $v, move \leftarrow v2, a$ 
       $\alpha \leftarrow$  MAX( $\alpha$ ,  $v$ )
    if  $v \geq \beta$  then return  $v, move$ 
return  $v, move$ 
```

Alpha-Beta Search Algorithm

```
function MIN-VALUE(game, state,  $\alpha$ ,  $\beta$ ) returns a (utility, move) pair
if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
 $v \leftarrow +\infty$ 
for each a in game.ACTIONS(state) do
     $v2, a2 \leftarrow$  MAX-VALUE(game, game.RESULT(state, a),  $\alpha, \beta$ )
    if  $v2 < v$  then
         $v, move \leftarrow v2, a$ 
         $\beta \leftarrow \text{MIN}(\beta, v)$ 
    if  $v \leq \alpha$  then return  $v, move$ 
return  $v, move$ 
```

Weaknesses for Games

- Large branching factor
- Difficult to define a heuristic function

Monte Carlo Tree Search (MCTS)

- Idea: It estimates **the average utility** over a number of **simulations** of complete games starting from the state
 - A simulation (also called a **playout** or **rollout**) chooses moves first for one player, then for the other, repeating until a terminal position is reached
 - For games, “average utility” is the same as “win percentage”

Rollout (S_i):


loop forever:

if S_i is a terminal state

return $\text{value}(S_i)$

$A_i = \text{random}(\text{available actions}(S_i))$

$S_i = \text{simulate}(A_i, S_i)$



Simulation Policy

- From what positions do we start the simulations?
- How many simulations do we allocate to each position?

Pure Monte Carlo Search

- Do N simulations starting from the current state of the game
- Track which of the possible moves from the current position has the highest win percentage
- Issue
 - Optimal play
 - Increase $N \rightarrow$ Computational resources \uparrow

Selection Policy

- Focus on the important parts of the game tree and balance two factors
 - **Exploration** of states that had few simulations
 - **Exploitation** of states that have done well in past simulations to get a **more accurate estimate** of their value

Monte Carlo Tree Search

Maintain a search tree and grow it on each iteration of the following four steps:


- Selection
- Expansion
- Simulation
- Back-propagation

Step 1: Selection

- Starting at the root node of the search tree, we choose a move (e.g., best win percentage), leading to a successor node, and repeat that process, moving down the tree to a leaf


- White: won 37 out of 100 playouts done so far

- Black: won 60/79 playouts




Step 2: Expansion

- We grow the search tree by generating a new child of the selected node




Step 3: Simulation

- We perform a playout from the newly generated child node




Step 4: Back-propagation

- We now use the result of the simulation to update all search tree nodes going up to the root



Update



Upper Confidence Bounds Applied to Trees (UCT)

- The effective selection policy ranks each possible move based on an upper confidence bound formula UCB1 as below

$$UCB1(n) = \underbrace{\frac{U(n)}{N(n)}}_{\text{Exploitation}} + C \times \underbrace{\sqrt{\frac{\log N(\text{PARENT}(n))}{N(n)}}}_{\text{Exploration}}$$

average win rate

how many times have the parent node been visited?

- $U(n)$: the total utility of all playouts that went through node n
 - $N(n)$: the number of playouts through node n
 - $\text{Parent}(n)$: the parent node of n in the tree
 - C : a constant that balances exploitation and exploration
- been tried?*
- how many times has this node*

UCT MCTS Algorithm

function MONTE-CARLO-TREE-SEARCH(*state*) **returns** *an action*

tree \leftarrow NODE(*state*)

while IS-TIME-REMAINING() **do**

leaf \leftarrow SELECT(*tree*)

child \leftarrow EXPAND(*leaf*)

result \leftarrow SIMULATE(*child*)

BACK-PROPAGATE(*result*, *child*)

return the move in ACTIONS(*state*) whose node has highest number of playouts

- UCB1 formula ensures that the node with most playouts is almost always the node with the highest win percentage, because the selection process favors win percentage more and more as the number of playouts goes up

\therefore exploration term $\rightarrow 0$ when $n \rightarrow \infty$

Course Topics

- Intelligent agents (AIMA Ch. 2)
- Search (AIMA Ch. 3, 4, 6, 5)
- Reasoning (AIMA Ch. 7-9, 12-15)
- Machine learning (AIMA Ch. 19-20)
- Deep learning (AIMA Ch. 22)
- Natural language processing (AIMA Ch. 24)
- Generative AI (Optional)